

# Cognitive Services

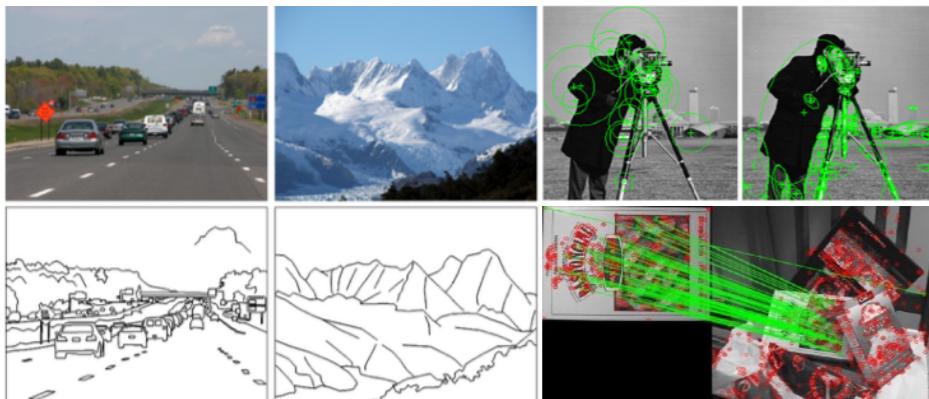
*LM Informatica & Data Science - 2<sup>nd</sup> semester - 6 CFU*

Foundations: edges, local invariant features, descriptors

Lamberto Ballan

# What we will learn today?

- Edge detection
- Image gradients
- Local invariant features
- DoG and SIFT descriptors



Background reading:  
*Forsyth and Ponce, Computer Vision, Chapter 8*  
*D.G. Lowe, SIFT paper, IJCV 2014*

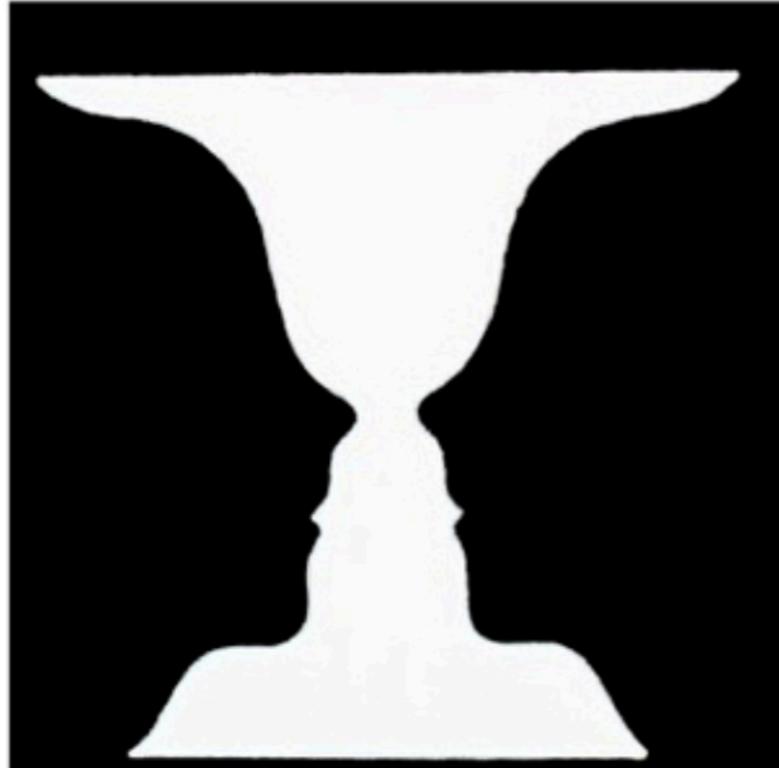
# Image filtering recap

- Compute a function of the local neighborhood at each pixel in the image
  - The function is specified by a “filter” saying how to combine values from neighbors
- Applications of image filtering:
  - extract information (edges, corners, blobs, ...)
  - detect patterns (template matching)
  - de-noising, super-resolution, in-painting





# What's an edge?



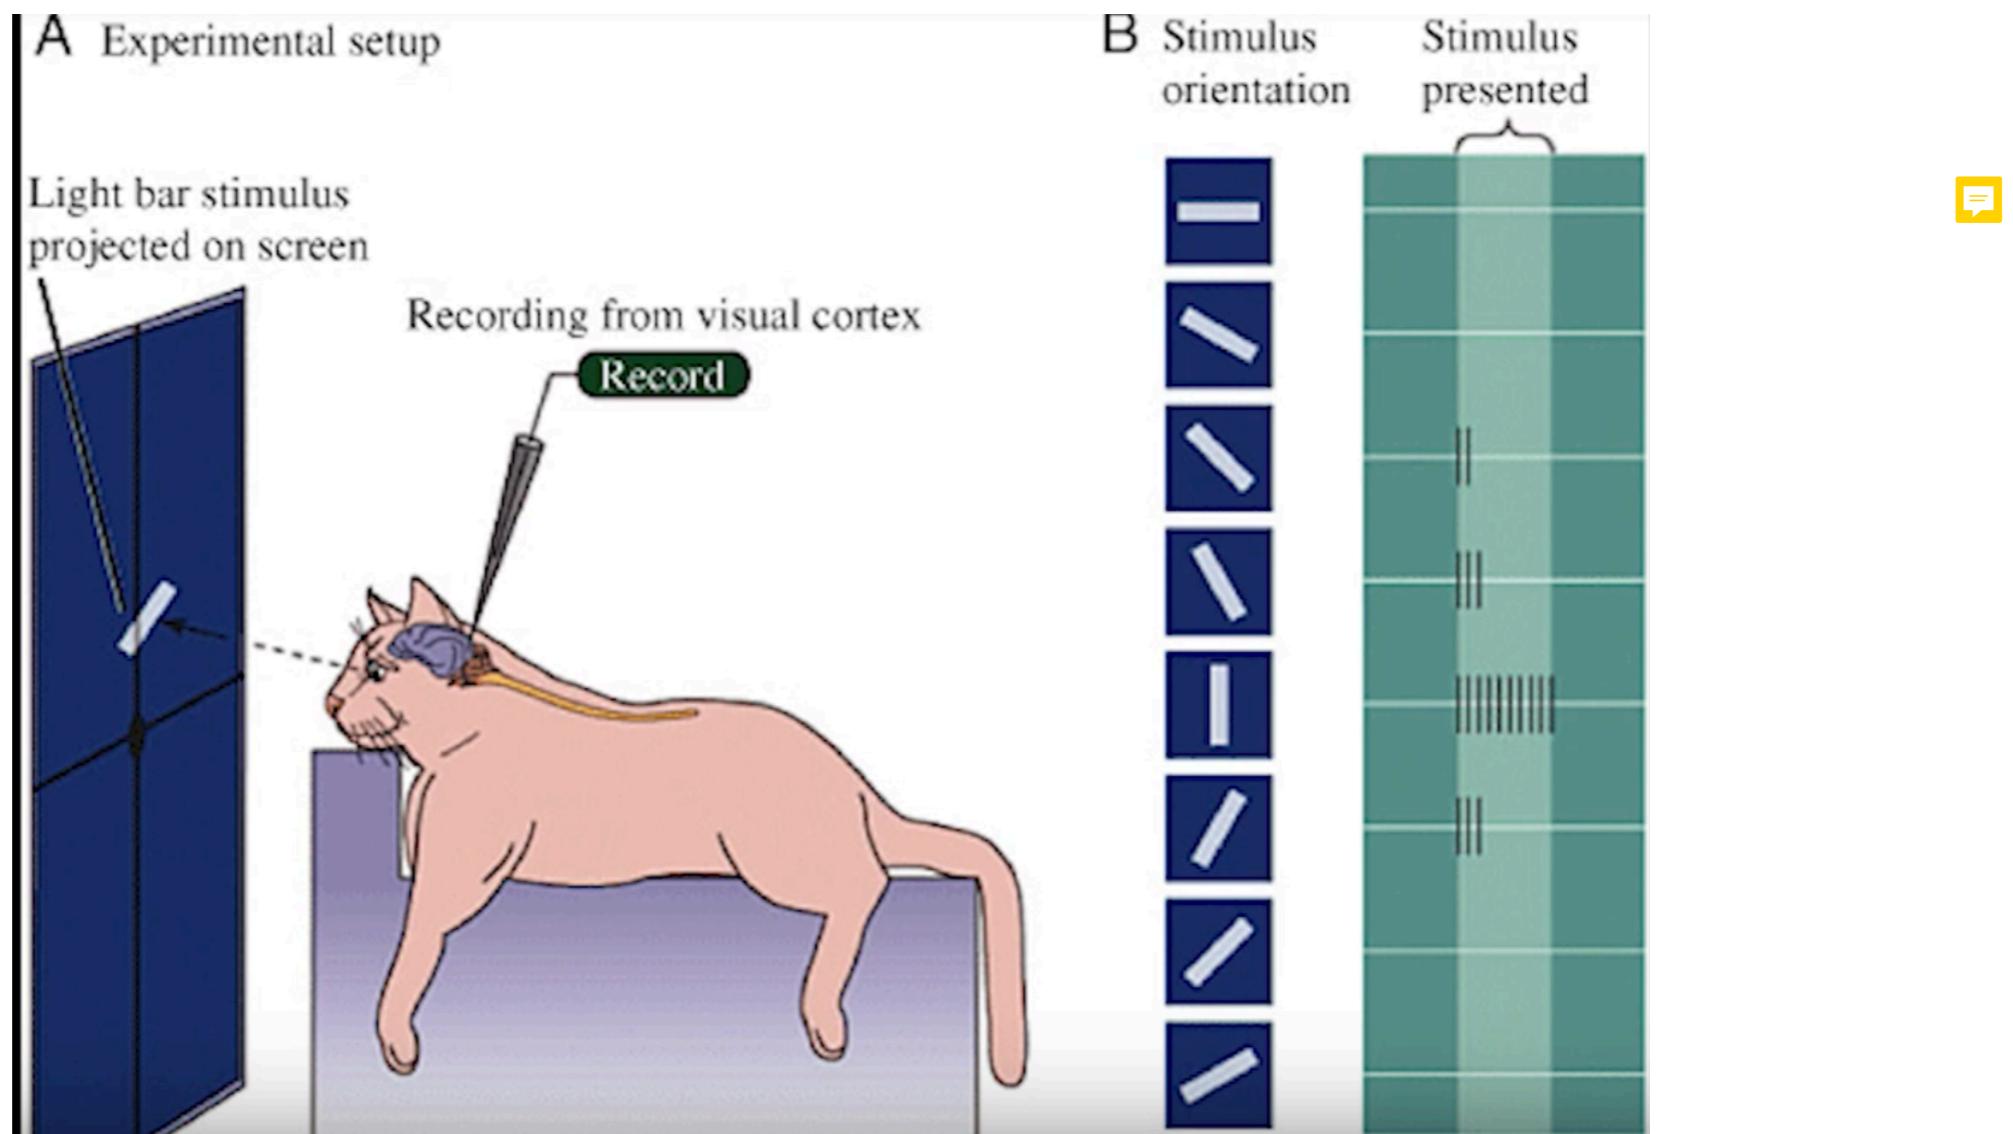
Cave painting (Chauvet, France, ~30,000 BC)



Shen Zhou's "Poet on a mountain top" (China, ~1,500 AD)

# What's an edge?

- We know edges are special from human (mammalian) vision studies



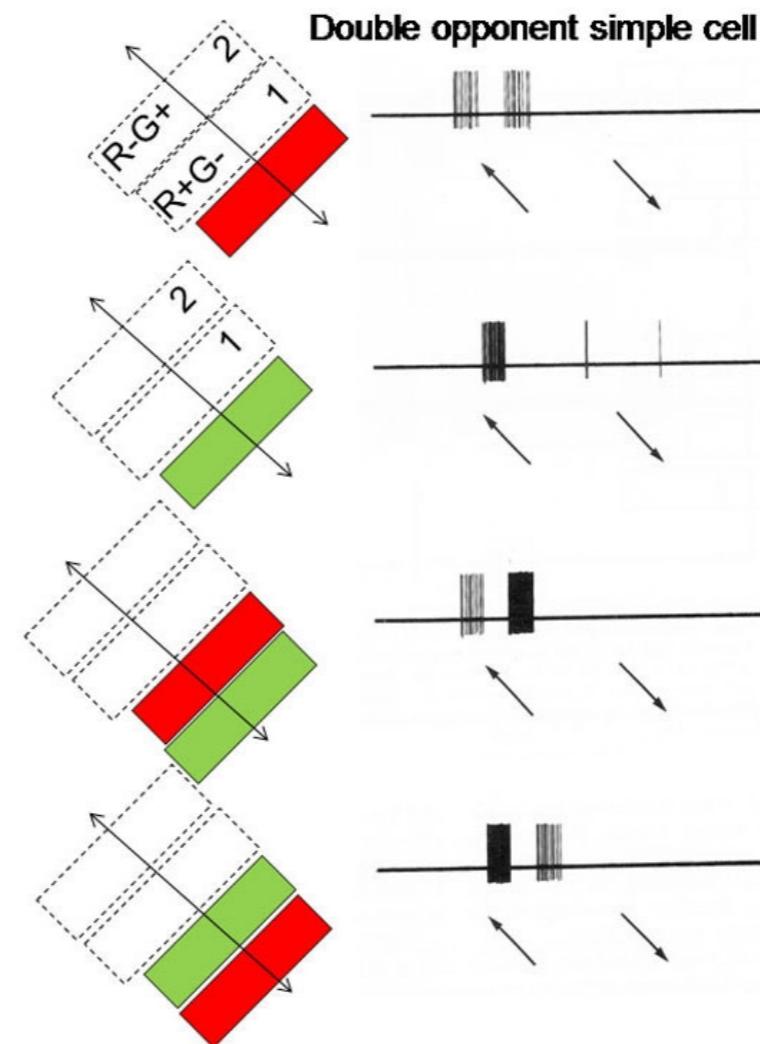
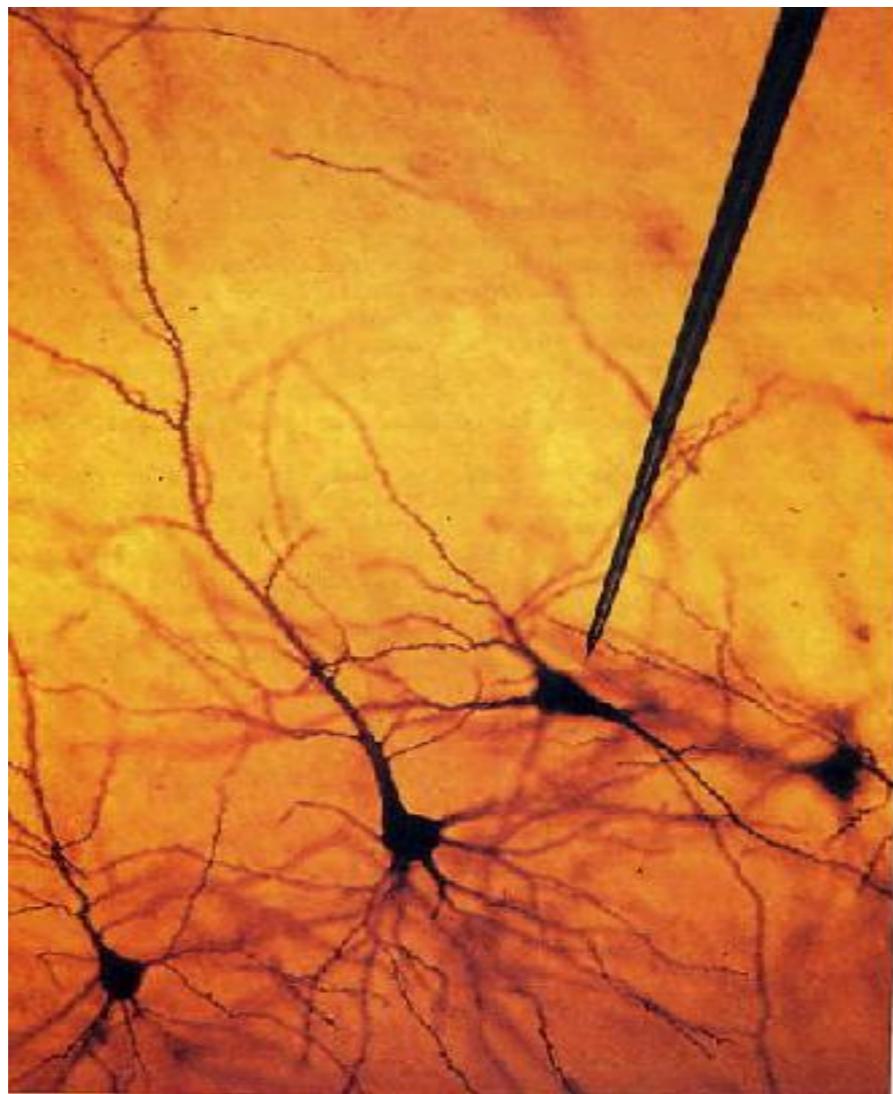
Hubel and Wiesel, 1960s

# What's an edge?



- We know edges are special from human (mammalian) vision studies

Hubel and Wiesel, 1960s



# Edge detection

- **Goal:** Identify sudden changes (discontinuities) in an image
  - ▶ Intuitively, most semantic and shape information from the image can be encoded in the edges
  - ▶ More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)

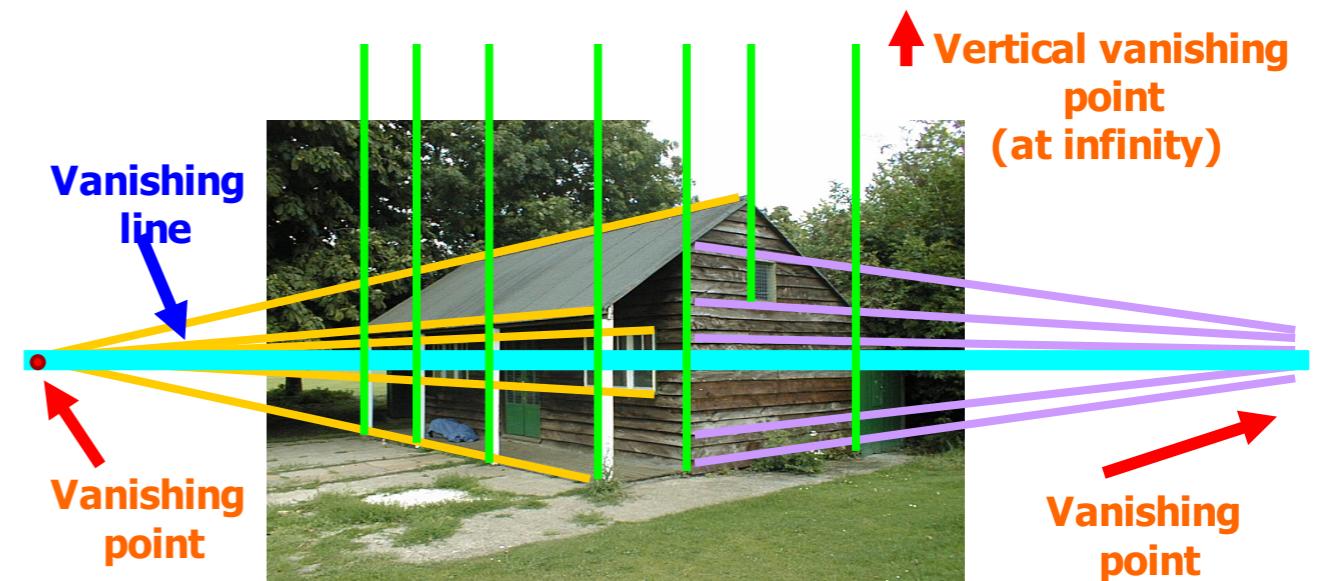


# Why do we care about edges?

- Extract information, recognize objects

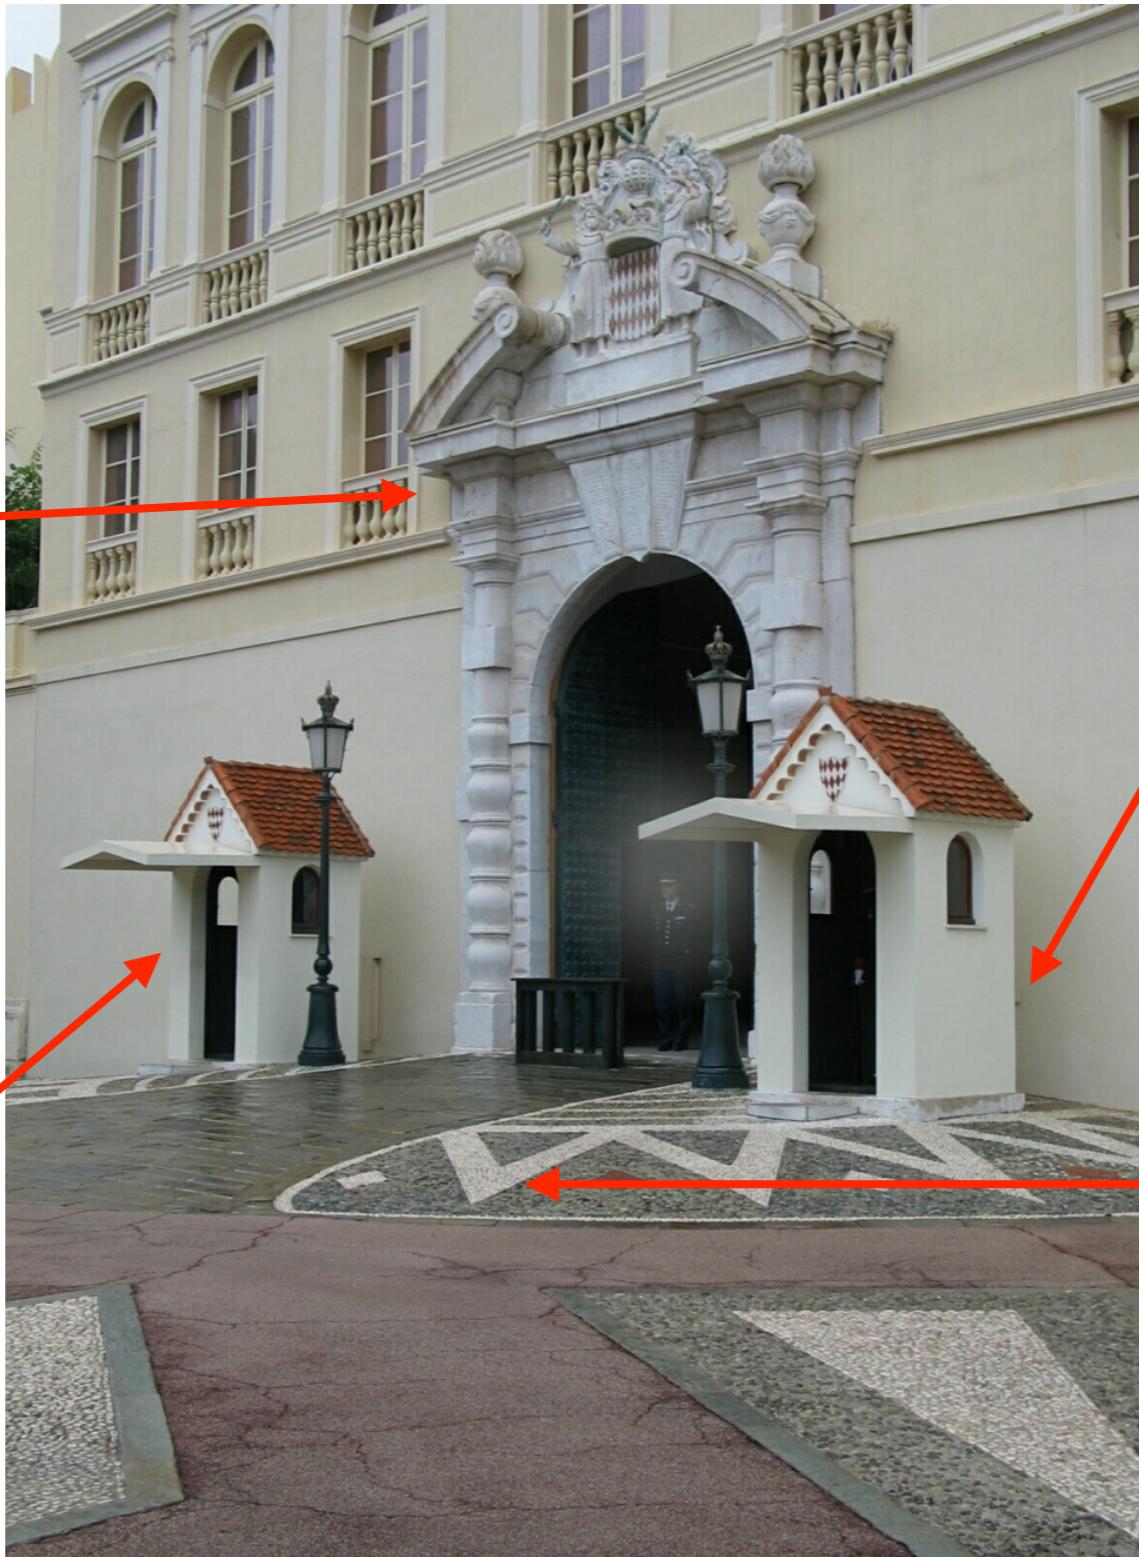


- Recover geometry and viewpoint

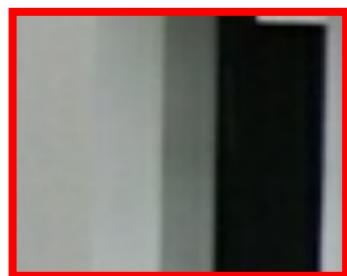


# What causes an edge?

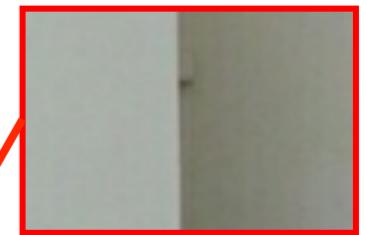
Illumination discontinuity:  
cast shadows



Change in surface  
orientation: shape



Depth discontinuity:  
object boundary



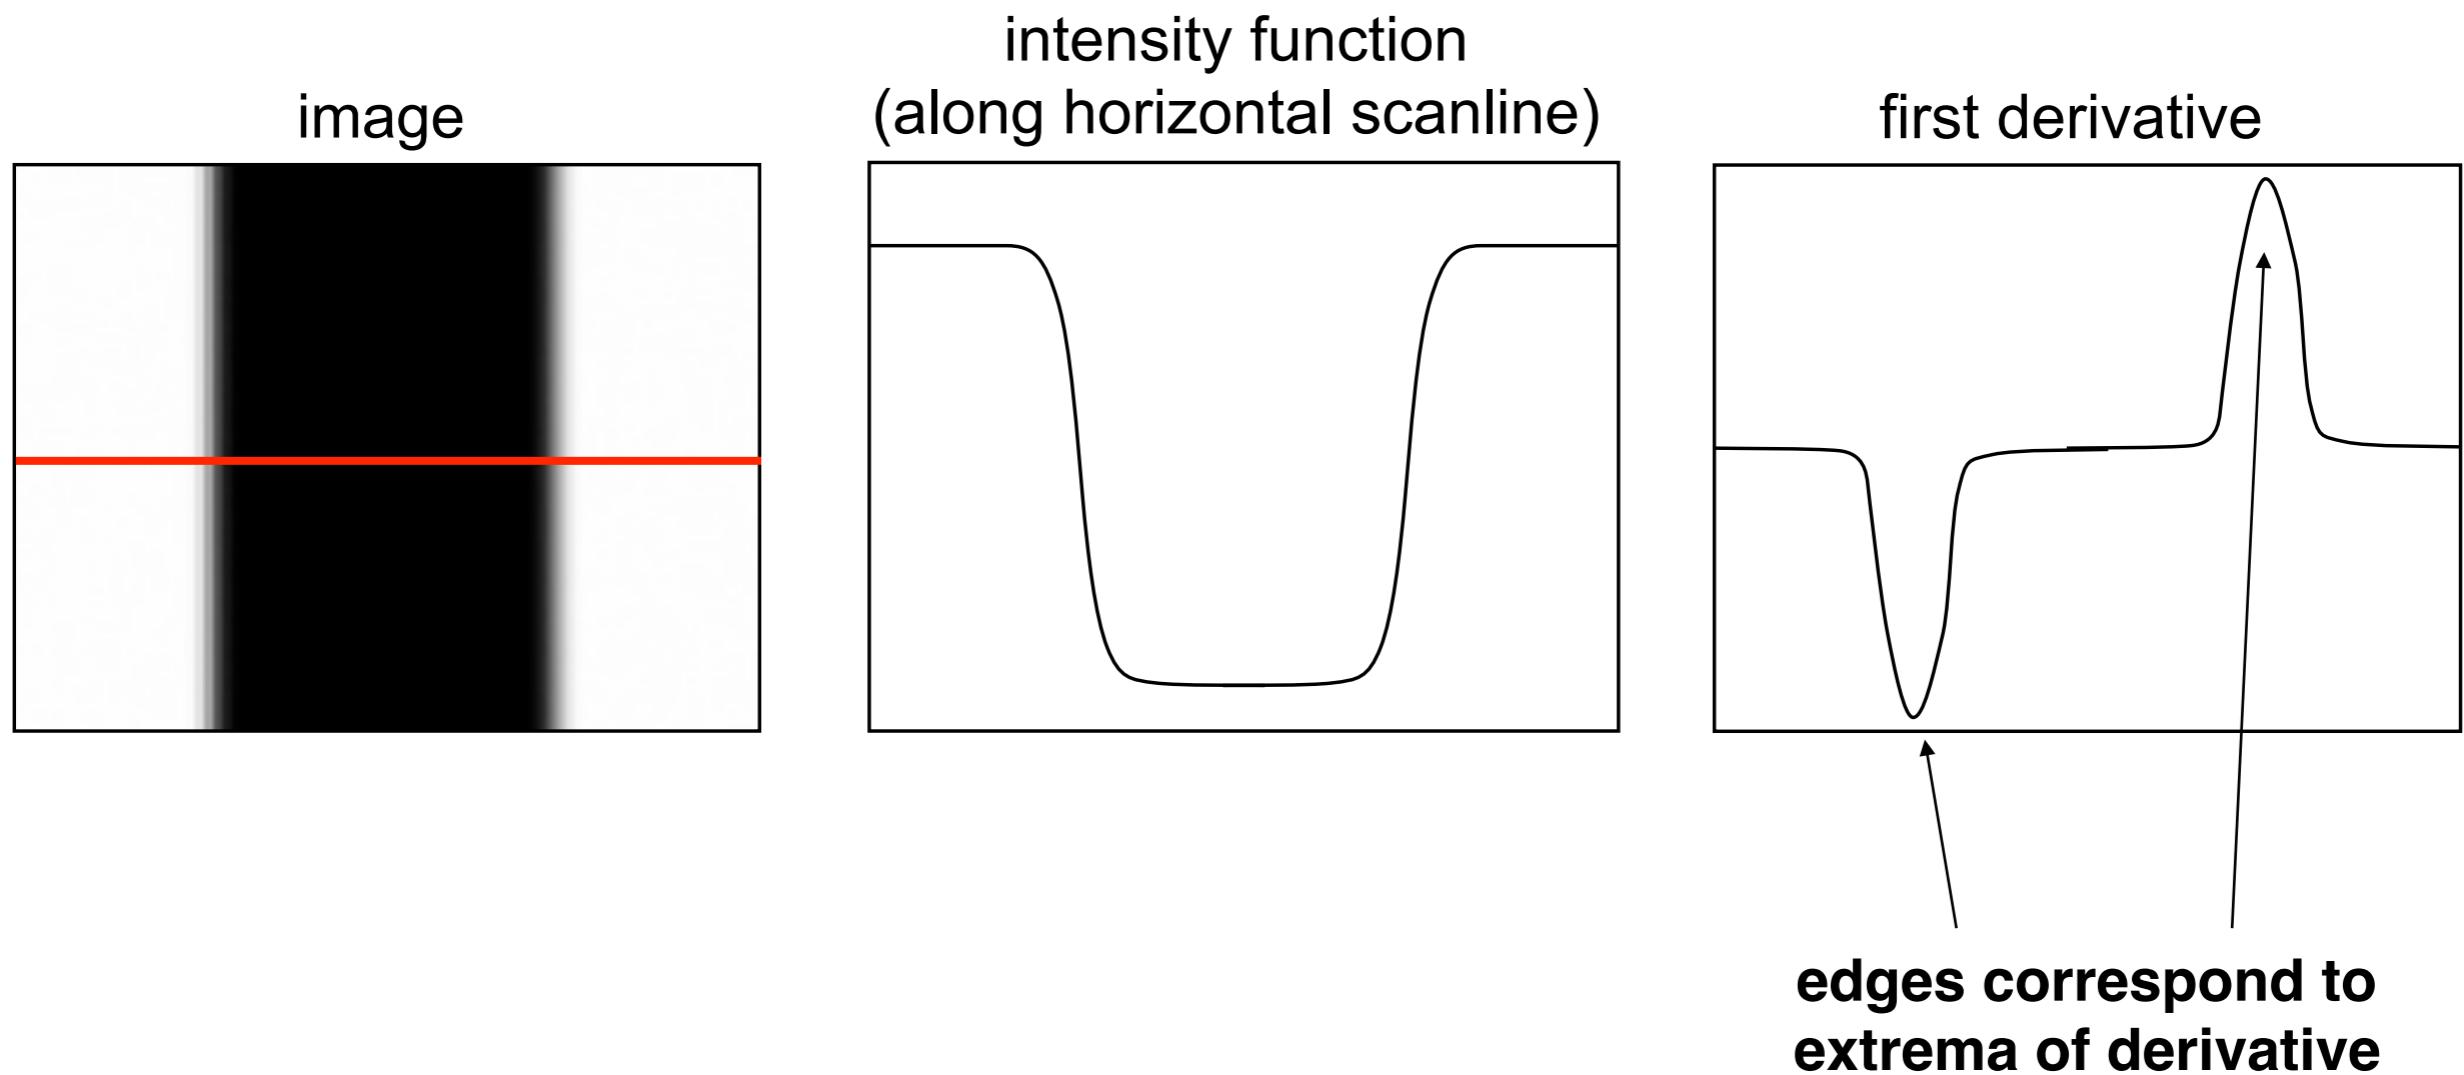
Surface color  
discontinuity



# Edges and image gradients



- An edge is a place of rapid change in the image intensity function



# Derivatives: a quick recap

- Derivative in 1D:

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x$$



*A simple example:*

$$y = x^2 + x^4$$
$$\frac{dy}{dx} = 2x + 4x^3$$

- Discrete derivative in 1D:

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x)$$



$$\frac{df}{dx} = \frac{f(x) - f(x - 1)}{1} = f'(x)$$

*we can approximate using finite differences*

$$\frac{df}{dx} = f(x) - f(x - 1) = f'(x)$$

# Derivatives: a quick recap



- Types of Discrete derivative in 1D:

*Filter*

- Backward:  $\frac{df}{dx} = f(x) - f(x-1) = f'(x)$  [ 0 1 -1 ]
- Forward:  $\frac{df}{dx} = f(x) - f(x+1) = f'(x)$  [-1 1 0]
- Central:  $\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$  [ 1 0 -1 ]

# Derivatives: a quick recap

- 1D discrete derivate example:



$$f(x) = 10 \ 15 \ 10 \ 10 \ 25 \ 20 \ 20 \ 20$$

$$f'(x) = \boxed{0} \ 5 \ -5 \ 0 \ 15 \ -5 \ 0 \ \boxed{0}$$

# 2D discrete derivative: image gradient



Given function

$$f(x, y)$$

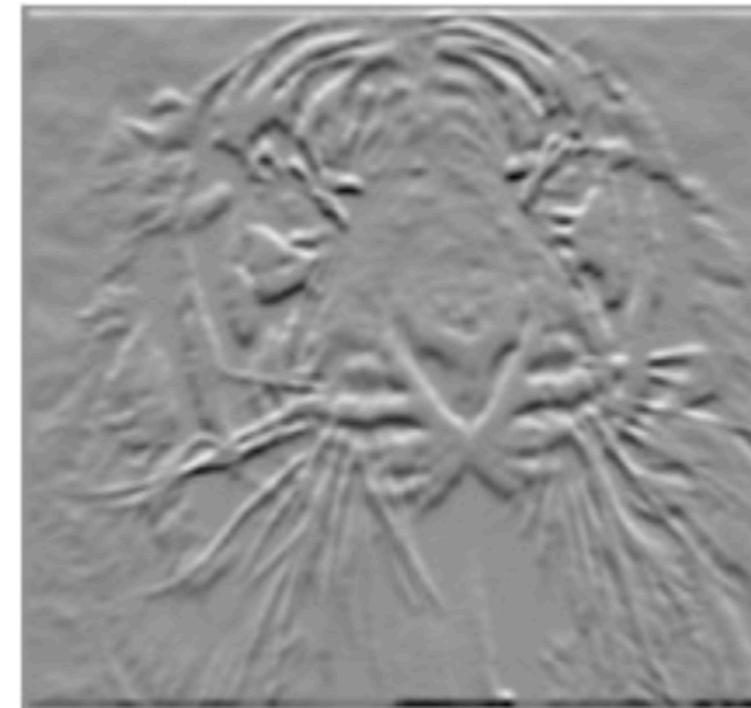
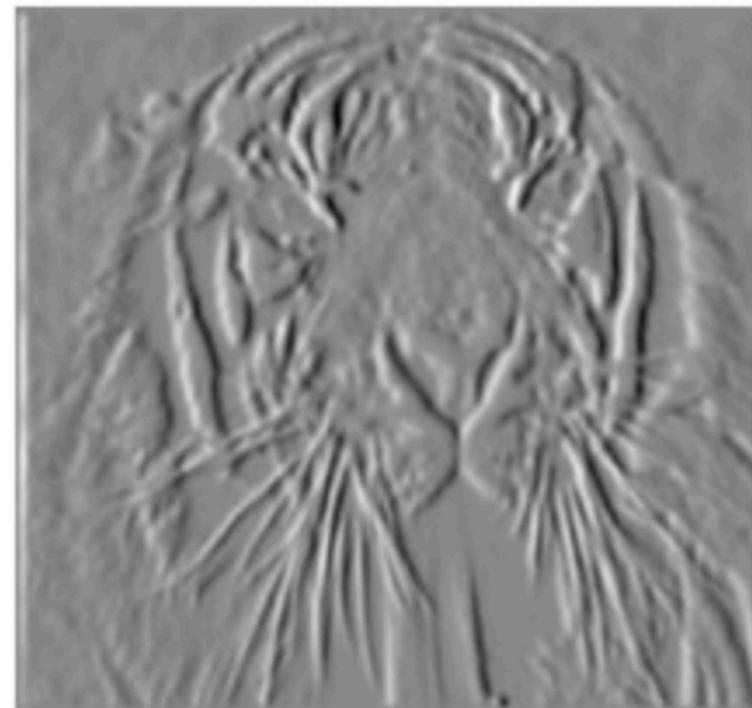
Gradient vector

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} \\ \frac{\partial f(x, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$$

*Which one is the gradient in the x-direction? How about y-direction?*

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$

x-direction



$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$

y-direction

# 2D discrete derivate example



$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \boxed{\text{Filter}} = ?$$
$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

What happens when we apply this filter?

# 2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \text{Filter} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

*Filter*

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} =$$

# 2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \boxed{\text{Filter}} = ?$$

The filter is defined as:

$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

What happens when we apply this filter?

# 2D discrete derivate example

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix} * \text{Filter} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

# Image gradient



- The gradient of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The first image shows a horizontal gradient with a red arrow pointing right, labeled  $\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$ . The second image shows a vertical gradient with a red arrow pointing down, labeled  $\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$ . The third image shows a diagonal gradient with a red arrow pointing up and to the right at an angle  $\theta$ , labeled  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$ .

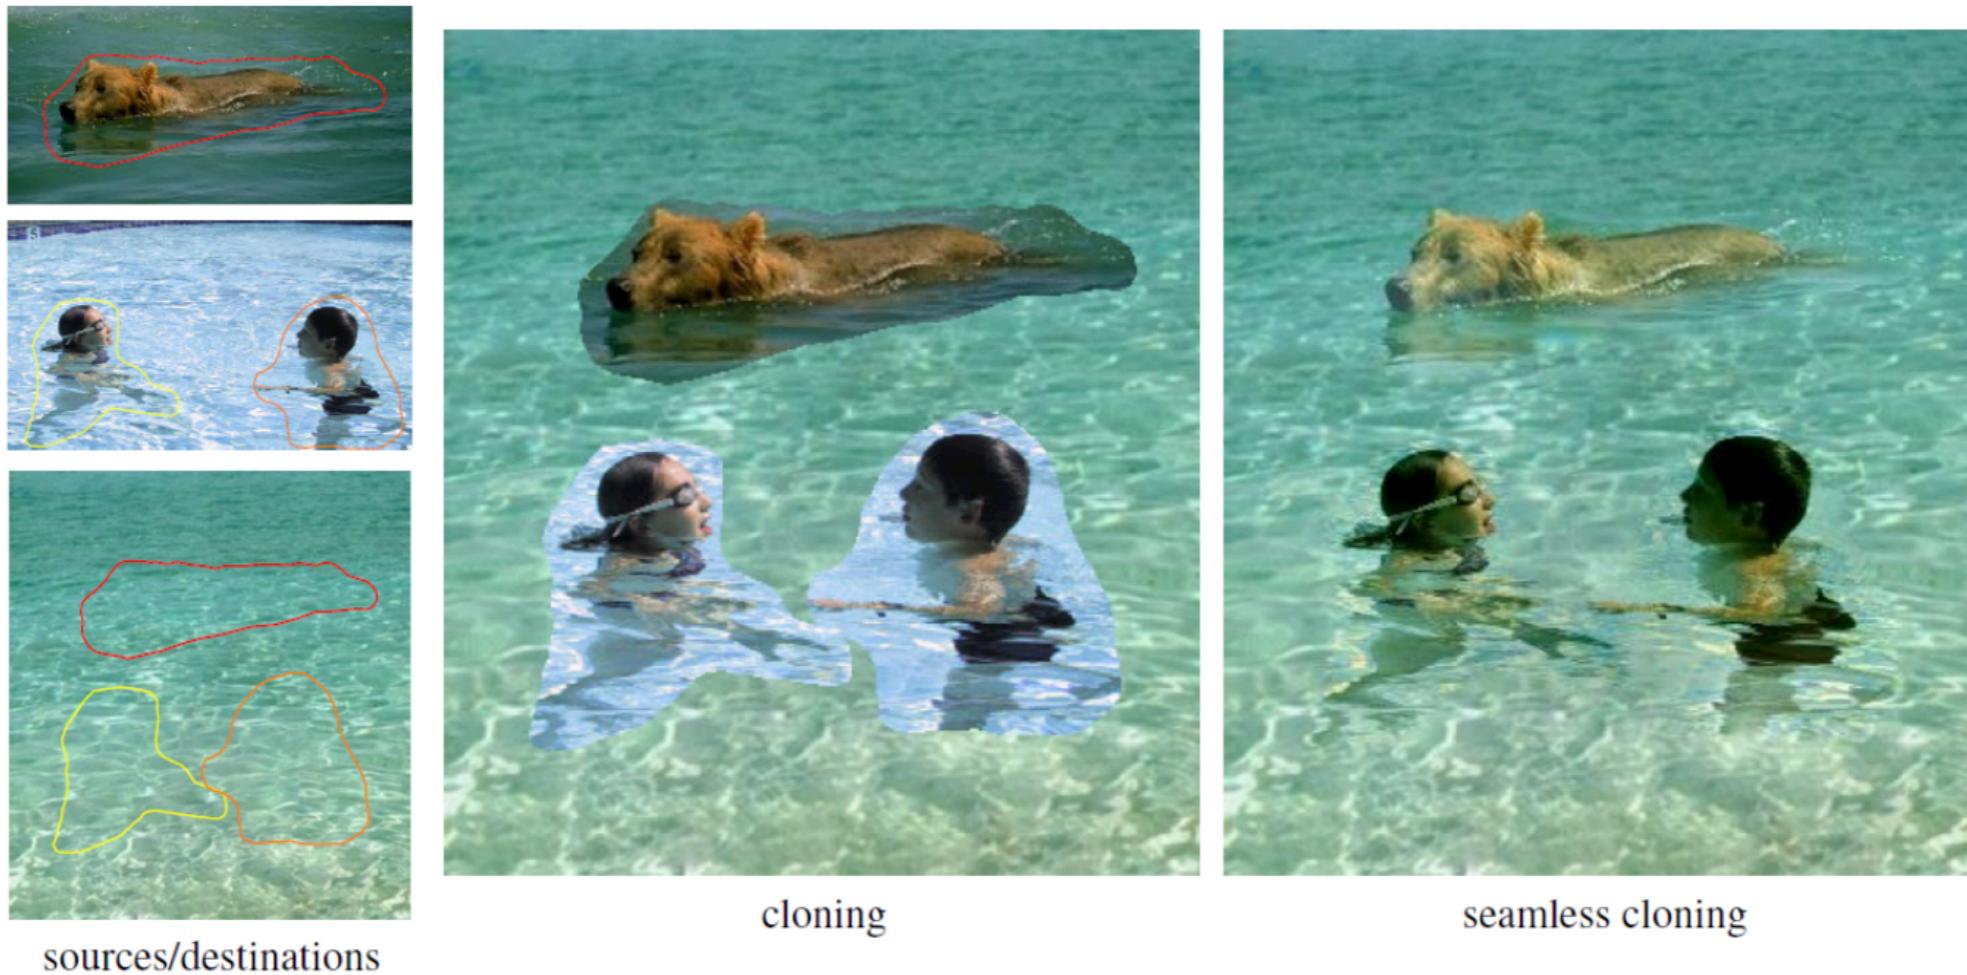
- The gradient vector points in the direction of most rapid increase in intensity
- The gradient direction is given by:  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$
- The *edge strength* is given by the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

# Applications: an example



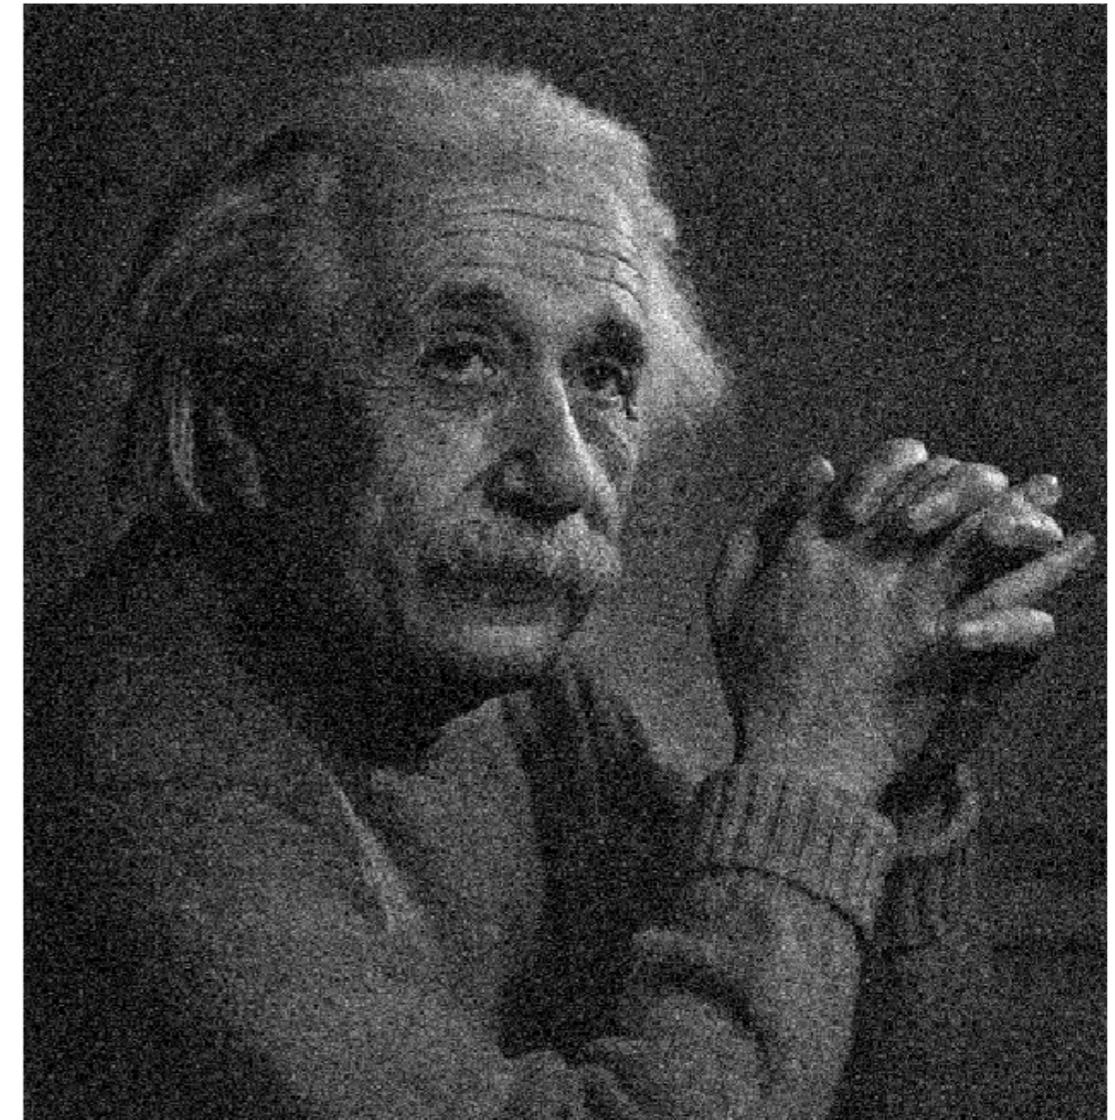
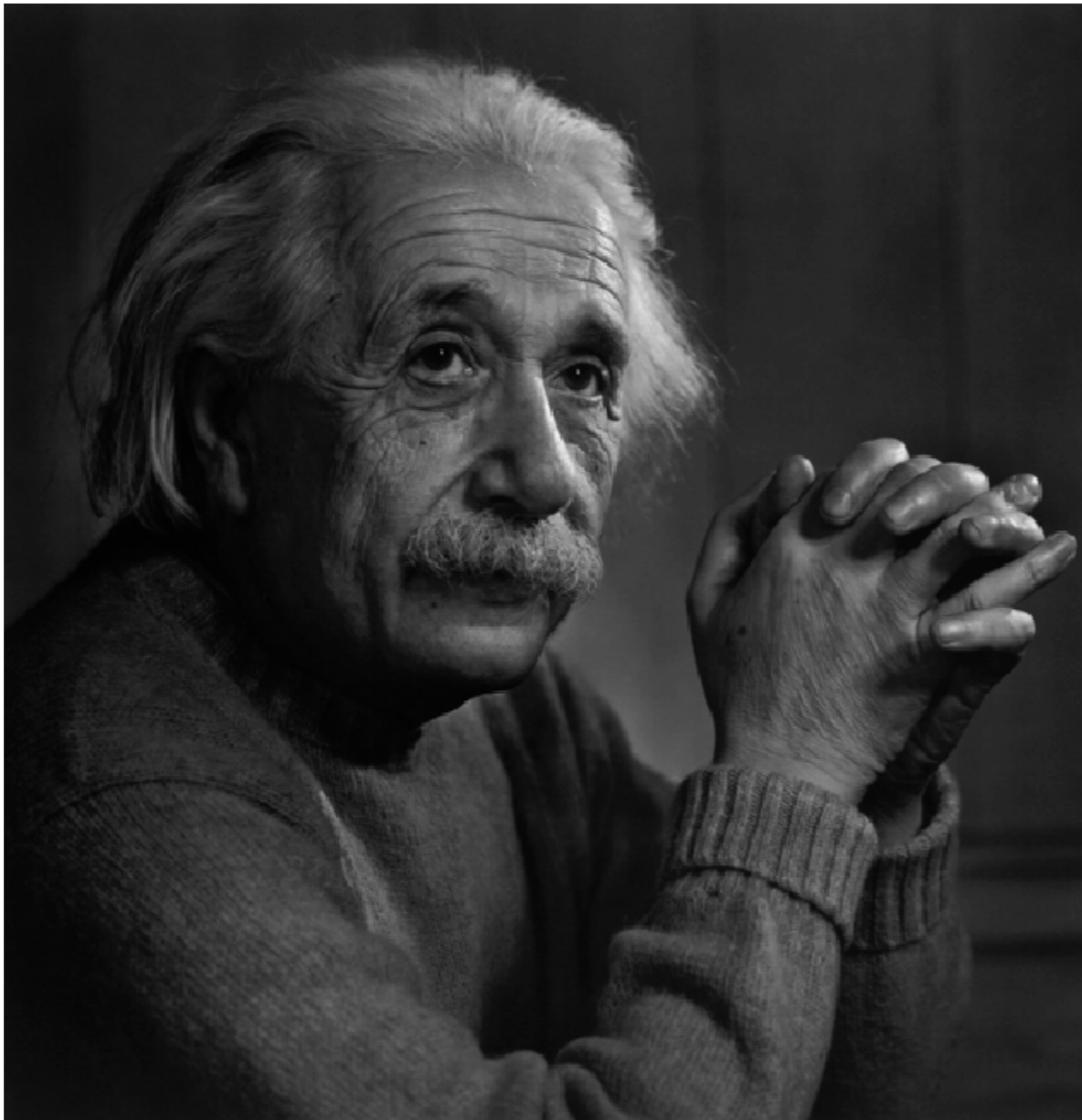
- Gradient-domain image editing
  - ▶ solve for pixel values in the target region to match gradients of the source region while keeping background pixels the same



P. Perez, M. Gangnet, A. Blake, [Poisson Image Editing](#), SIGGRAPH 2003

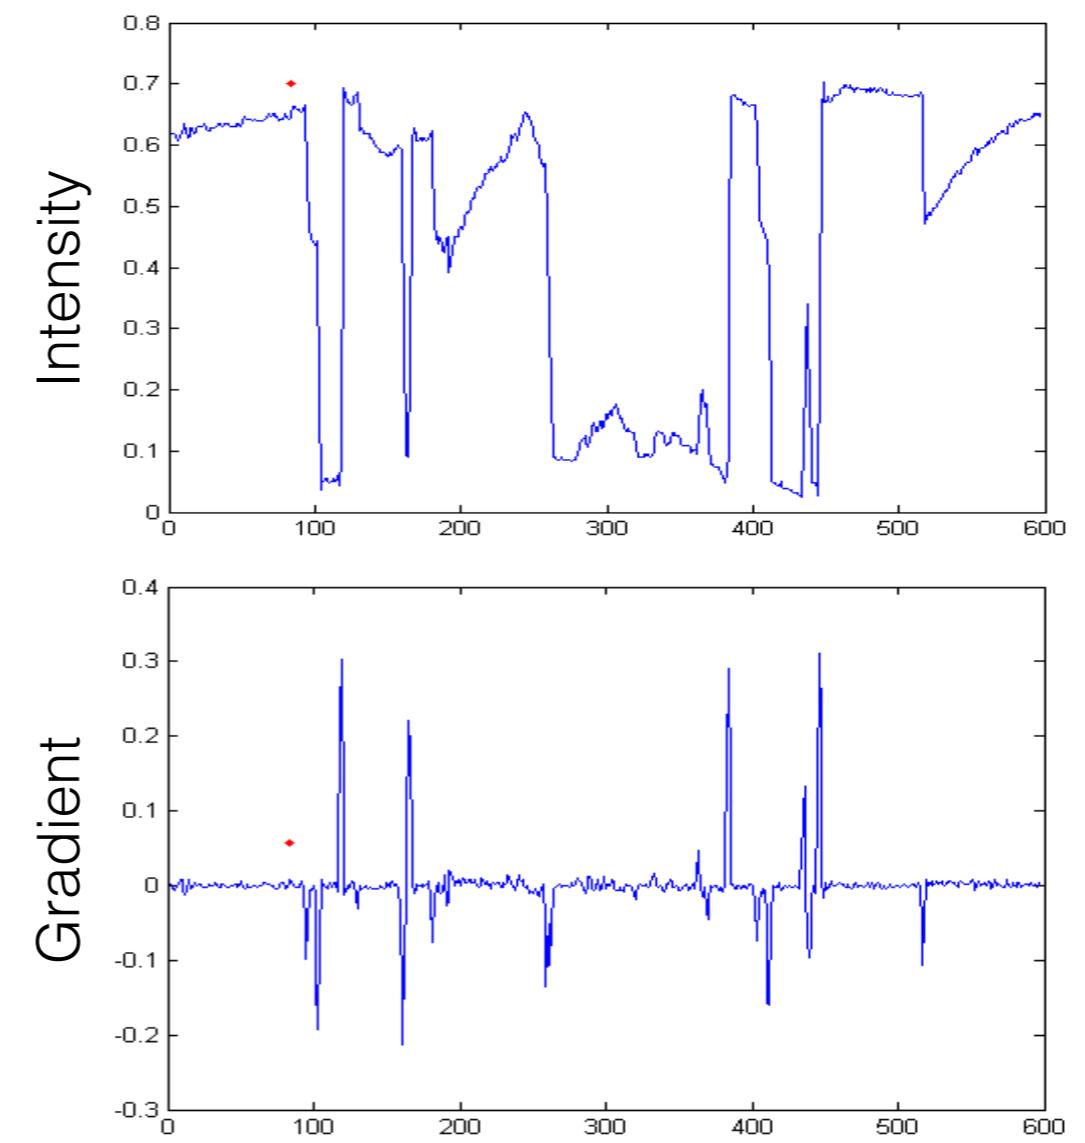


# Effects of noise



# Intensity profile

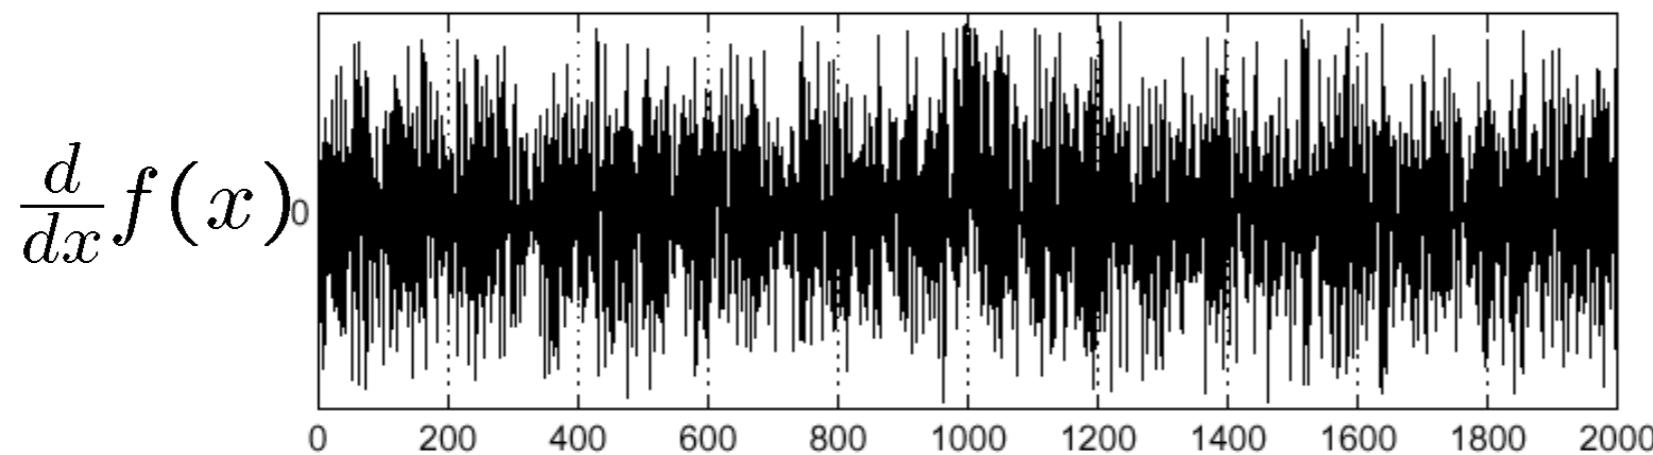
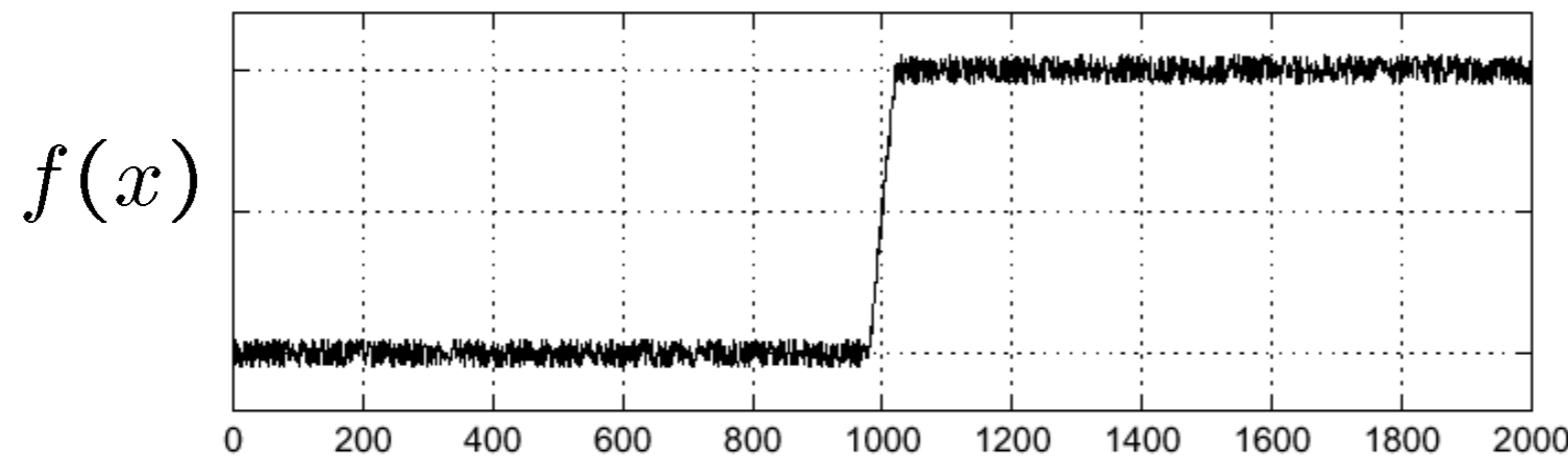
- Consider a single row or column of the image
  - ▶ Plotting intensity as a function of position gives a signal



# Effects of noise



- Consider a single row or column of the image
  - ▶ Plotting intensity as a function of position gives a signal



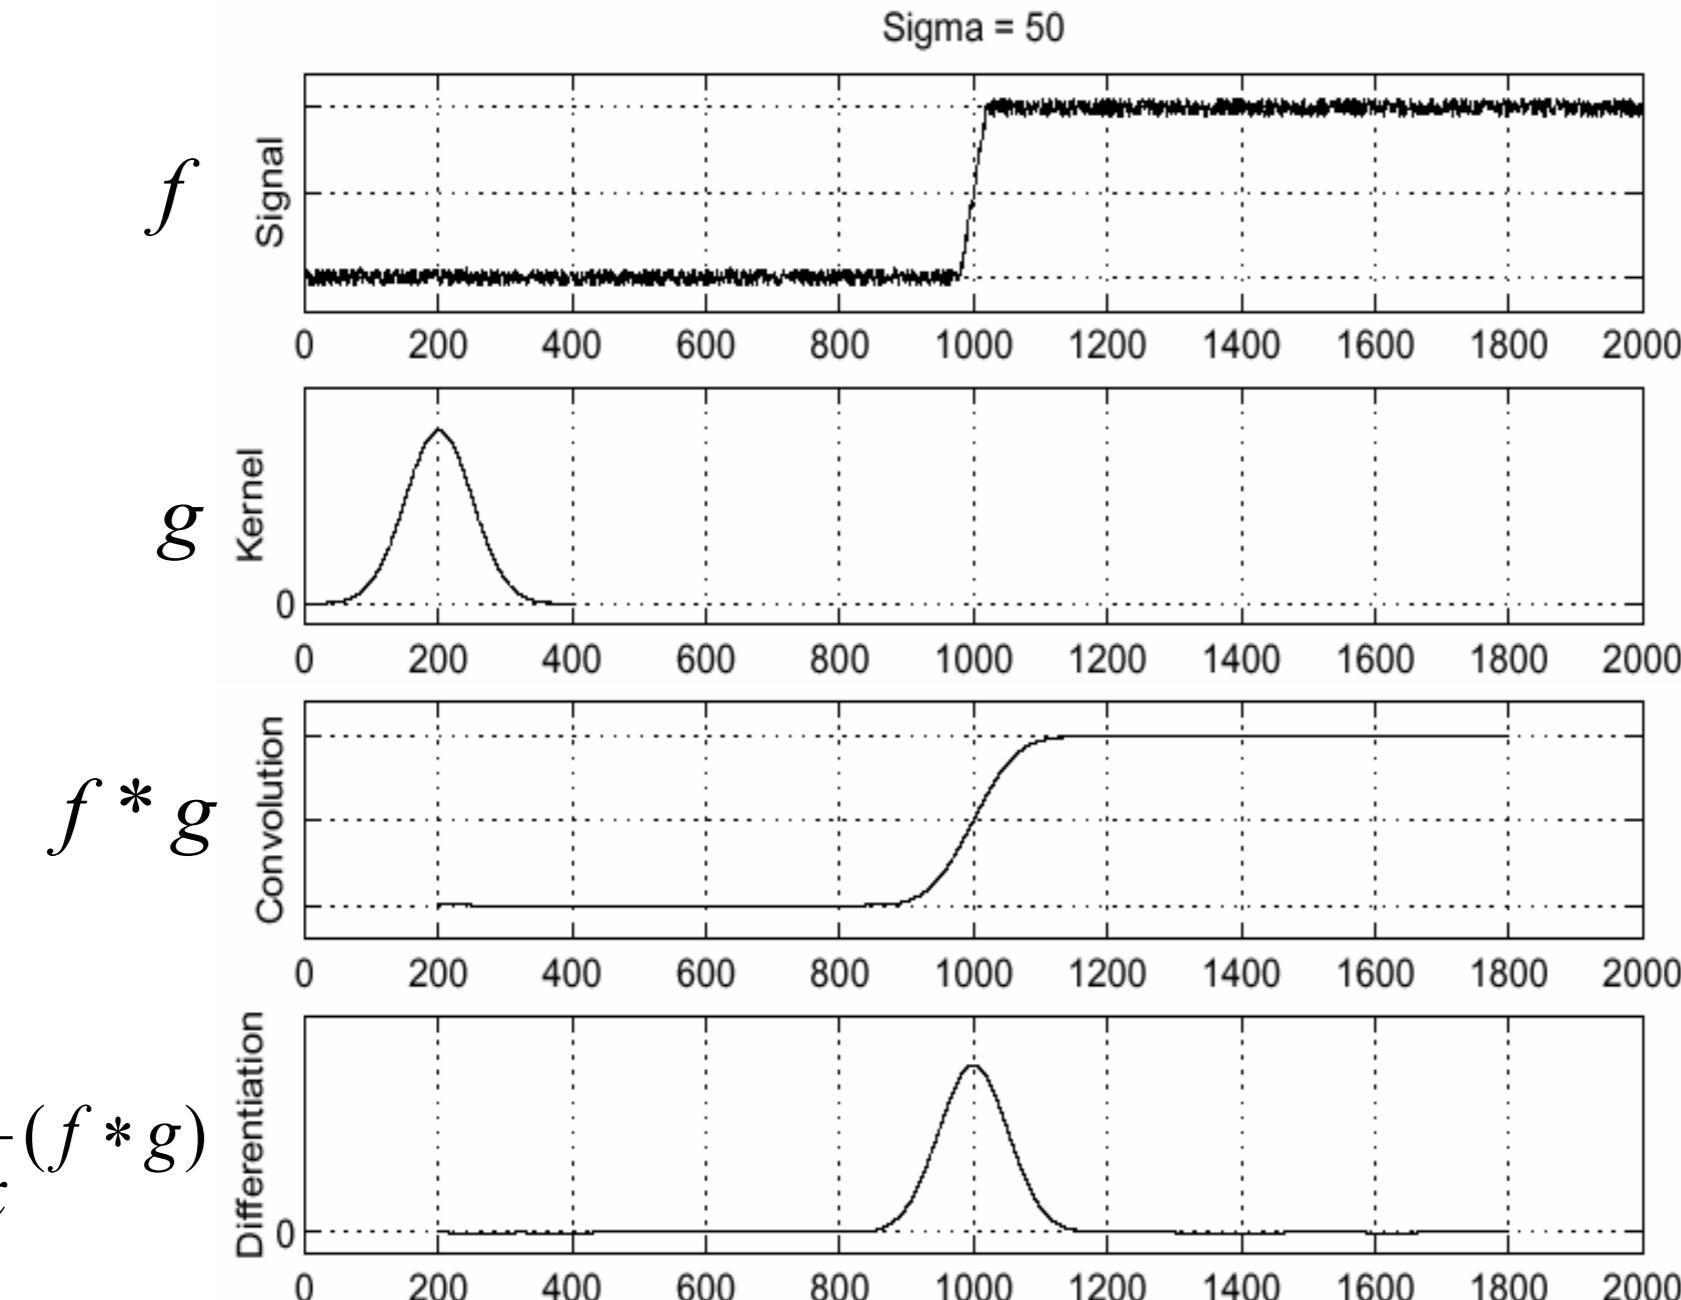
*Where is the  
edge?*

# Effects of noise



- Finite difference filters respond strongly to noise
  - ▶ Image noise results in pixels that look very different from their neighbors
  - ▶ Generally, the larger the noise the stronger the response
- What is to be done?
  - ▶ Smoothing the image should help, by forcing pixels different to their neighbors (=noise pixels?) to look more like neighbors

# Solution: smooth first



Where is the  
edge?

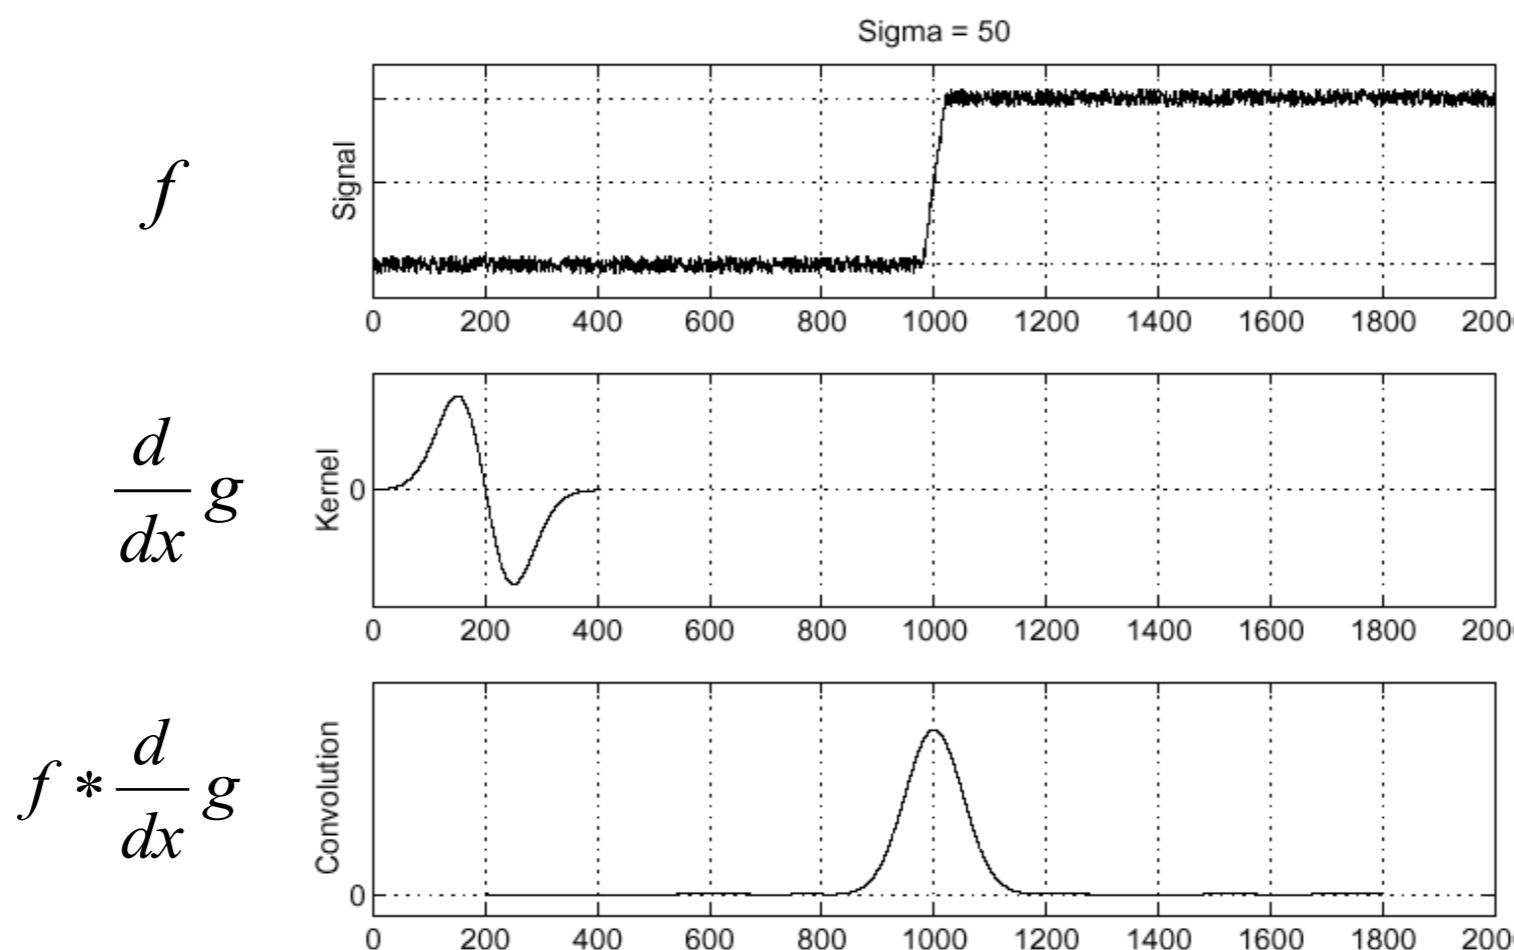
Look for peaks in  $\frac{d}{dx}(f * g)$

# Derivative theorem of convolution

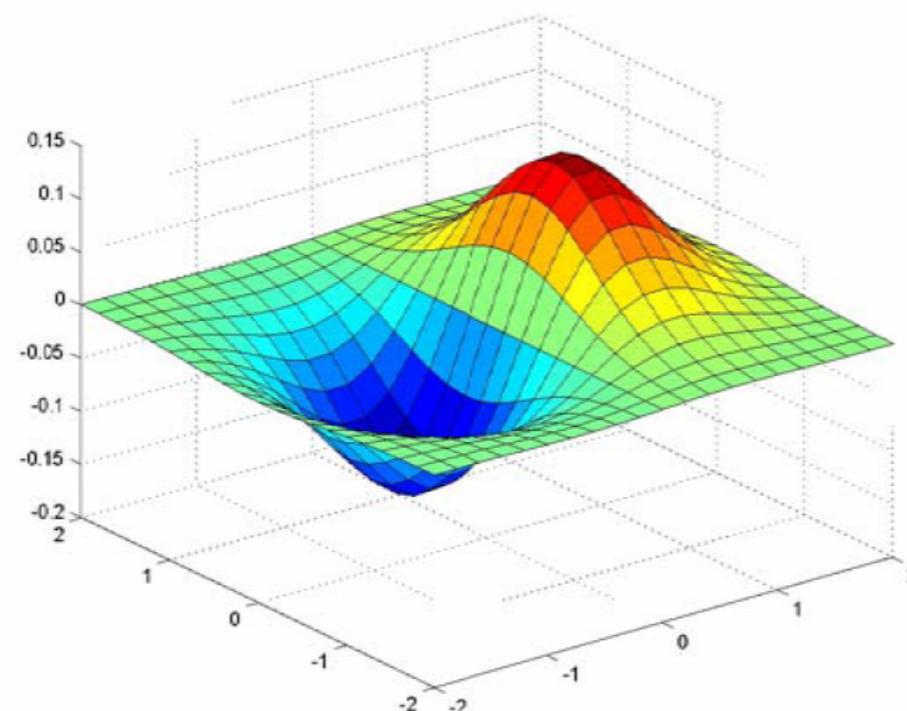
- Differential property of convolution:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

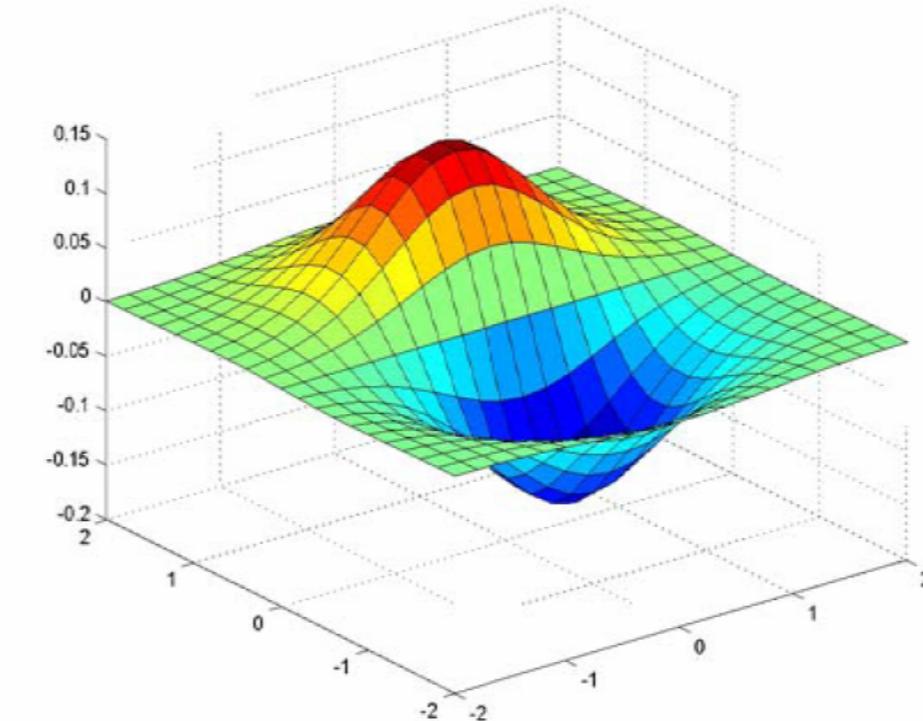
- ▶ It gives us a useful property and save us one operation:



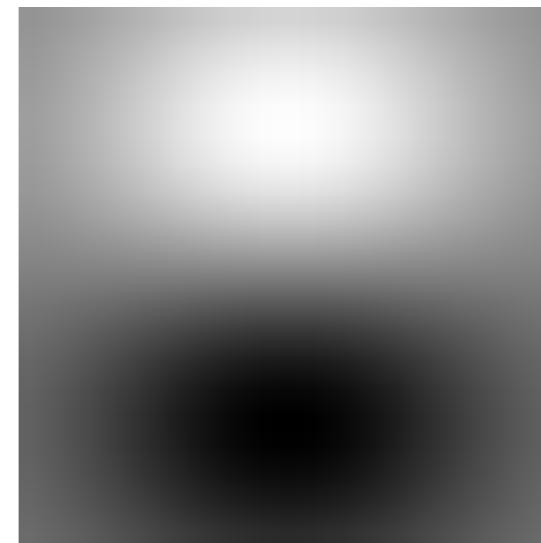
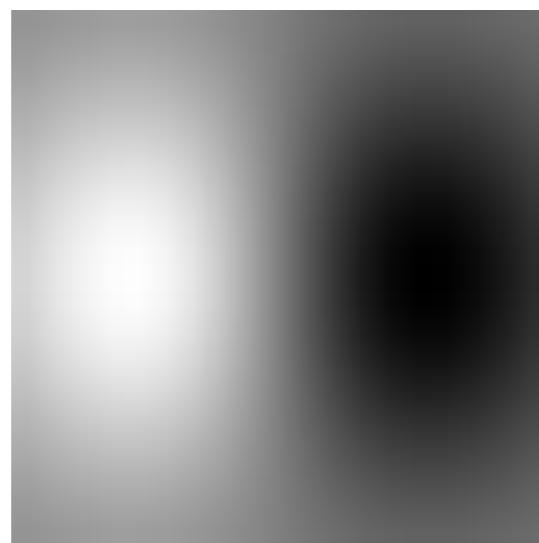
# Derivative of Gaussian filter



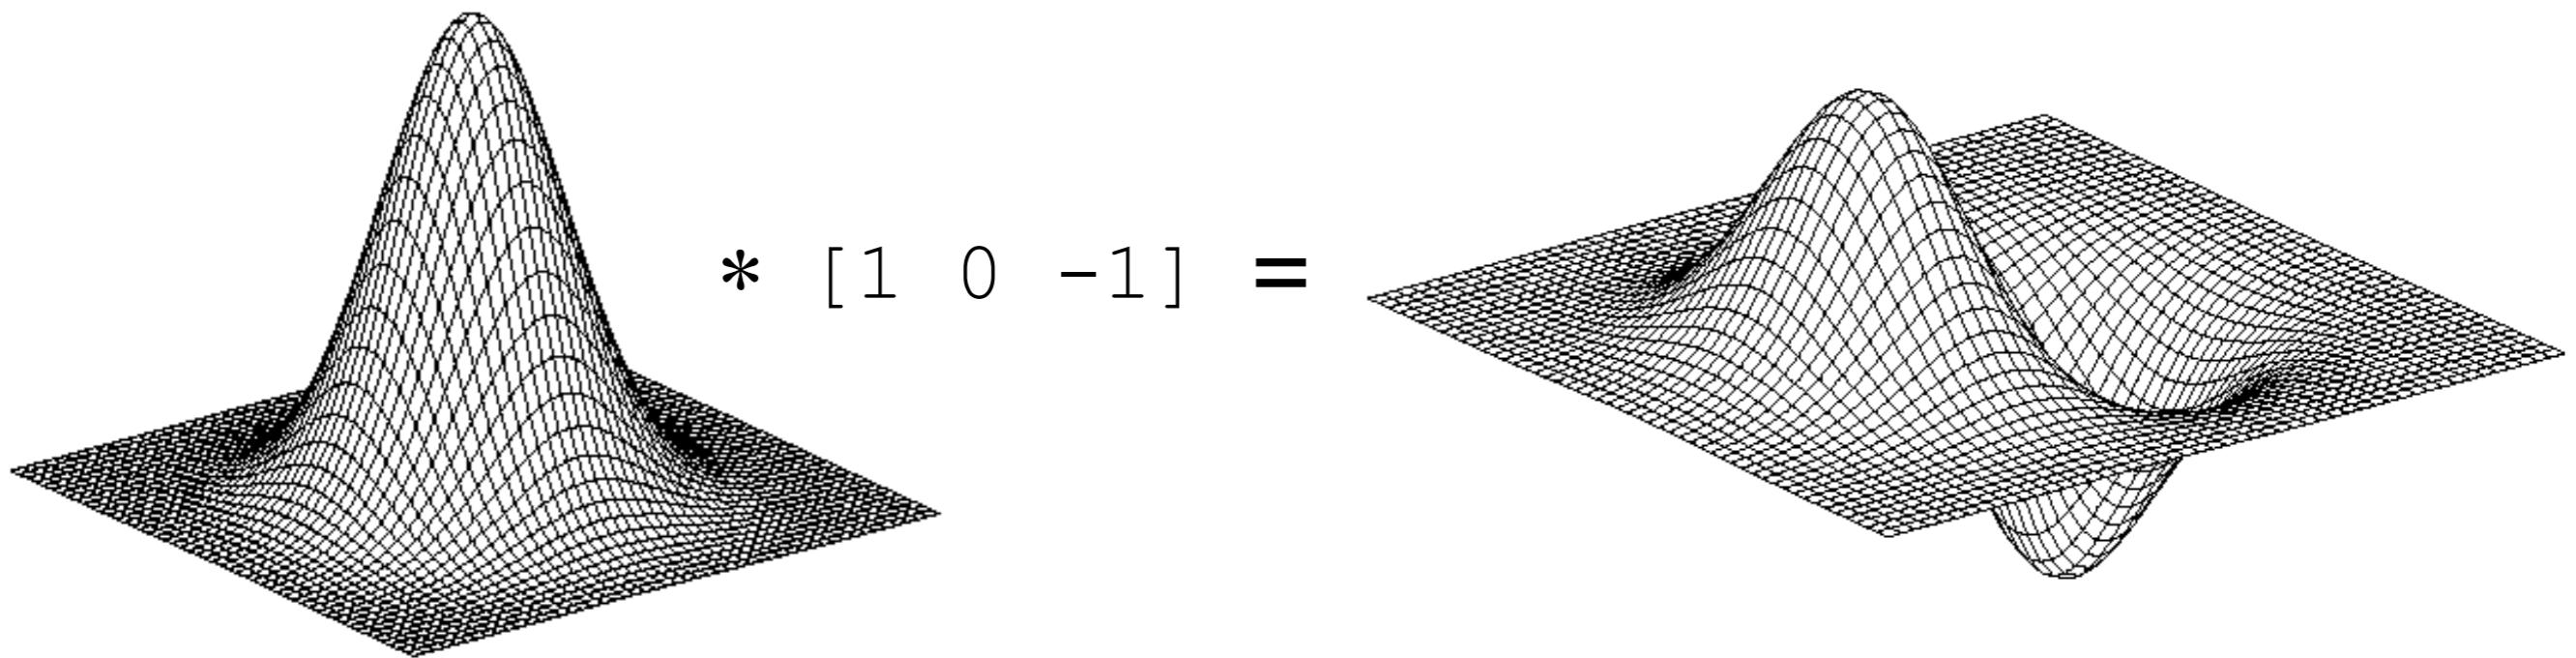
x-direction



y-direction



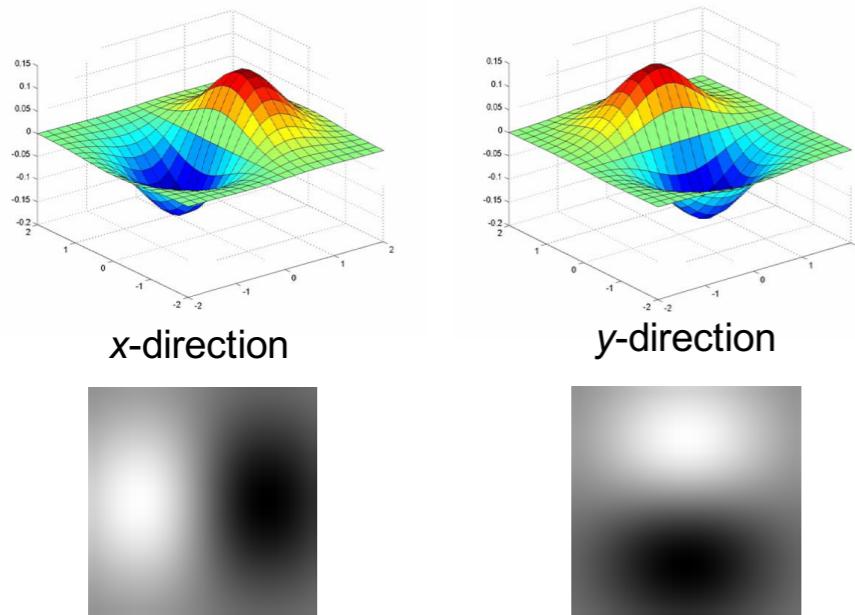
# Derivative of Gaussian filter



2D Gaussian

x-derivative

# Derivative of Gaussian filter

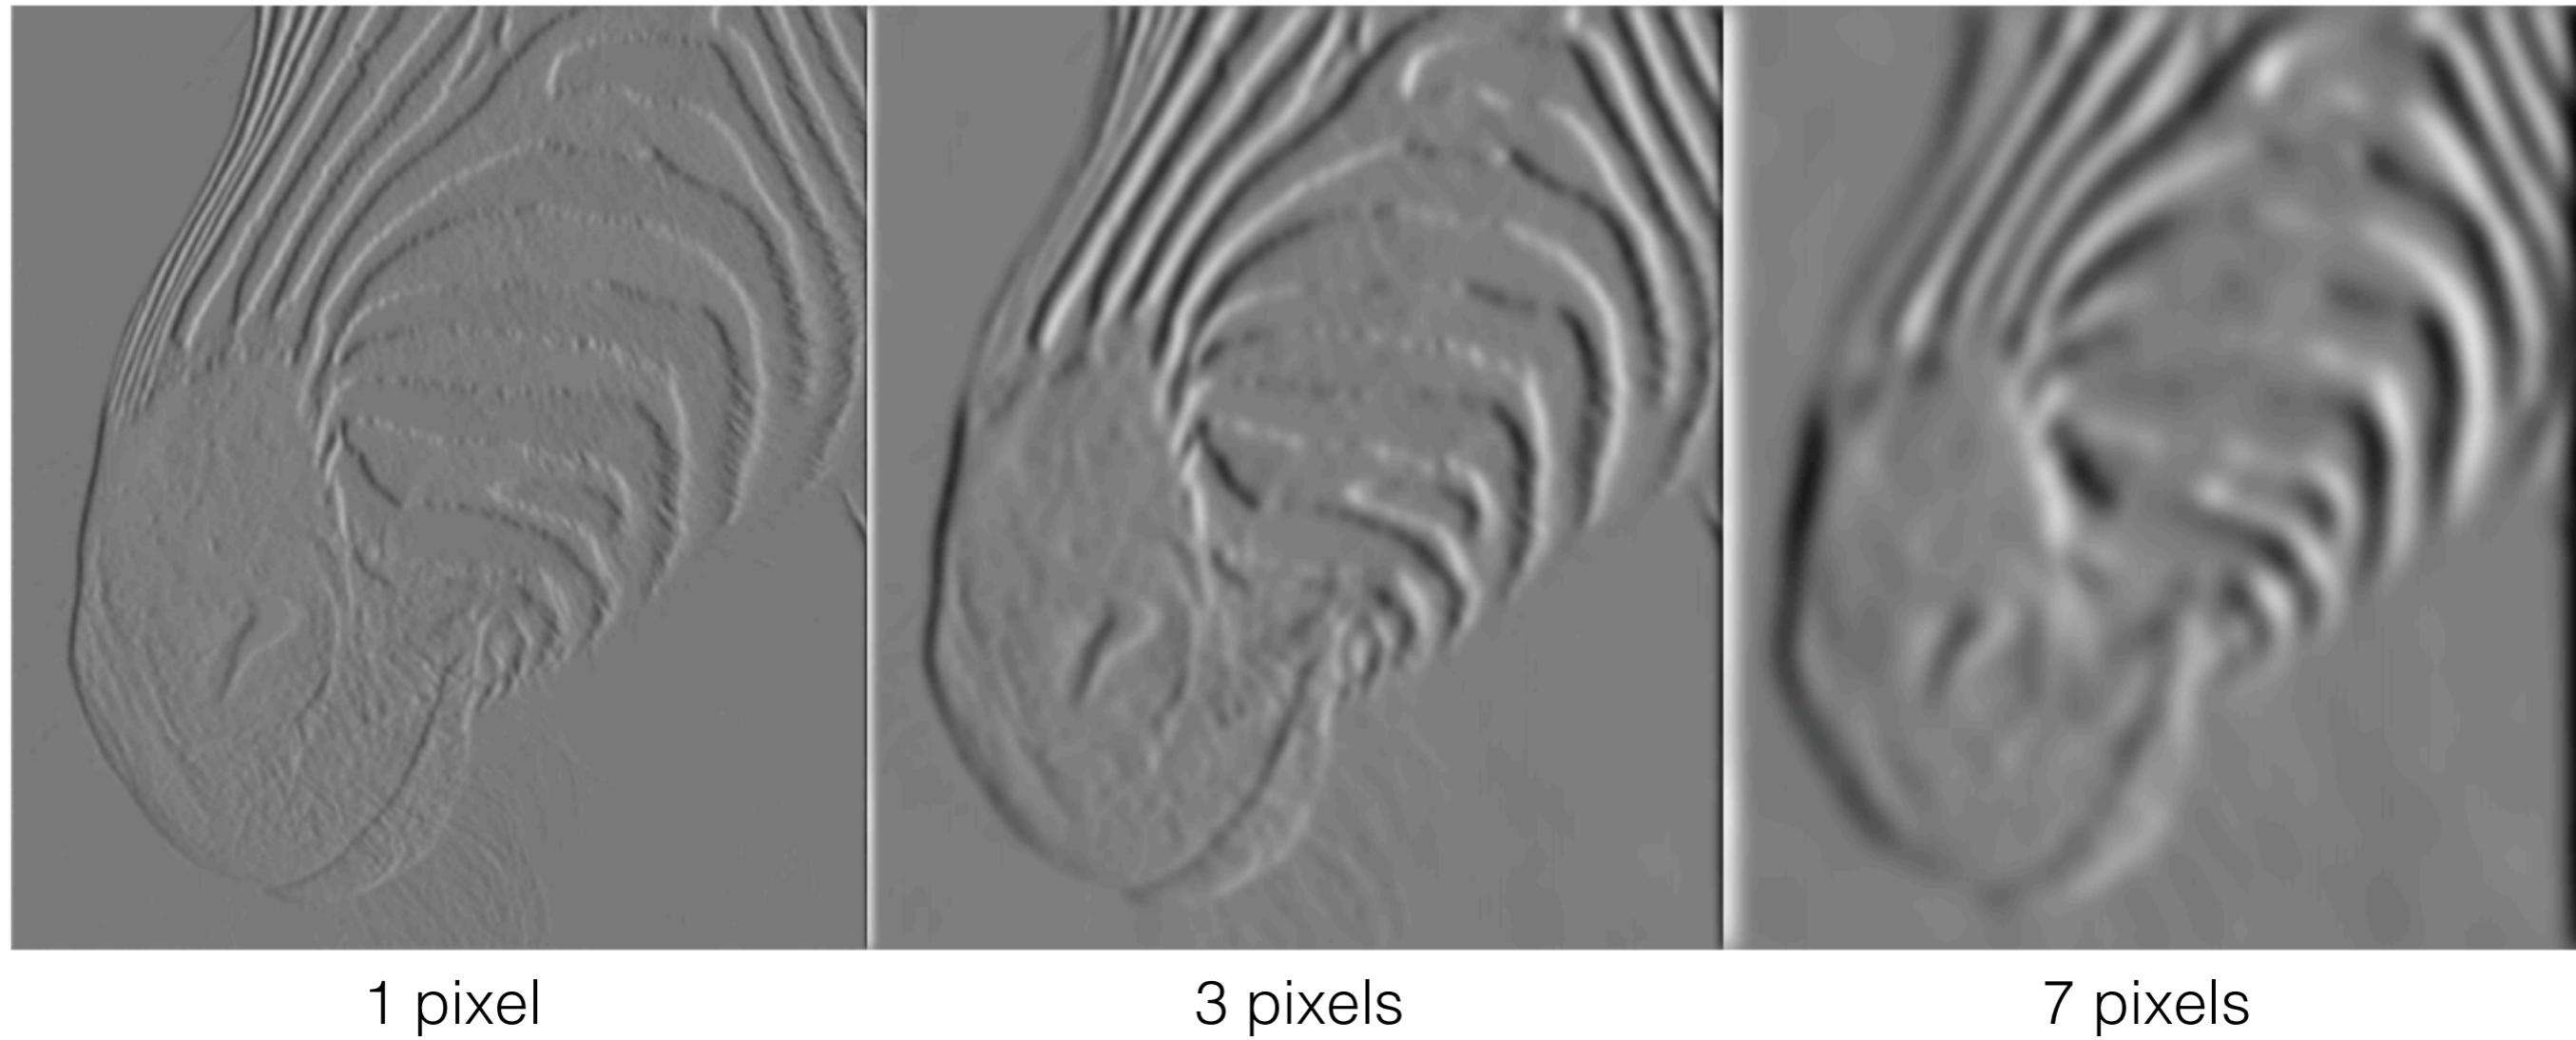


Note: separability of the Gaussian Filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

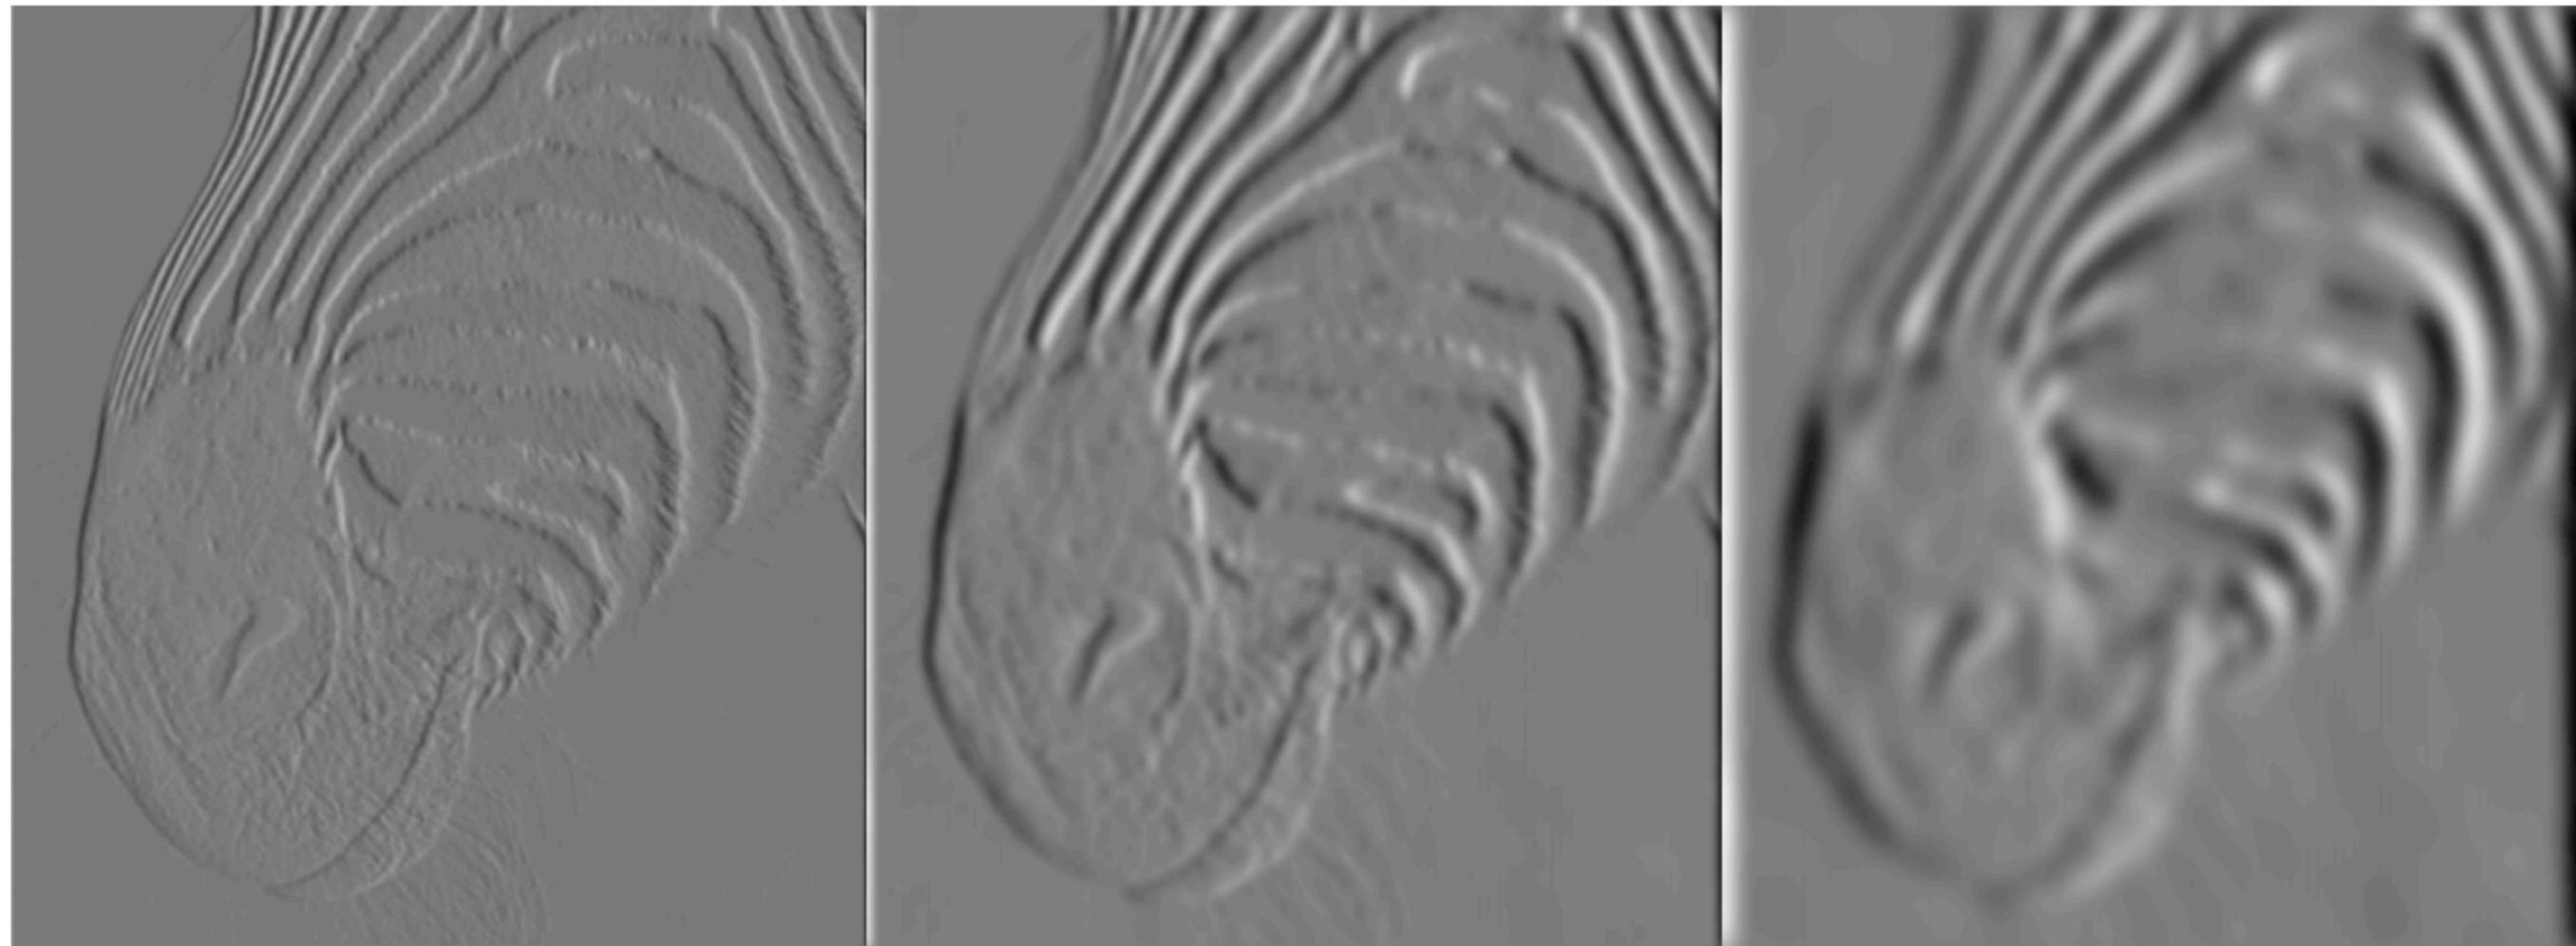
- ▶ The 2D Gaussian can be expressed as the product of 2 functions (one a function of x and the other a function of y)
- ▶ In this case the two functions are the (identical) 1D Gaussian

# Scale of Gaussian derivative filter



- ▶ Note:  $\sigma$  is the *scale / width / spread* of the Gaussian kernel
- ▶ Smoothed derivative remove noise, but blurs edges

# Scale of Gaussian derivative filter



1 pixel

3 pixels

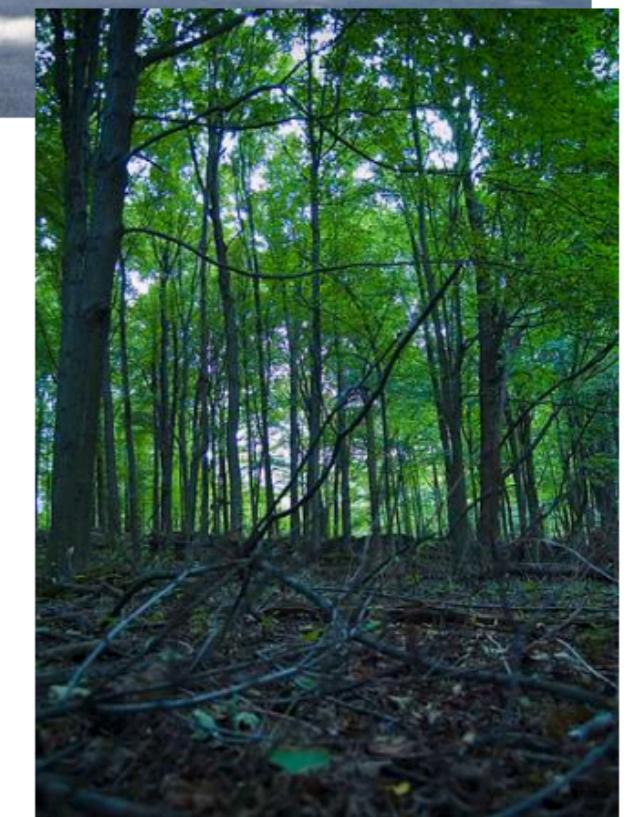
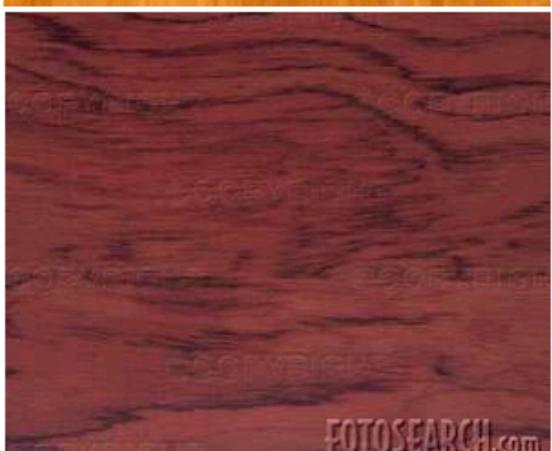
7 pixels

- ▶ The apparent structures differ depending on scale  $\sigma$ 
  - ▶ larger values: larger scale edges detected
  - ▶ smaller values: finer features detected



# So, what scale to choose?

- It depends what we are looking for...

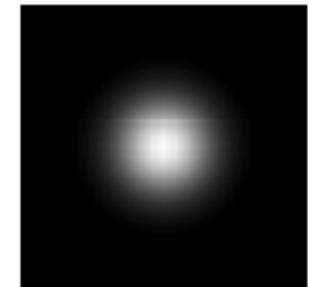


# Review: smoothing vs derivative filters



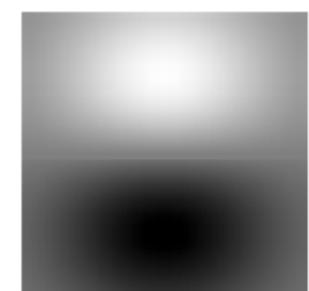
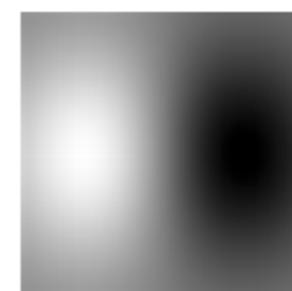
- Smoothing filters

- Gaussian: remove high-frequency components
- Can the values of a smoothing filter be negative?
- What should the values sum to?
  - **One**: constant regions are not affected by the filter



- Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
  - **Zero**: no response in constant regions
- High absolute value at points of high contrast



# Building an edge detector



Original (input) image



Output

# Edge detection: Sobel filter



- A derivative filter + some smoothing



*Input image*



1	0	-1
2	0	-2
1	0	-1

*Left Sobel*

*Filter returns large response on vertical or horizontal lines?*

A: Responds to vertical lines

*Is the output always positive?*

A: Output can be positive or negative

# Edge detection: Sobel filter

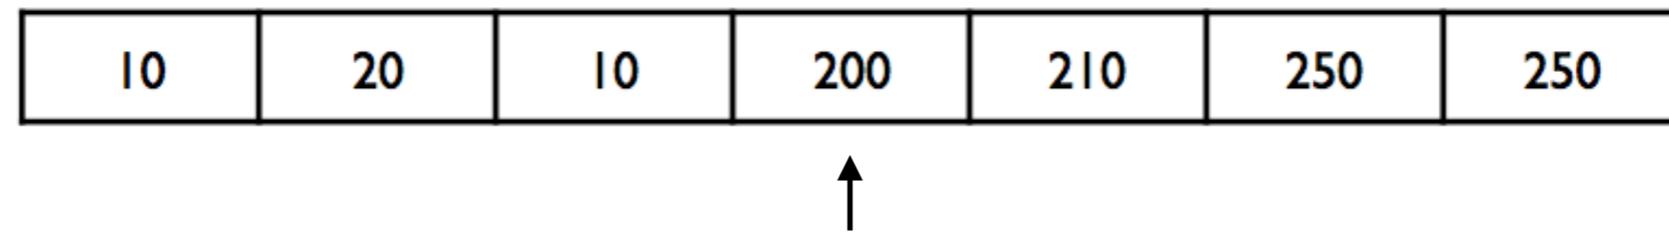


- Where does this filter come?
  - ▶ Remember: (central) discrete derivative in 1D

$$\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$$

$[1 \ 0 \ -1]$   
*1D derivative filter*

- ▶ Example:



# Edge detection: Sobel filter

- Decomposing the Sobel Filter

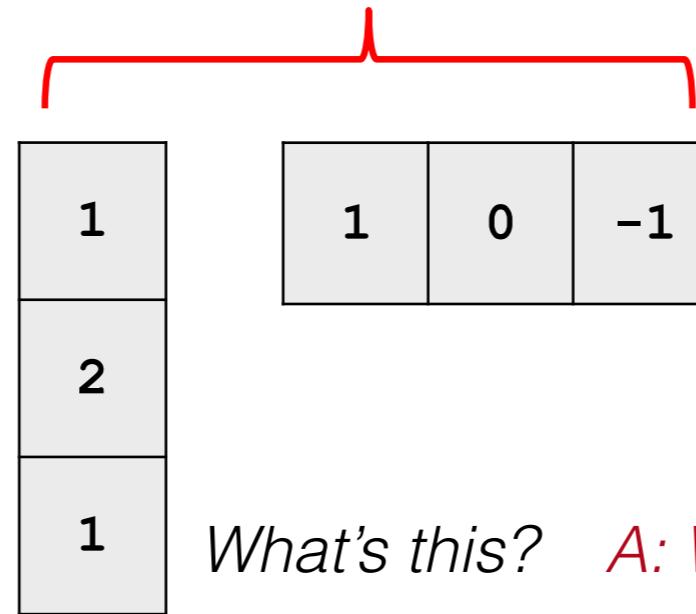


\*

1	0	-1
2	0	-2
1	0	-1

*Left Sobel*

*Input image*



# Edge detection: derivative filters

- Common derivative filters:

**Sobel**

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

**Scharr**

3	0	-3
10	0	-10
3	0	-3

3	10	3
0	0	0
-3	-10	-3

**Prewitt**

1	0	-1
1	0	-1
1	0	-1

1	1	1
0	0	0
-1	-1	-1

**Roberts**

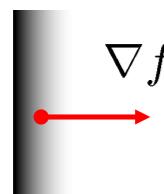
0	1
-1	0

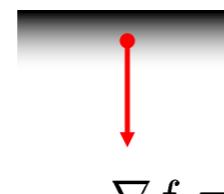
1	0
0	-1

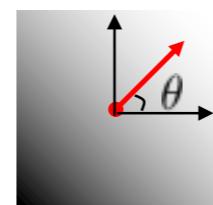
# Building an edge detector

- The Canny Edge detector
  - ▶ Filter image with derivative of Gaussian
  - ▶ Find magnitude and orientation of gradient
  - ▶ ...

## ***Image gradient***

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$


$$\theta = \tan^{-1} \left( \frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

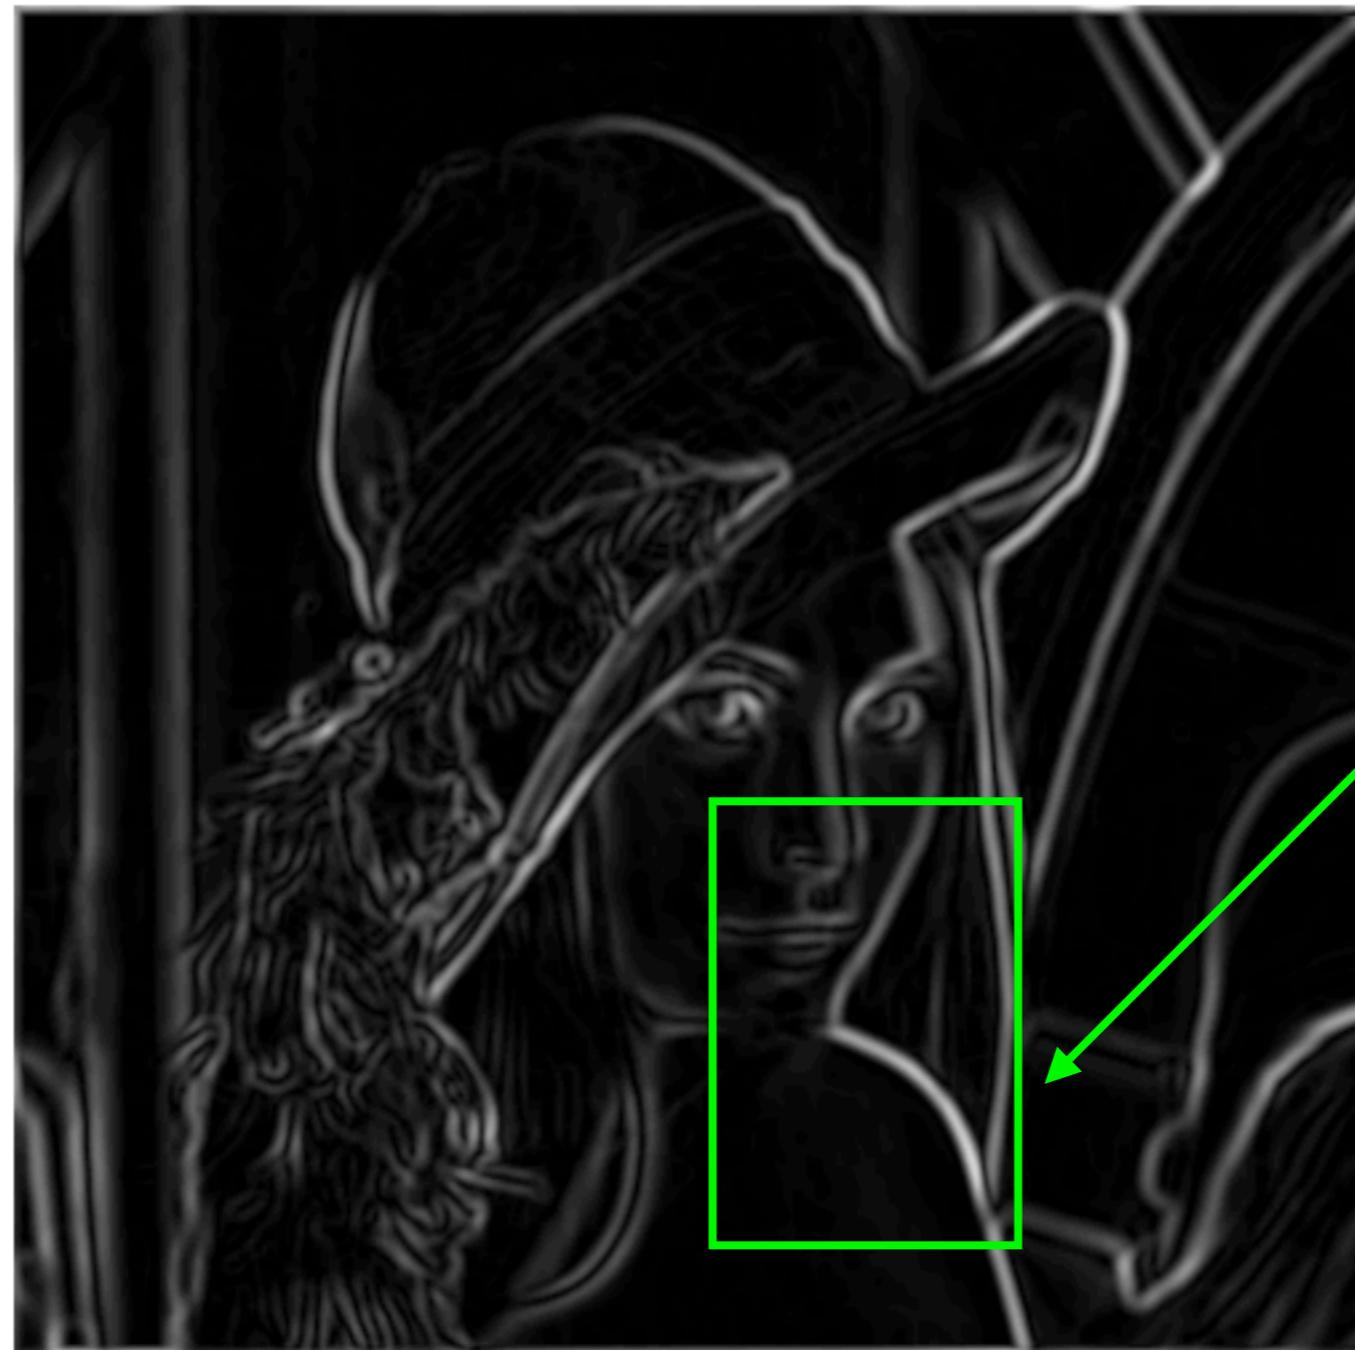
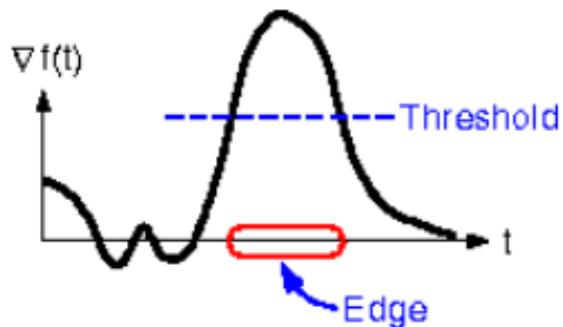
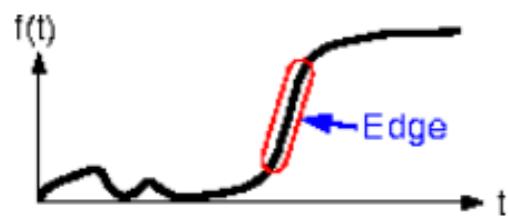
# Building an edge detector: Canny



*Gradient magnitude*

# Building an edge detector: Canny

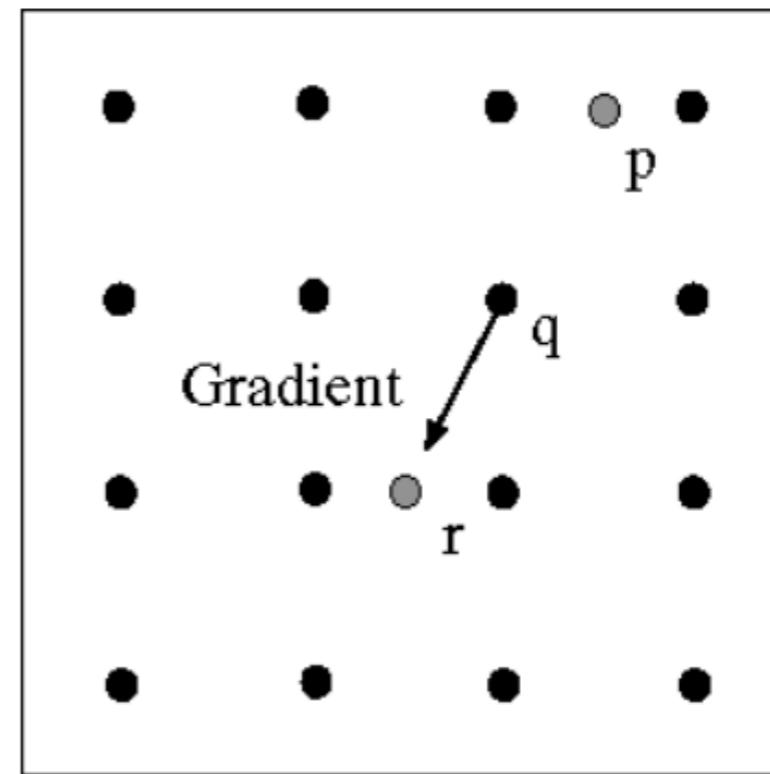
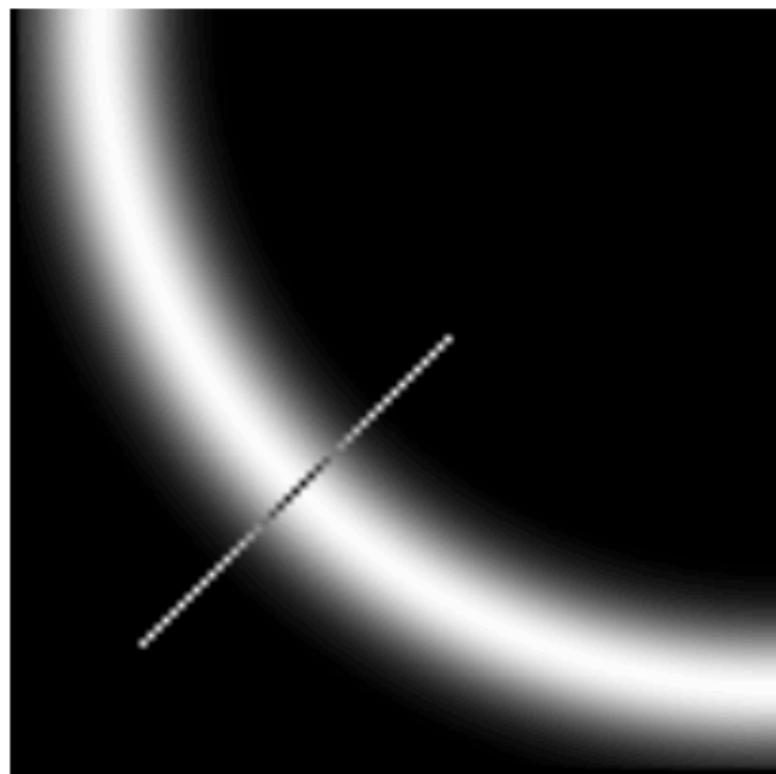
## Thresholding



*Gradient magnitude*

# Building an edge detector: Canny

- Non-maximum suppression
  - ▶ Check if pixel is local max along gradient direction
  - ▶ Select single maximum across width of the edge  
(requires checking interpolated pixels p and r)



# Building an edge detector: Canny



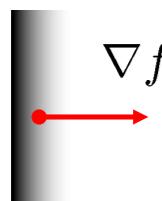
*Problem: pixels  
along this edge  
didn't survive the  
thresholding*

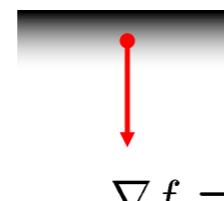
*Non-maximum suppression: thinning*

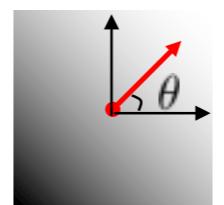
# Building an edge detector: Canny

- The Canny Edge detector
  - ▶ Filter image with derivative of Gaussian
  - ▶ Find magnitude and orientation of gradient
  - ▶ Non-maximum suppression
    - ▶ Thin wide “ridges” down to single pixel width

## ***Image gradient***

$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


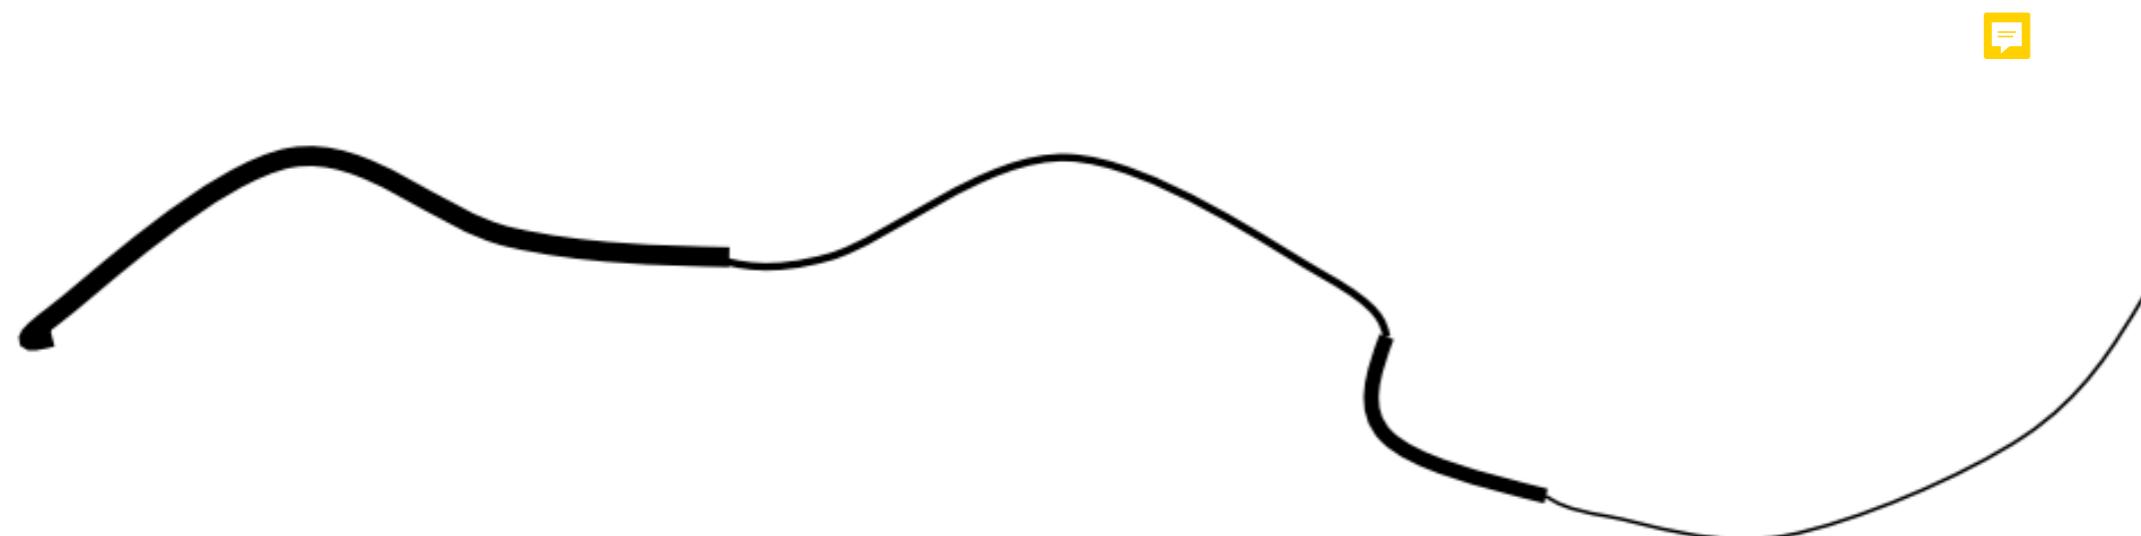
$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$


$$\theta = \tan^{-1} \left( \frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

# Building an edge detector: Canny

- Hysteresis thresholding: use a high threshold to start edge curves, and a low threshold to continue them



# Building an edge detector: Canny

- Hysteresis thresholding



*Original image*



*high threshold (strong edges)*



*low threshold (weak edges)*



*Hysteresis threshold*

# Building an edge detector: Canny

- The Canny Edge detector
  - ▶ Filter image with derivative of Gaussian
  - ▶ Find magnitude and orientation of gradient
  - ▶ Non-maximum suppression
    - ▶ Thin wide “ridges” down to single pixel width
  - ▶ Linking and thresholding (hysteresis)
    - ▶ Two thresholds: use the high threshold to start edge curves and the low threshold to continue them

Matlab: `edge(image, 'canny');`

J. Canny, “A Computational Approach To Edge Detection”, TPAMI 1986

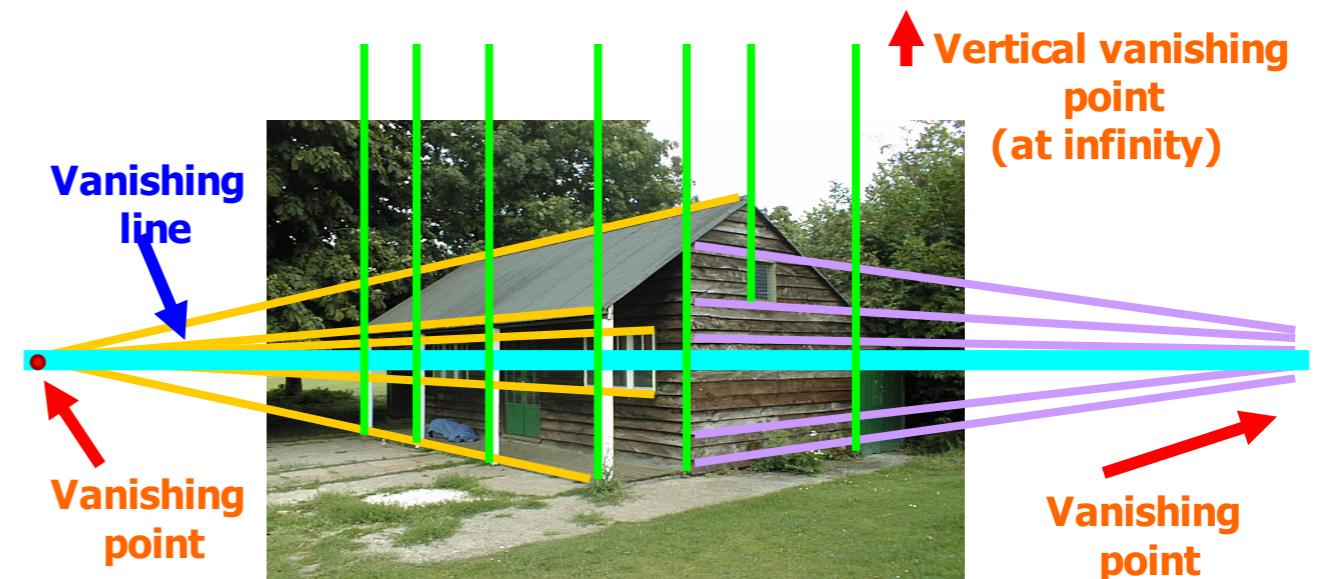
# Recall: why do we care about edges?

- Extract information, recognize objects

Edges are great, but they are not very robust to a number of transformations...

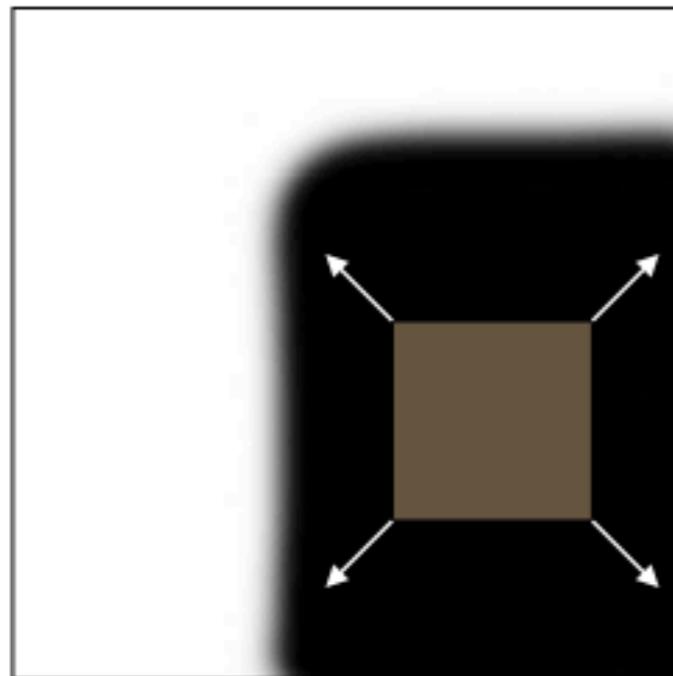


- Recover geometry and viewpoint

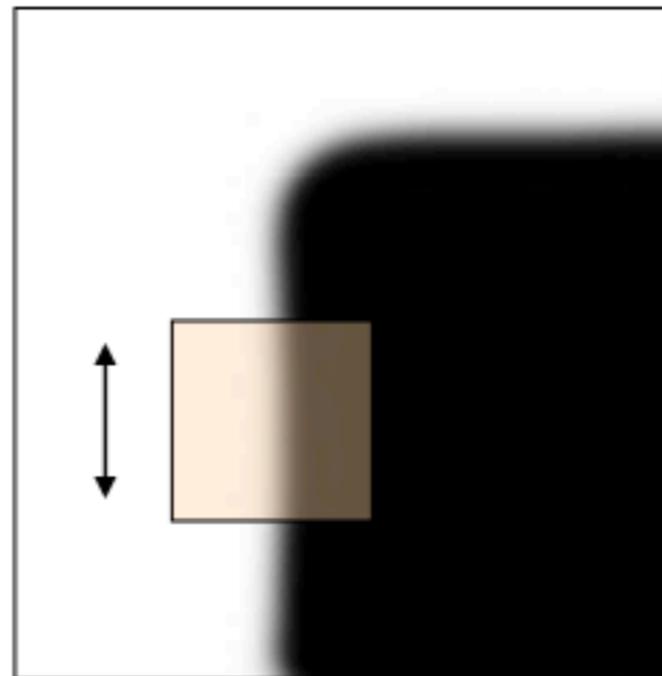


# Edges, corners and blobs (keypoints)

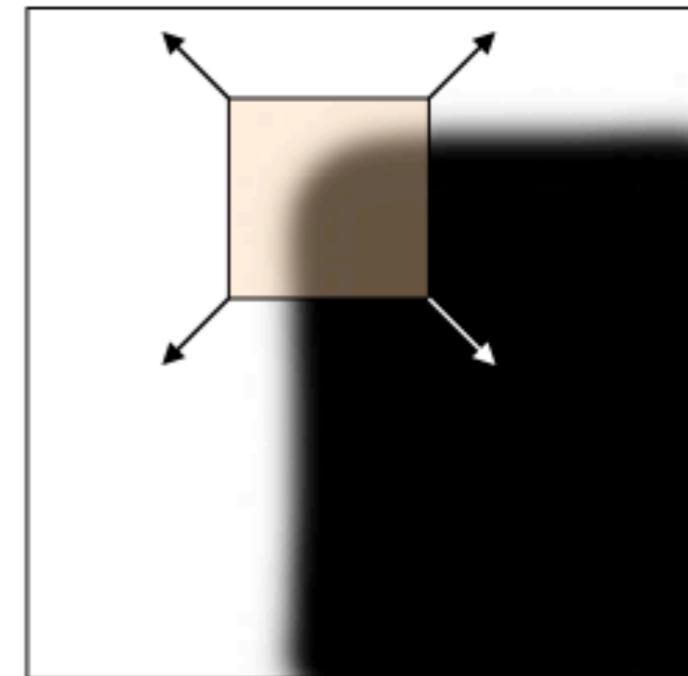
- Low-level image features: a basic classification



**“flat” region:**  
no change in all  
directions



**“edge”:**  
no change along  
the edge direction

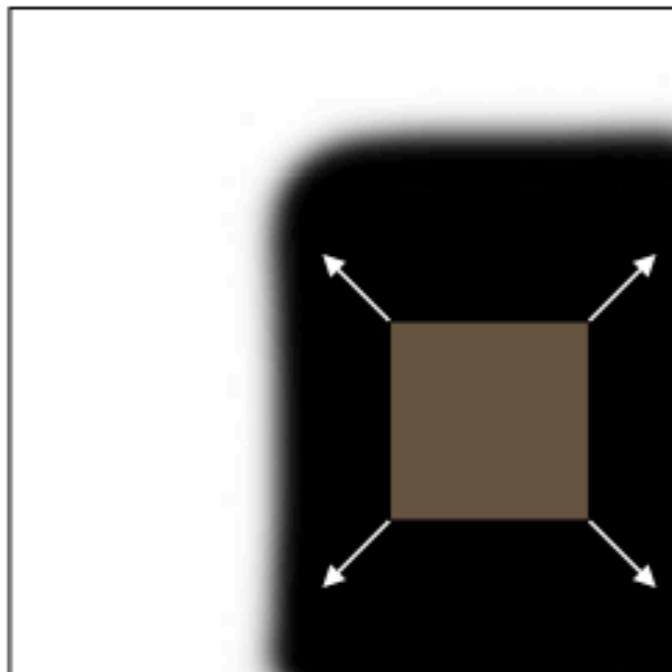


**“corner”:**  
significant change  
in all directions

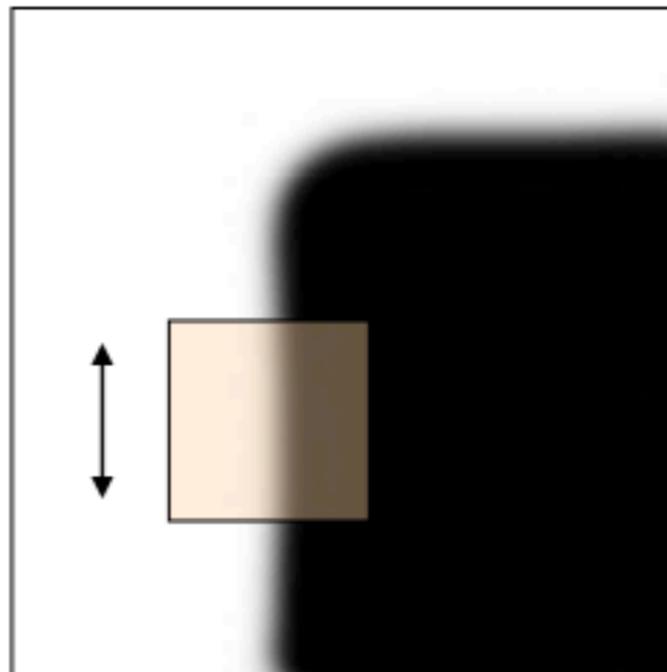
# Corner detection: basic idea



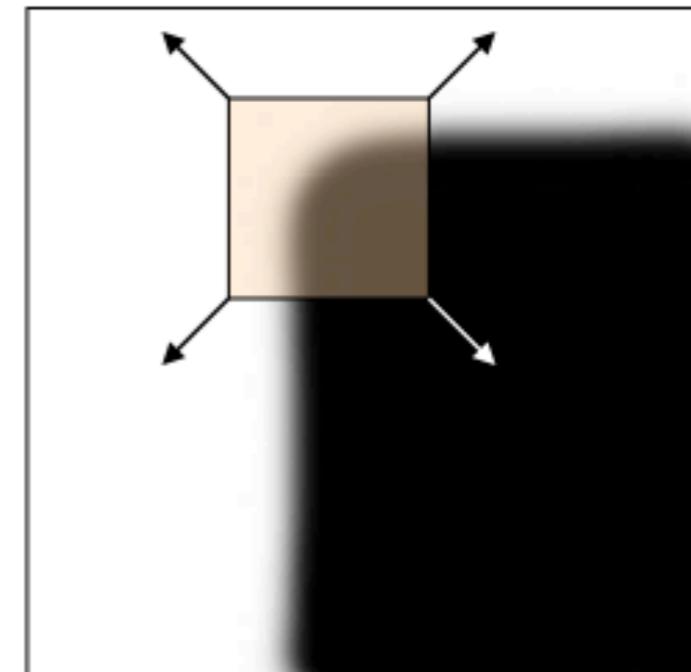
- We should easily recognize the point by looking through a small window (locally)
  - Shifting a window in any direction should give a large change in intensity



**“flat” region:**  
**no change in all directions**



**“edge”:**  
**no change along the edge direction**



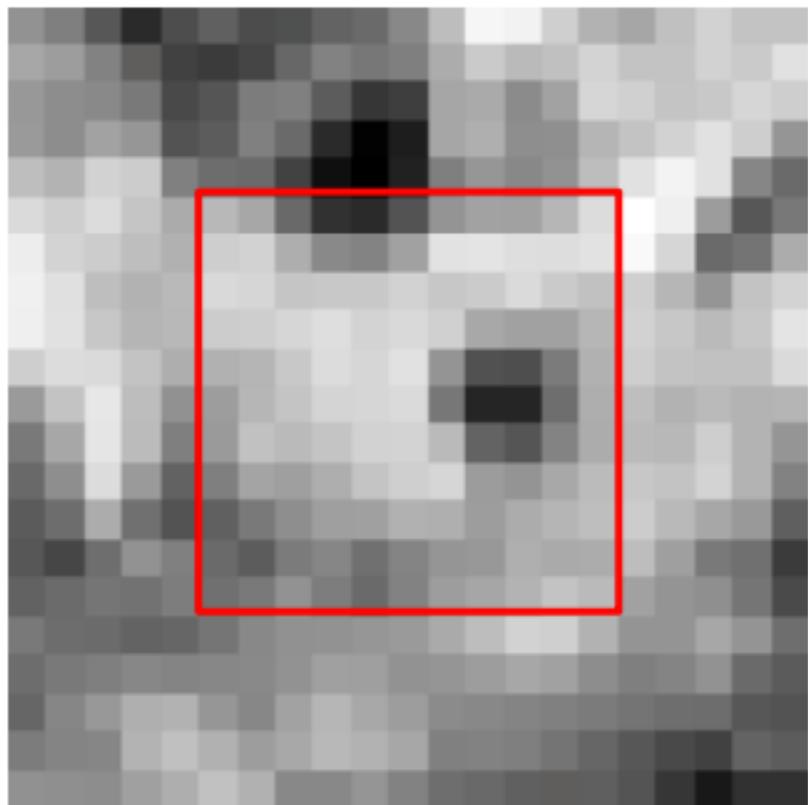
**“corner”:**  
**significant change in all directions**

# Corner detection: basic idea

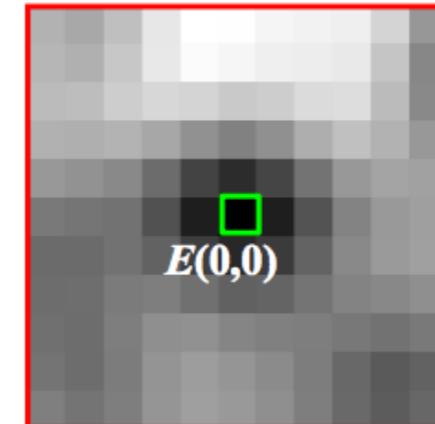
- Given an image  $I$ , Captures change in appearance of window  $W$  for the shift  $[u, v]$ :

$$E(u, v) = \sum_{(x,y) \in W} [I(x+u, y+v) - I(x, y)]^2$$

$I(x, y)$



$E(u, v)$



# Corners and keypoints



- Corners (e.g. Harris corner detector) are a good example of keypoints / interest points



Harris corners example

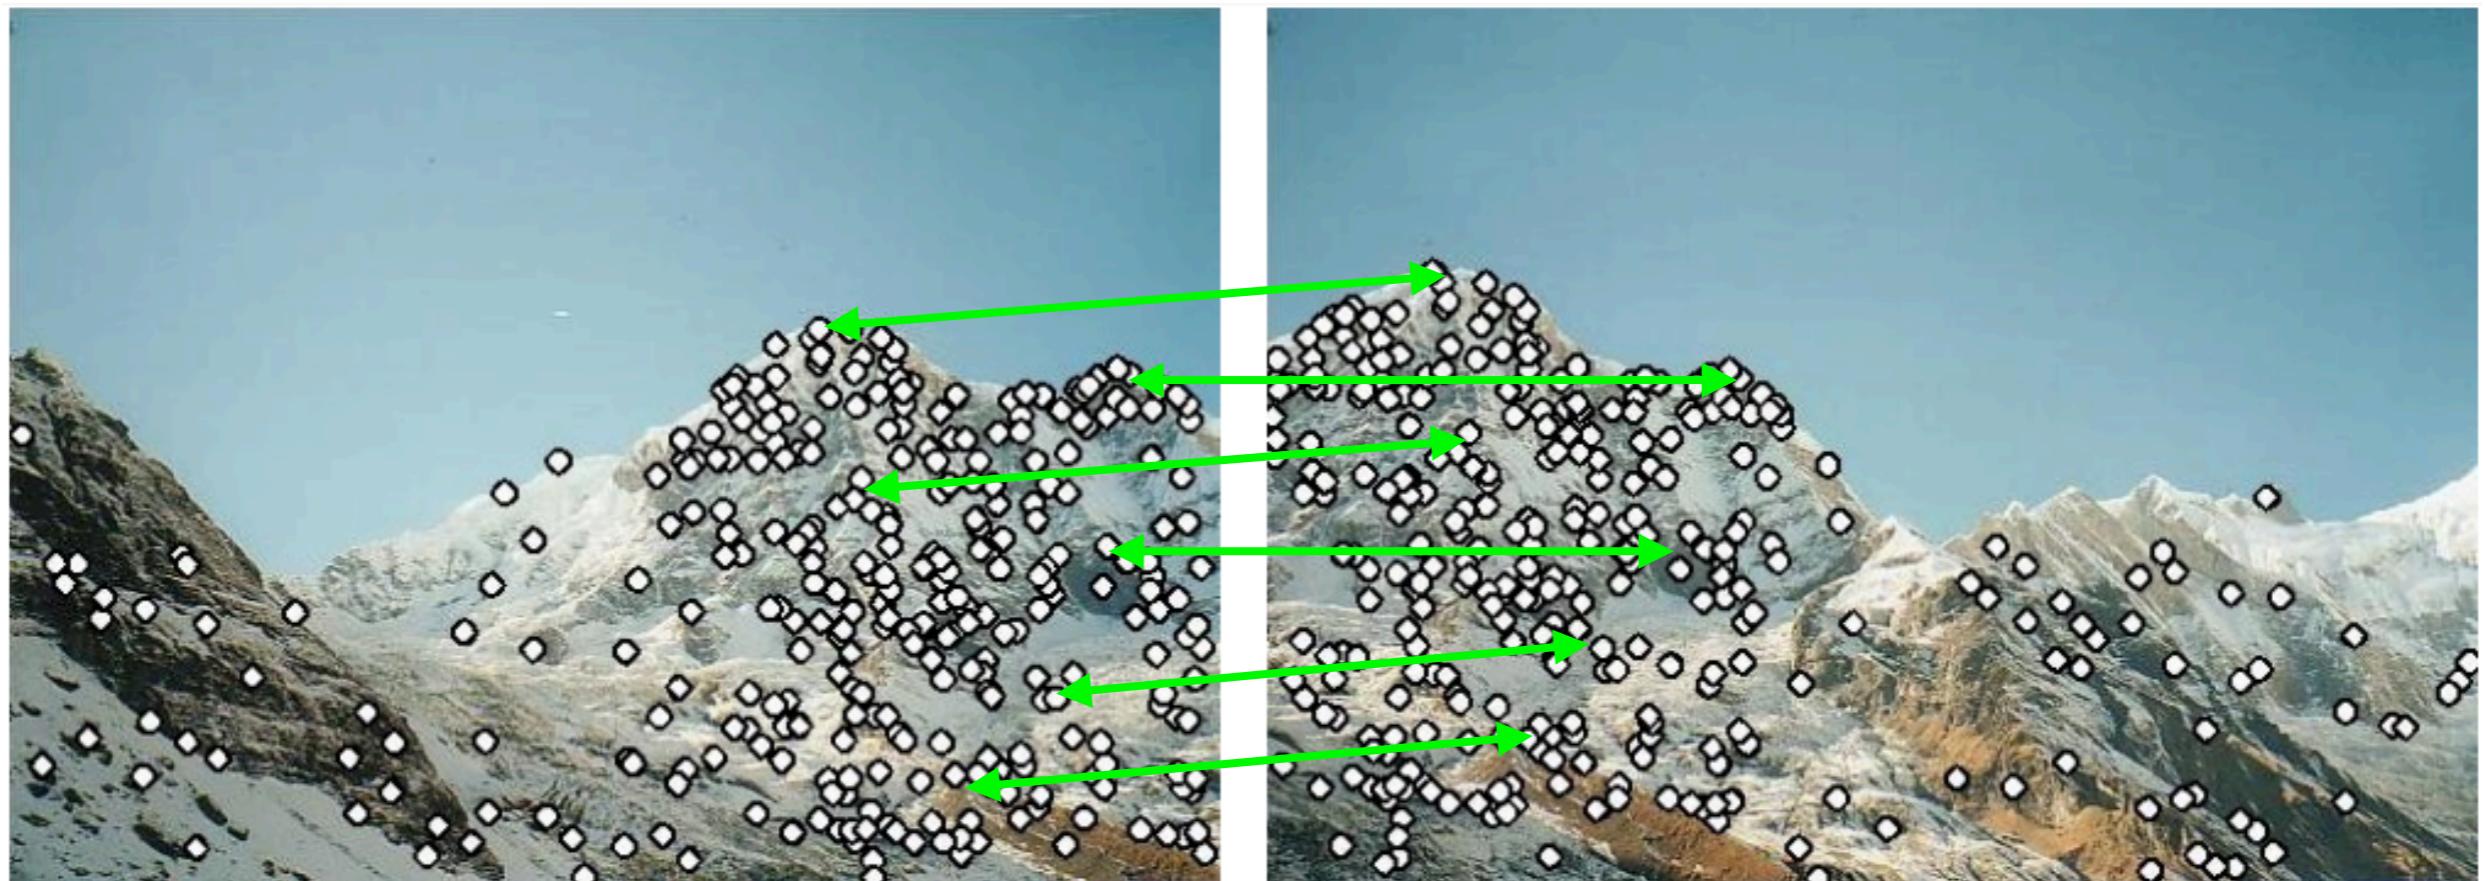
# Why extract keypoints from an image?

- Motivation: a key example is panorama stitching
  - ▶ i.e. we have two images, how to combine them? 



# Why extract keypoints from an image?

- Motivation: a key example is panorama stitching
  - ▶ i.e. we have two images, how to combine them?

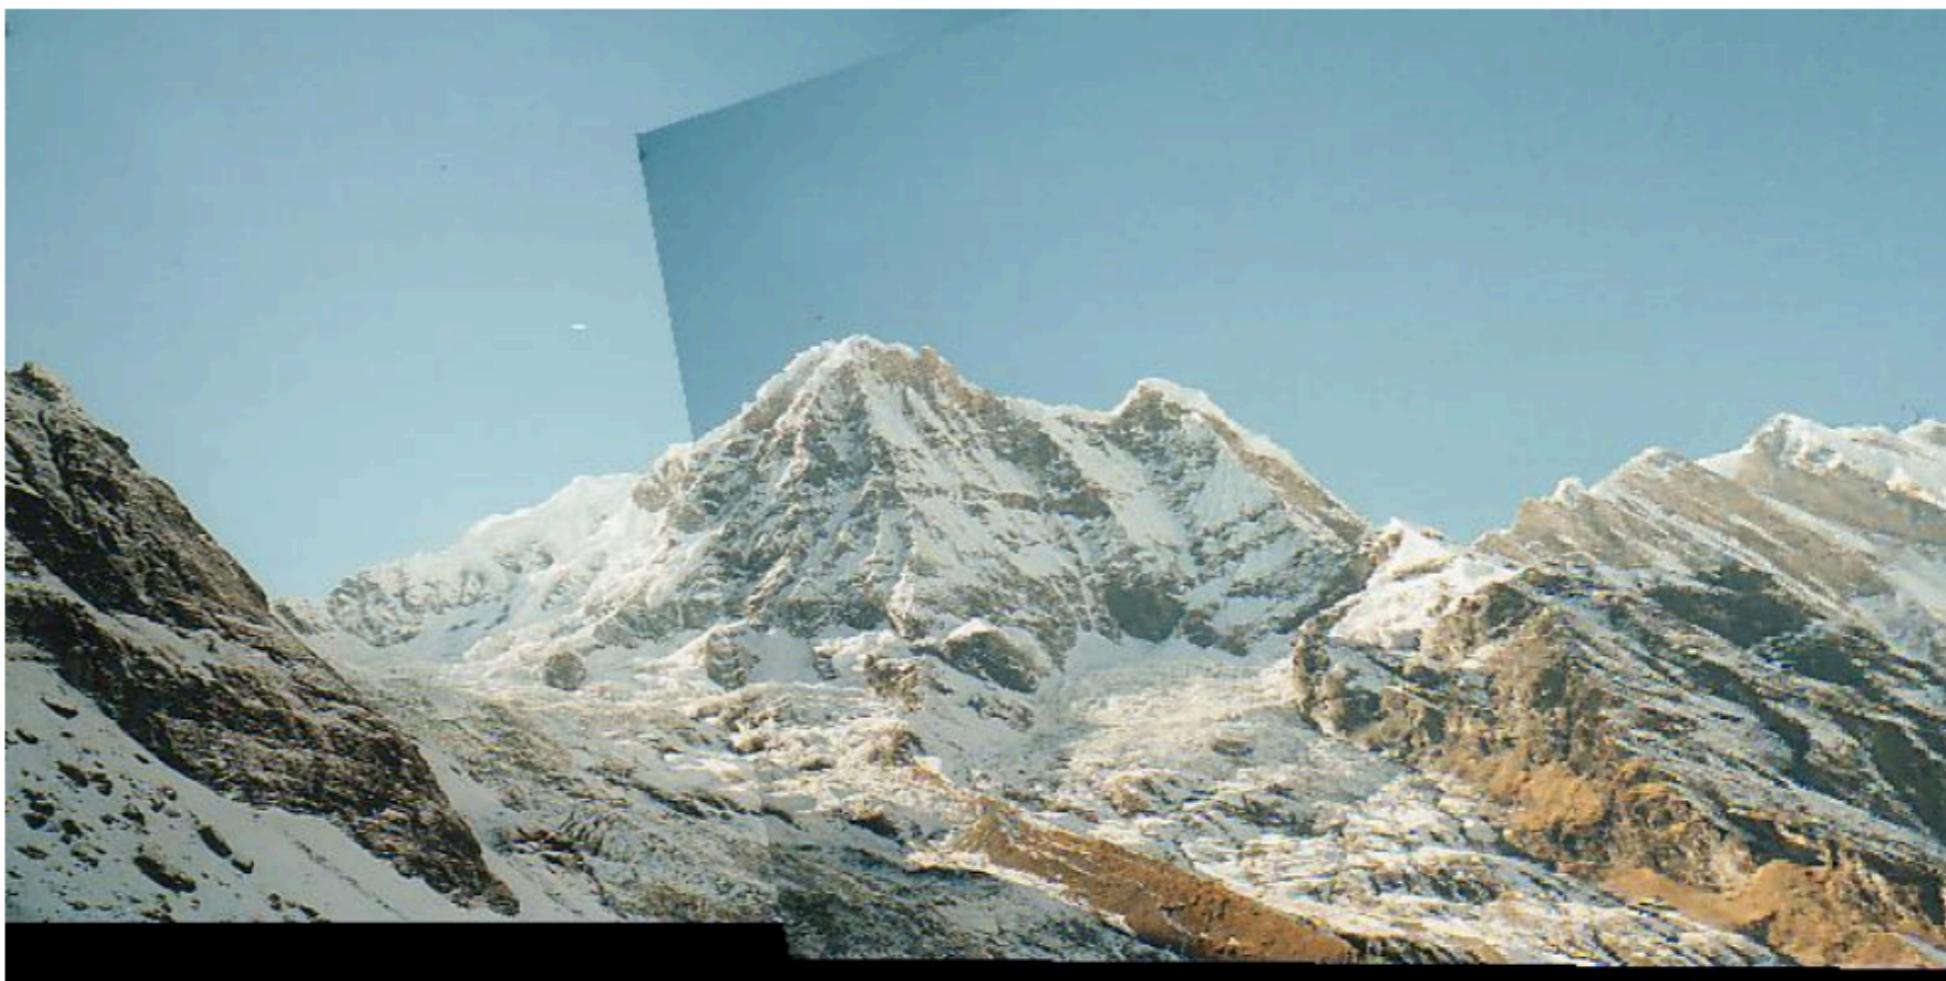


**Step 1:** extract keypoints

**Step 2:** match keypoint features

# Why extract keypoints from an image?

- Motivation: a key example is panorama stitching
  - i.e. we have two images, how to combine them?



**Step 1:** extract keypoints

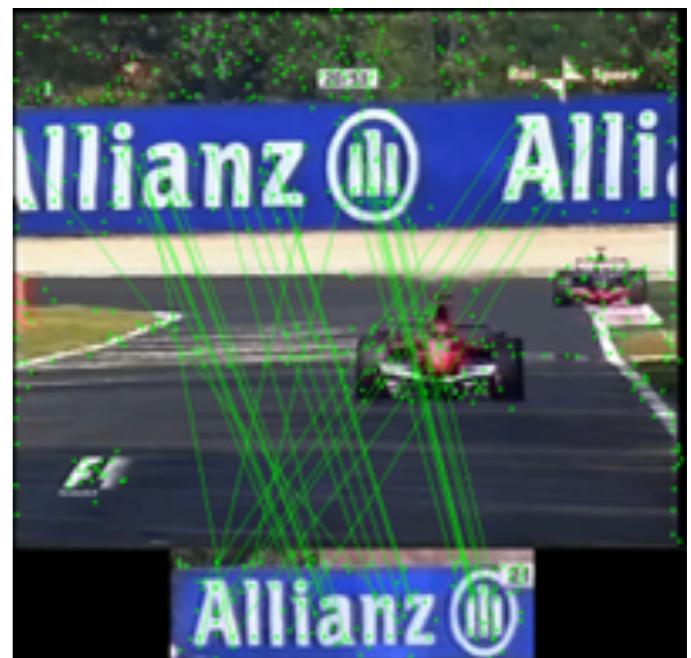
**Step 2:** match keypoint features

**Step 3:** align images

# Keypoints: applications



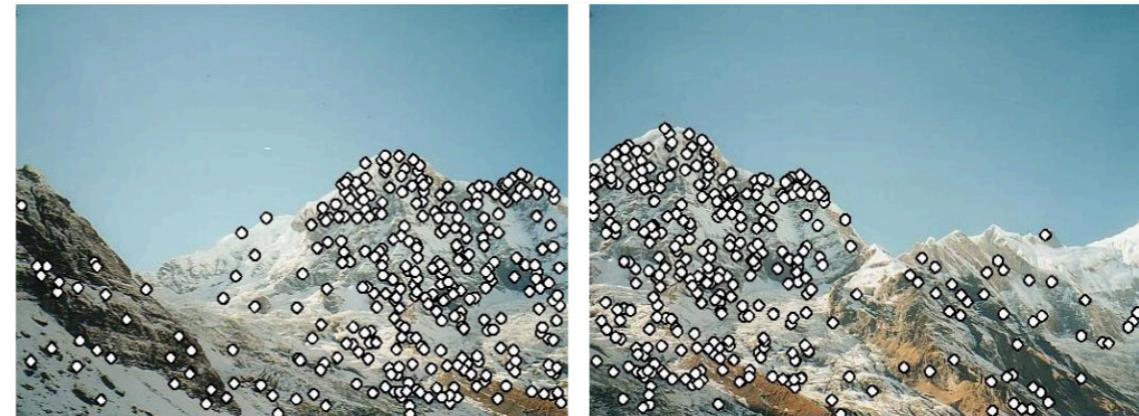
- Keypoints are used for:
  - ▶ Image alignment
  - ▶ 3D reconstruction
  - ▶ Motion tracking
  - ▶ Robot navigation
  - ▶ Indexing and database retrieval
  - ▶ Object recognition



# Characteristics of Good Keypoints



- Repeatability
  - Can be found despite geometric/photometric transformations
- Salience
  - Each keypoint is distinctive
- Compactness and efficiency
  - Many fewer keypoints than image pixels
- Locality
  - Occupies small area of the image: robust to clutter & occlusion

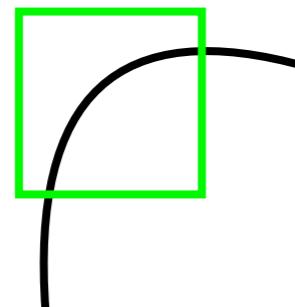


# An example: Harris corner detector

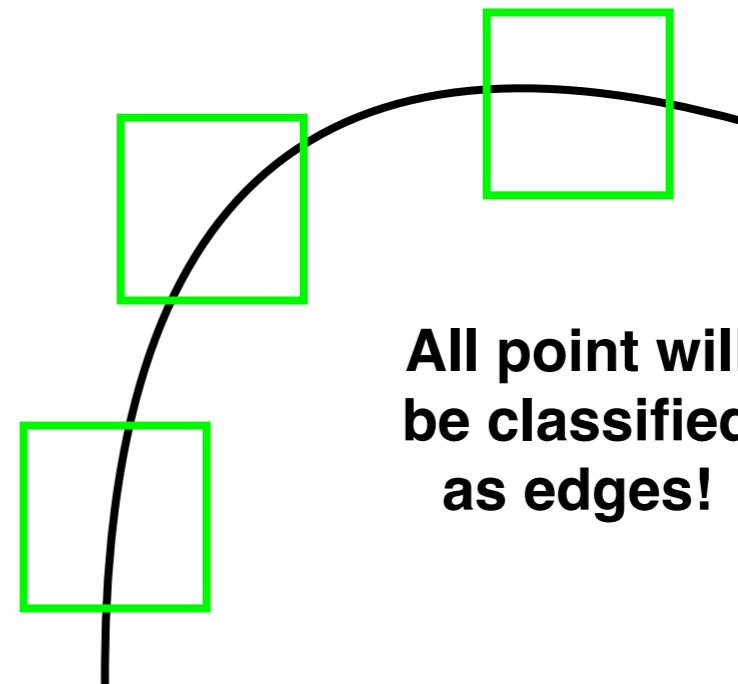
- Properties (i.e. geometric invariances):



- Translation invariance? Yes
- Rotation invariance? Yes
- Scale invariance? No



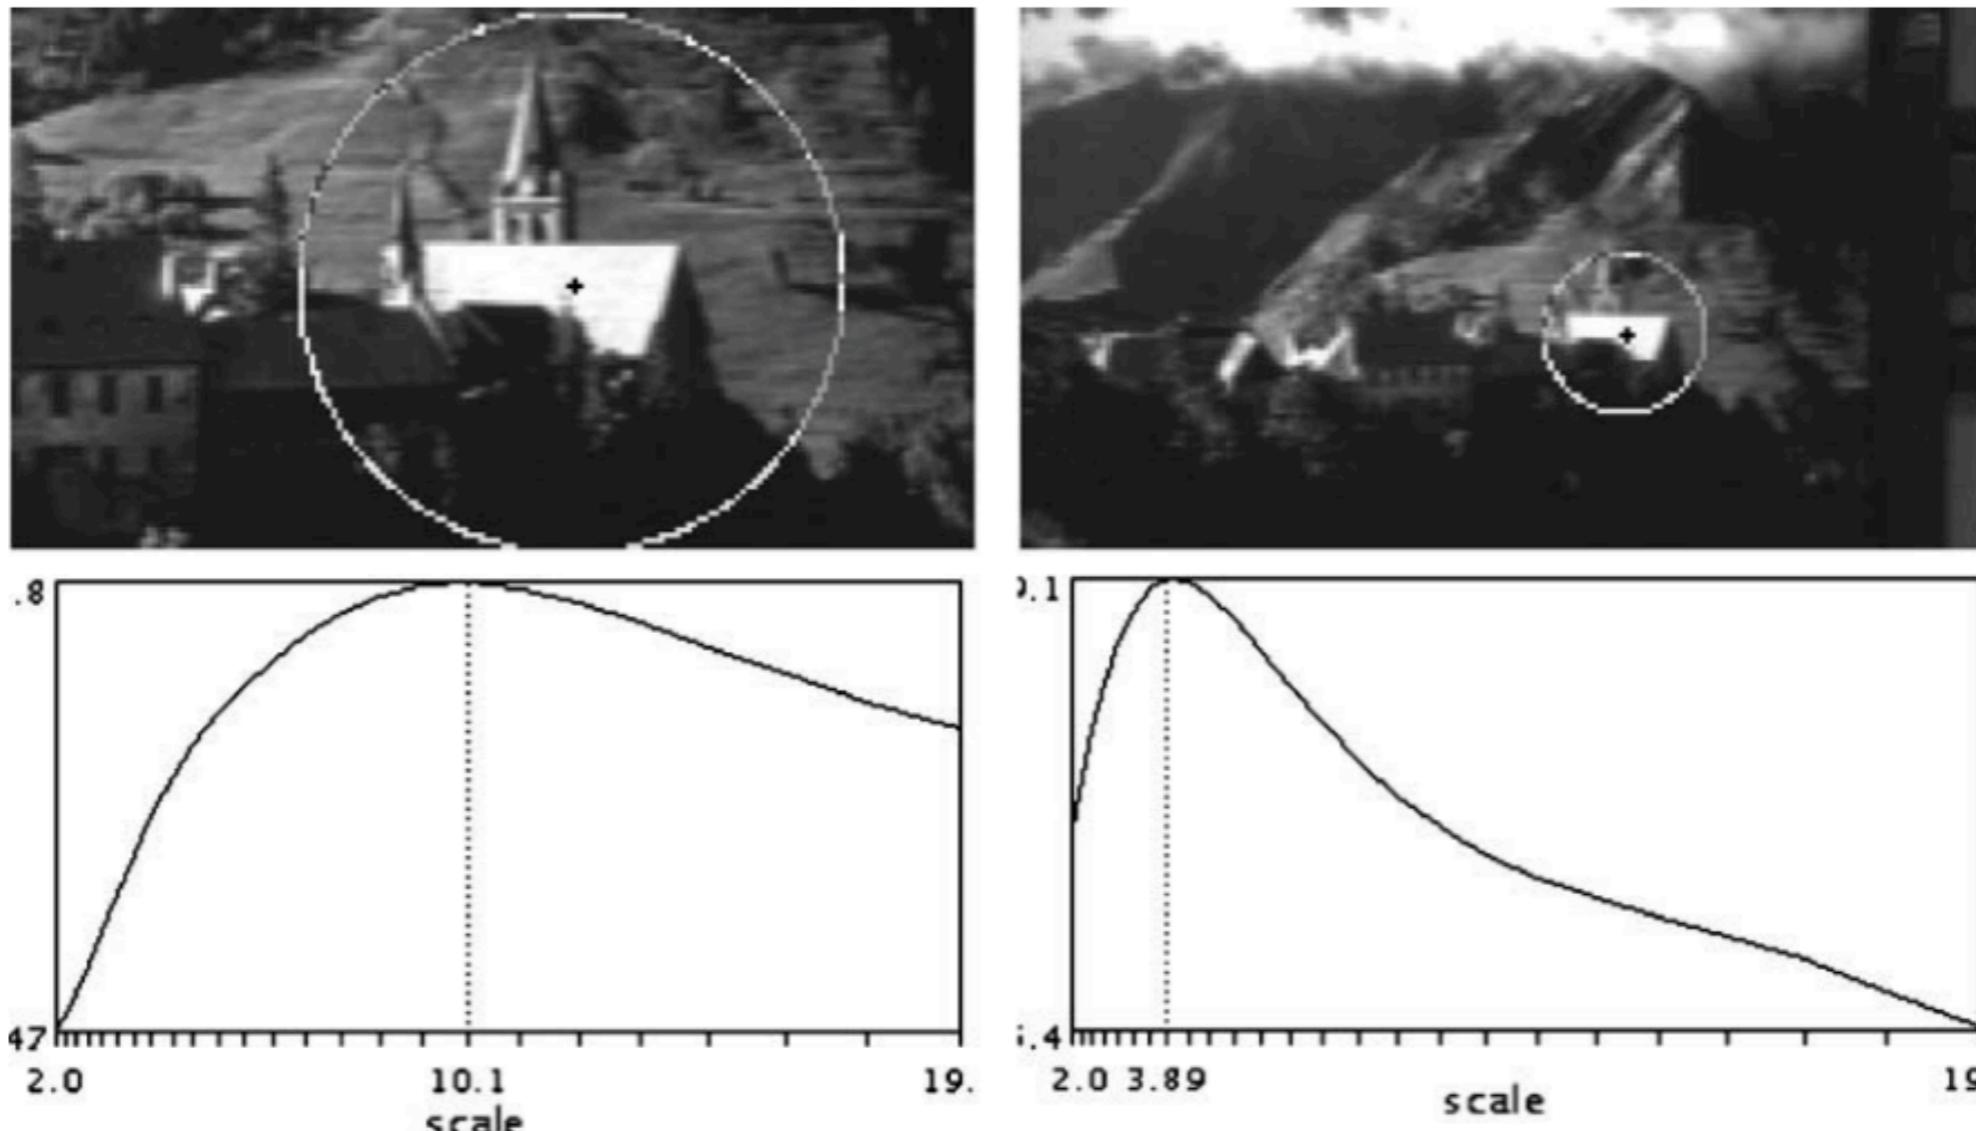
Corner



All point will  
be classified  
as edges!

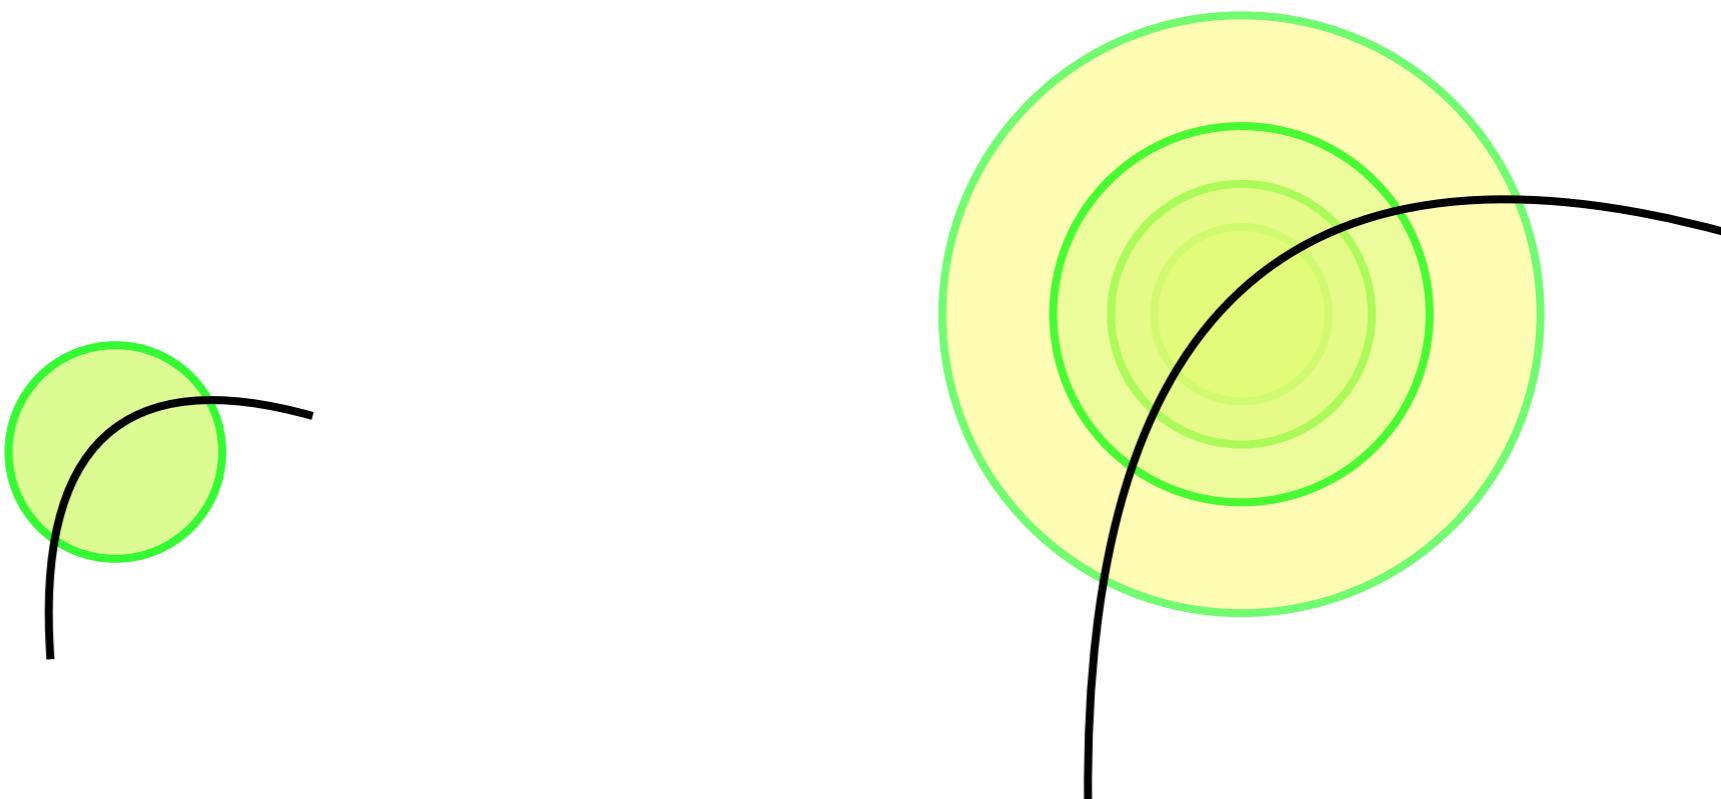
# Feature detection with scale selection

- We want to extract features with characteristic scale that is invariant (*covariant*) with the image transformation



# Scale invariant detection

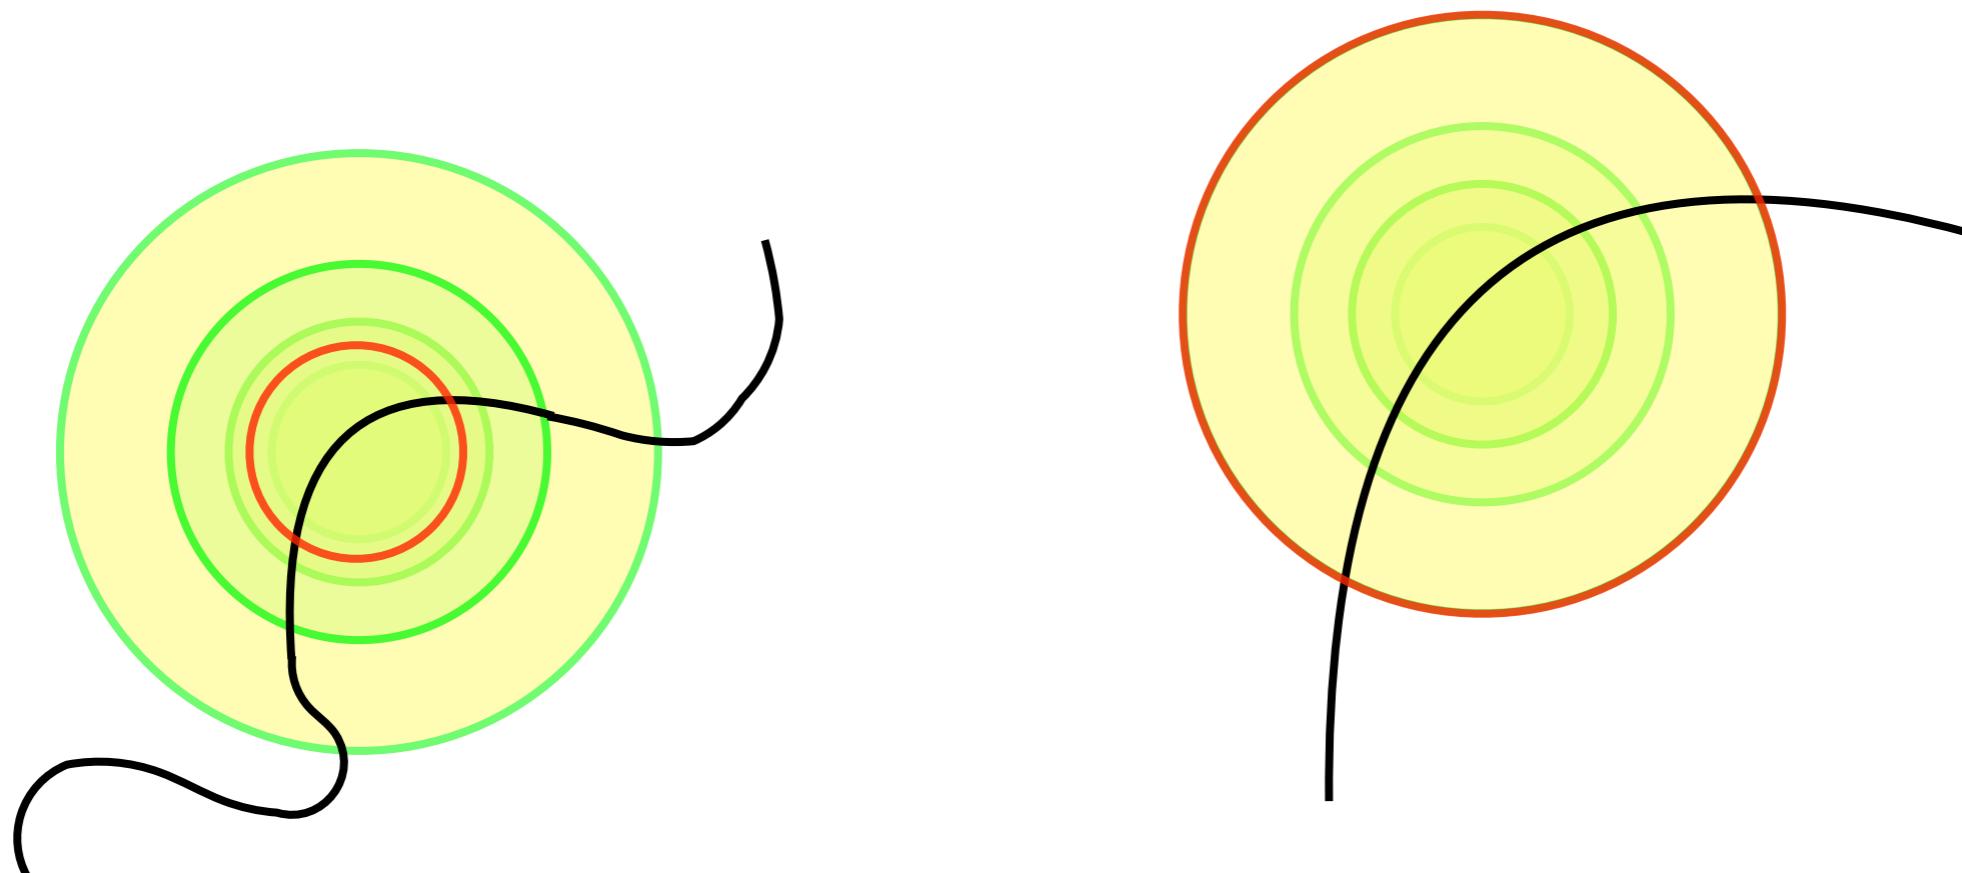
- Consider regions (e.g. circles) of different sizes around a point
- How to find regions of corresponding sizes that will look the same in both images?



# Scale invariant detection



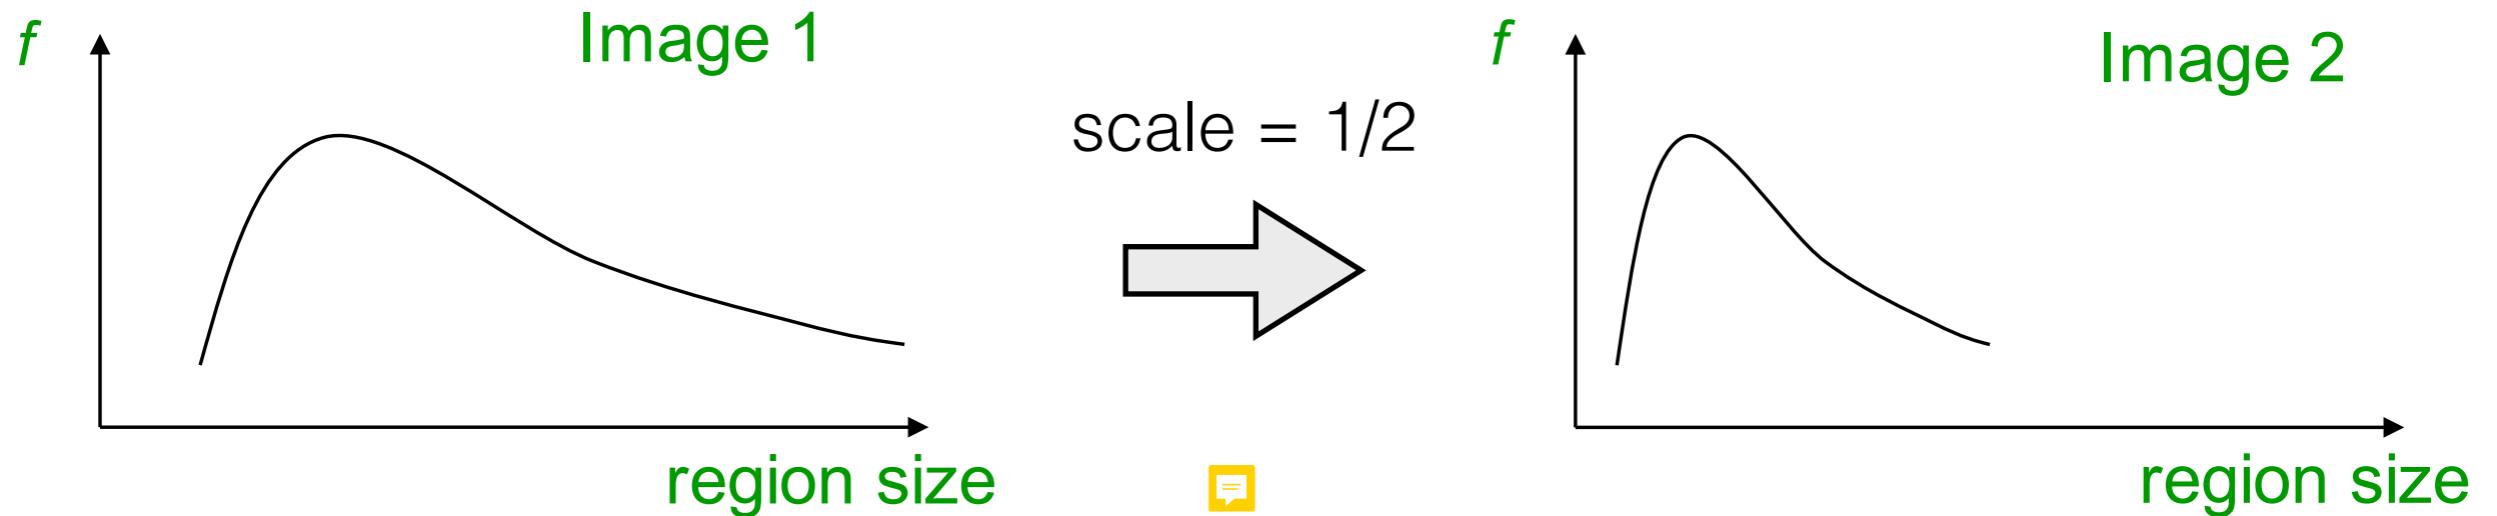
- The problem: how do we choose corresponding circles **independently** in each image?





# Scale invariant detection

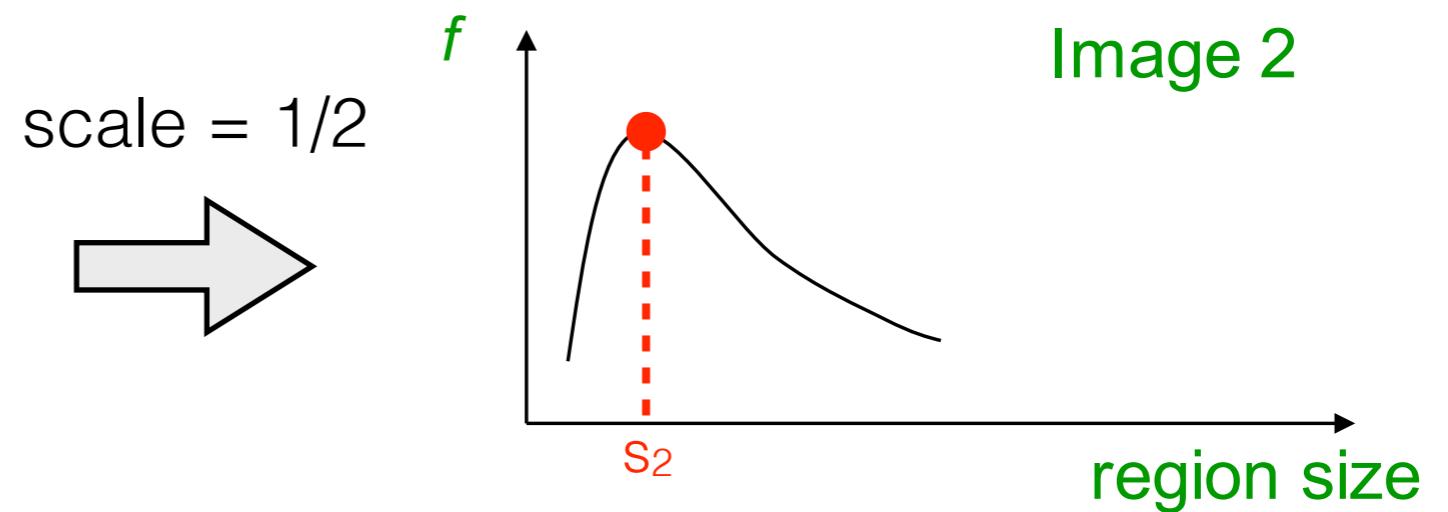
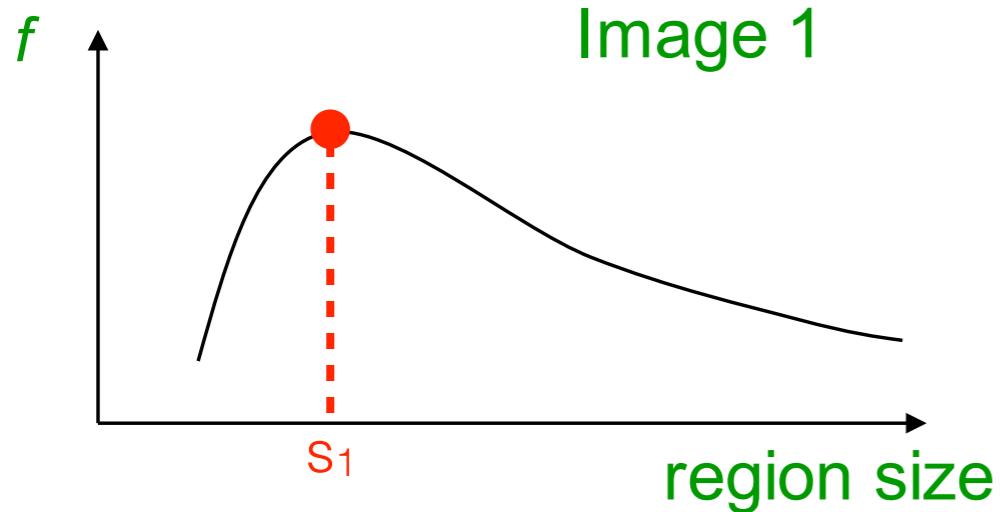
- Solution: design a function on the region (circle) which is scale invariant (*i.e.* the same for corresponding regions, even if they are at different scales)
  - ▶ e.g. average intensity (same even for different sizes)
- For a point in one image, we can consider it as a function of region size (circle radius)



# Scale invariant detection

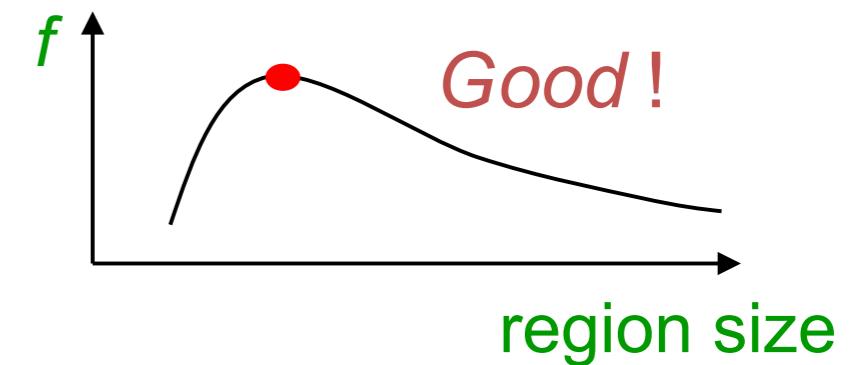
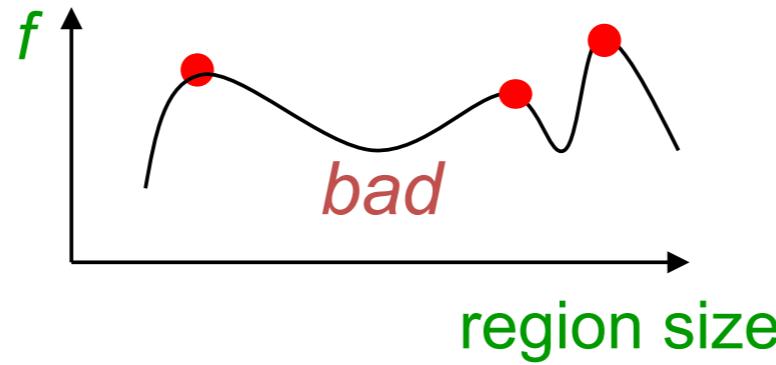
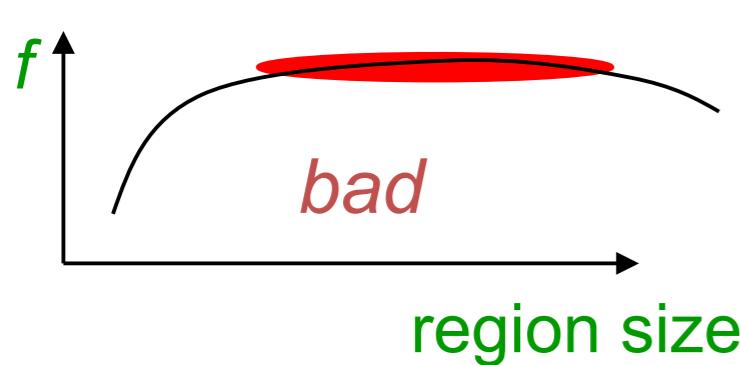
- Common approach: take a local max of this function
- Observation: region size for which the maximum is achieved should be *invariant (co-variant)* with scale

Important: this scale invariant region size is found in each image independently!



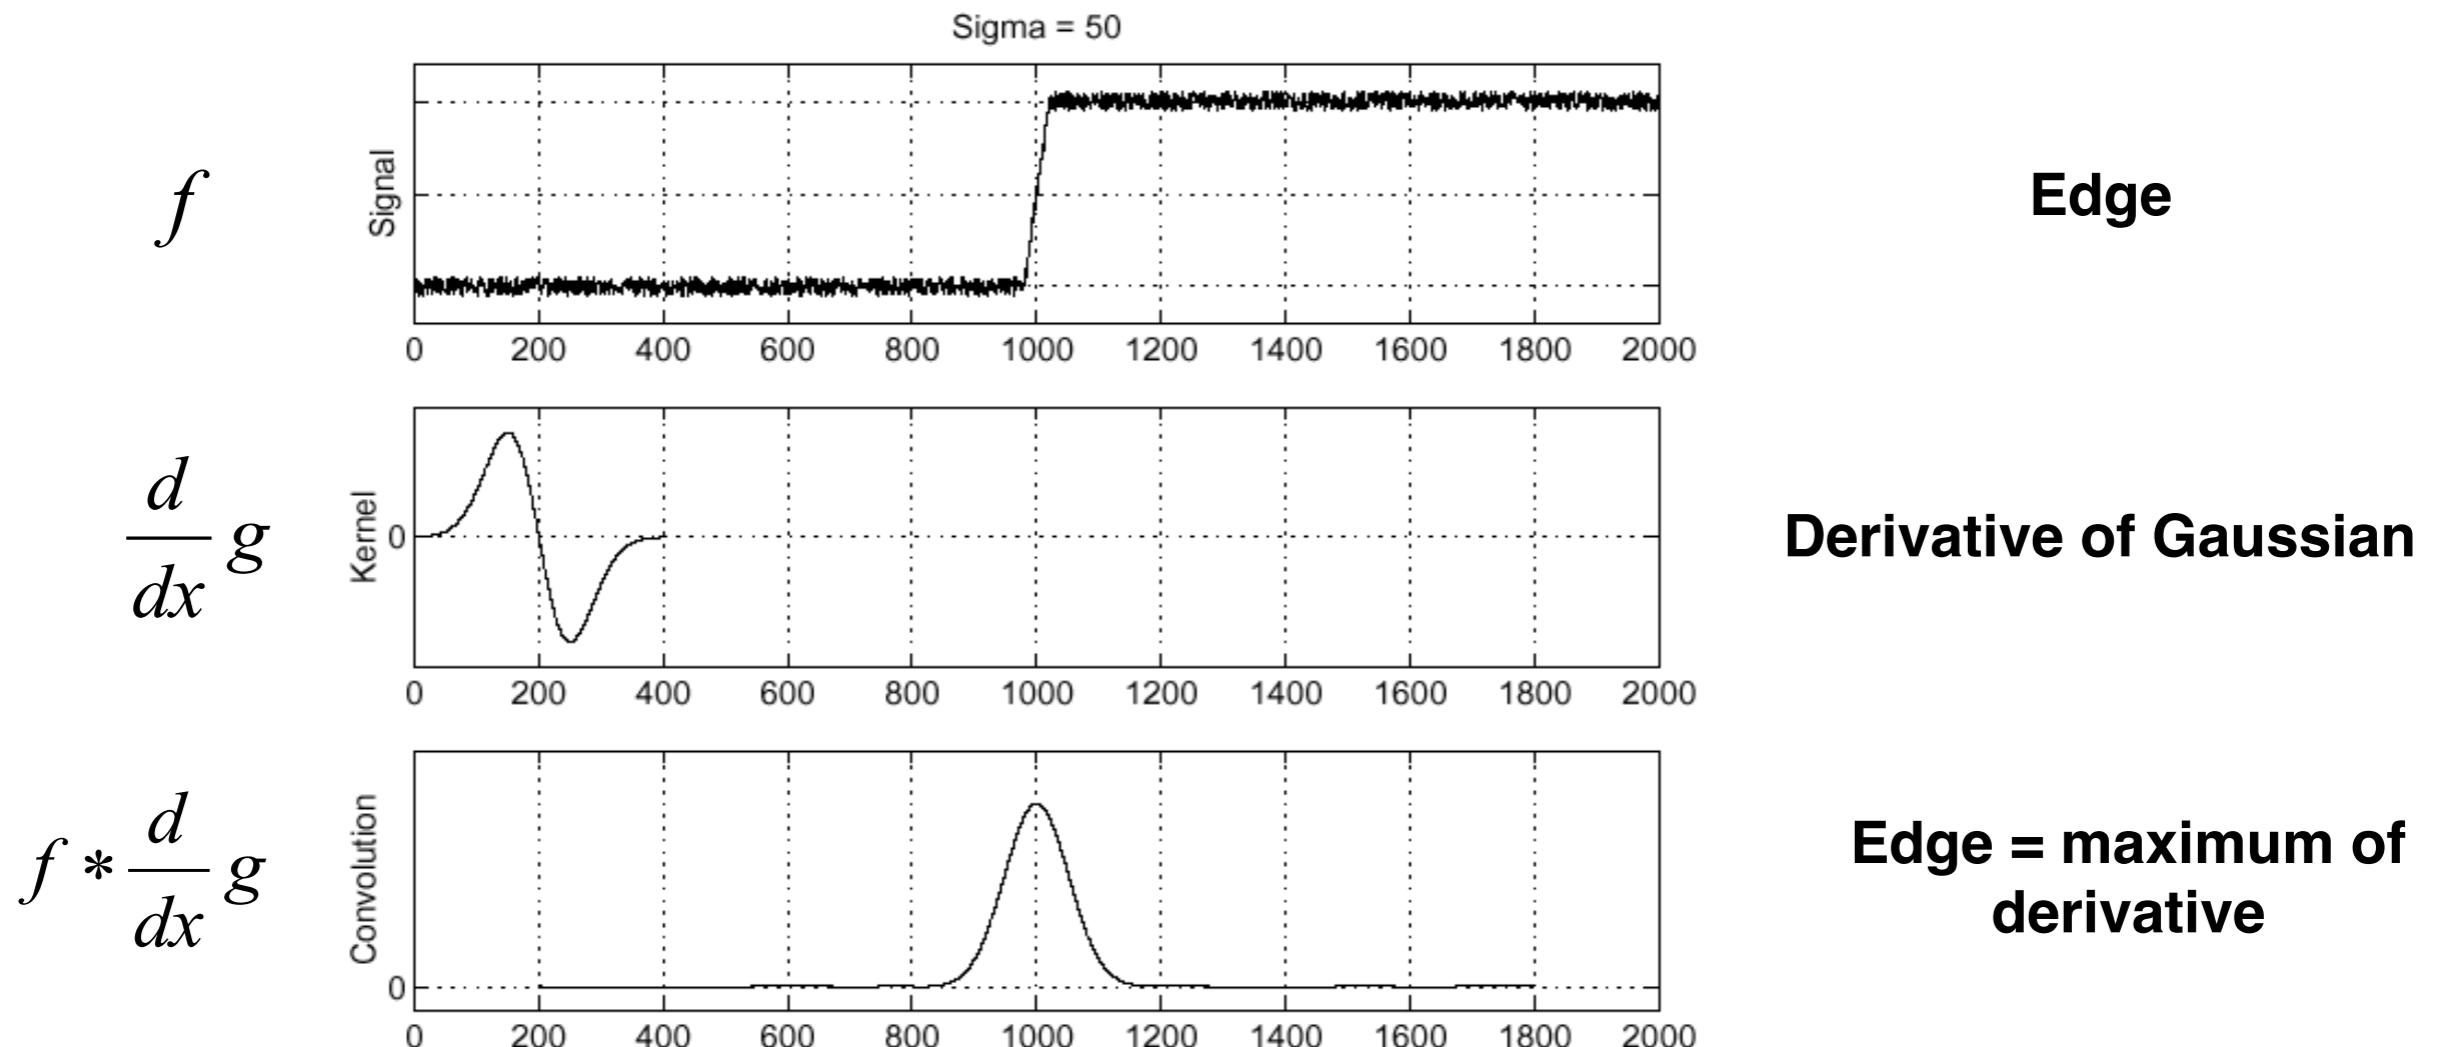
# Scale invariant detection

- A good function for scale detection has one stable sharp peak

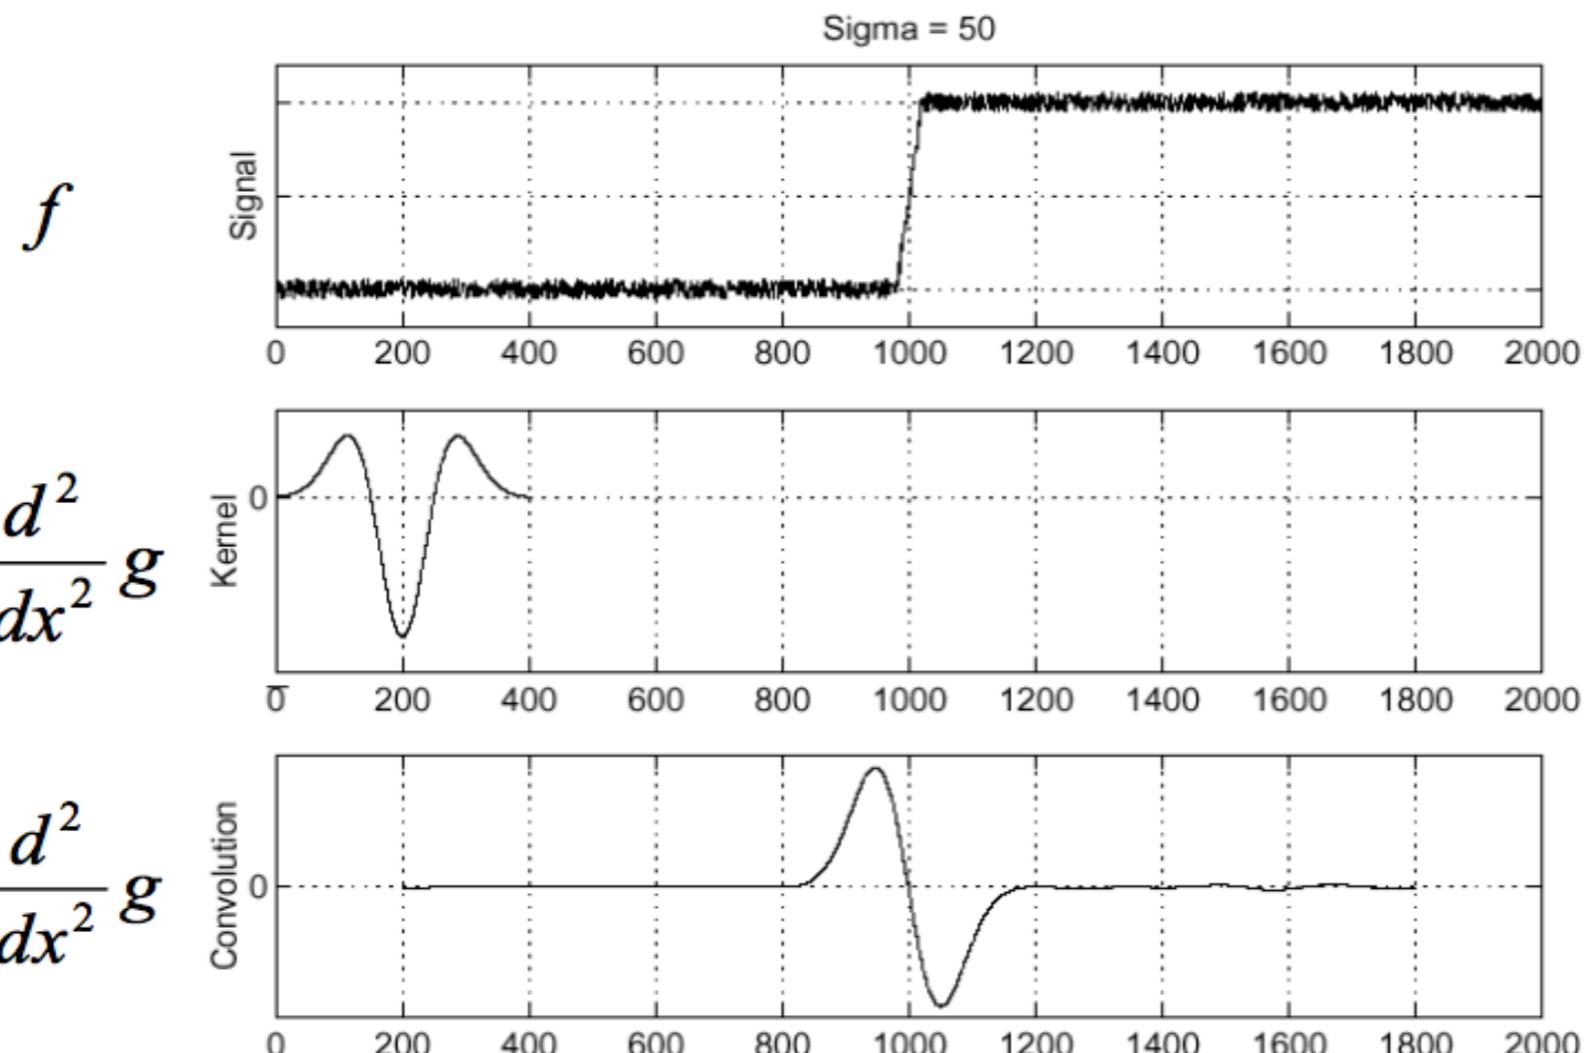


- So a good function would be one which responds to contrast (sharp local intensity change)

# Recall: edge detection



# A bit more on edge detection



**Edge**

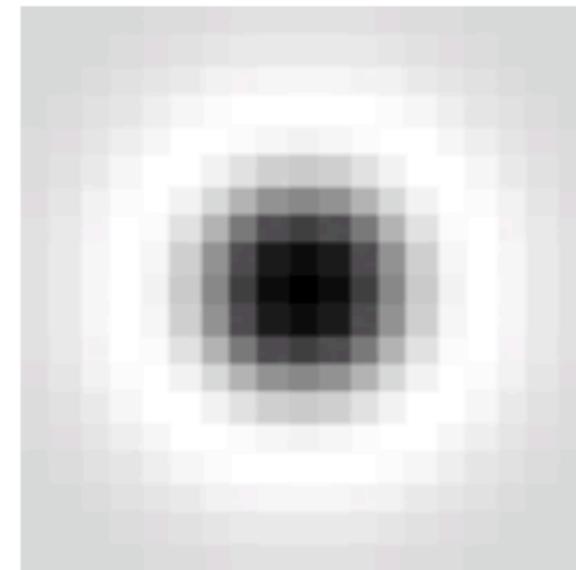
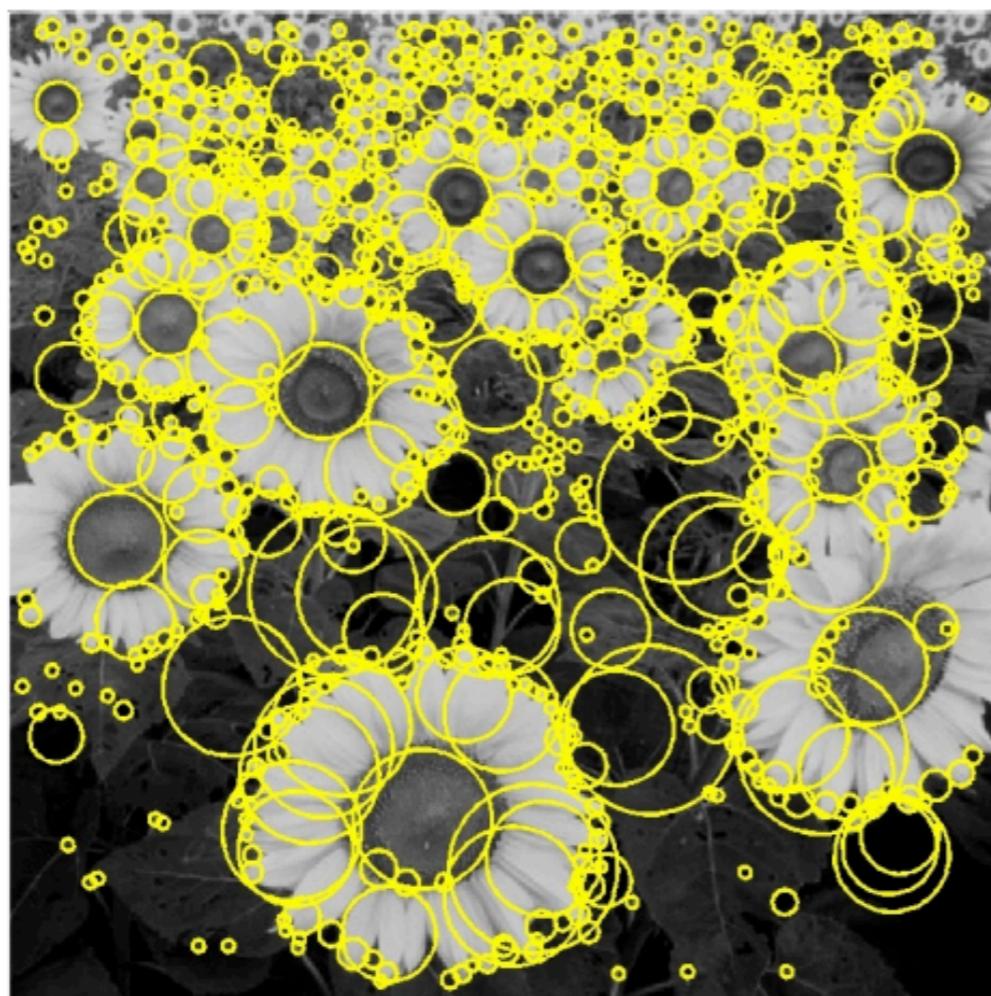
**Second Derivative of Gaussian (Laplacian)**

*Mexican hat filter*

**Edge = zero crossing of second derivative**

# Blob detection: basic idea

- To detect “blobs”, convolve the image with a blob filter at multiple scales and look for extrema of filter response in the resulting *scale space*



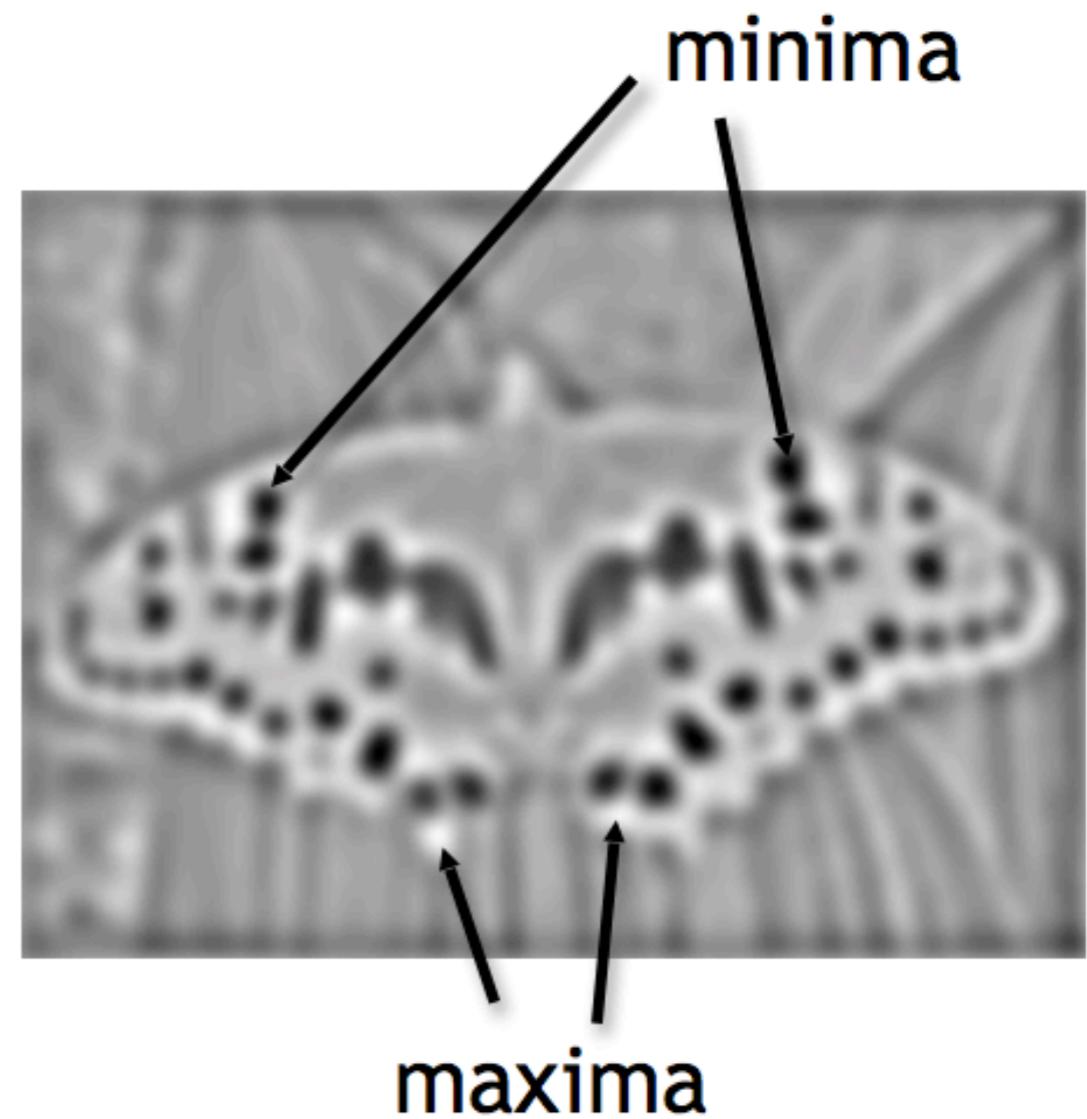
*Blob filter*

# Blob detection: basic idea

- Find maxima and minima of blob filter response in space and scale



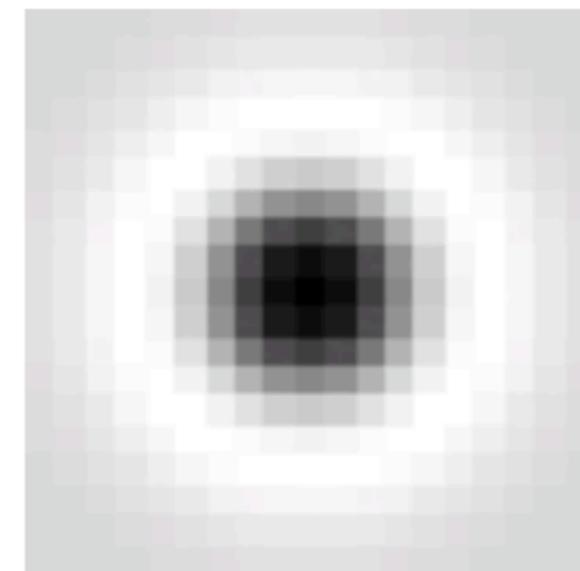
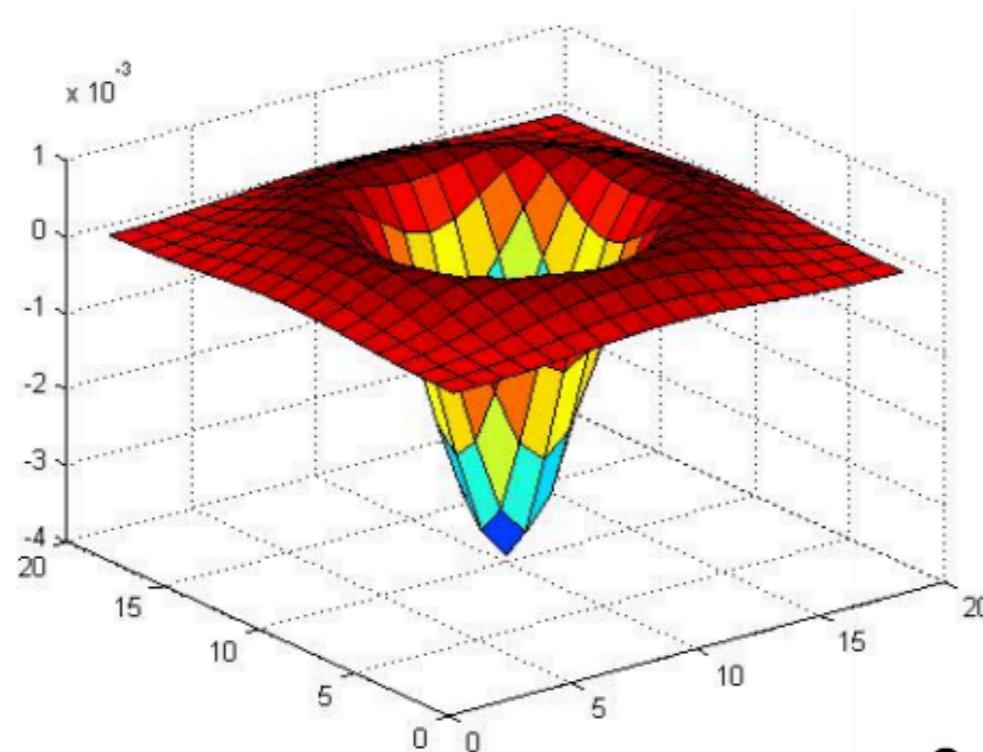
$$* \quad \bullet =$$



# Blob filter



- **Laplacian of Gaussian (LoG):** circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

# Scale invariant detection

- Functions for determining scale:  $f = \text{Kernel} * \text{Image}$

**Kernels:**



$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

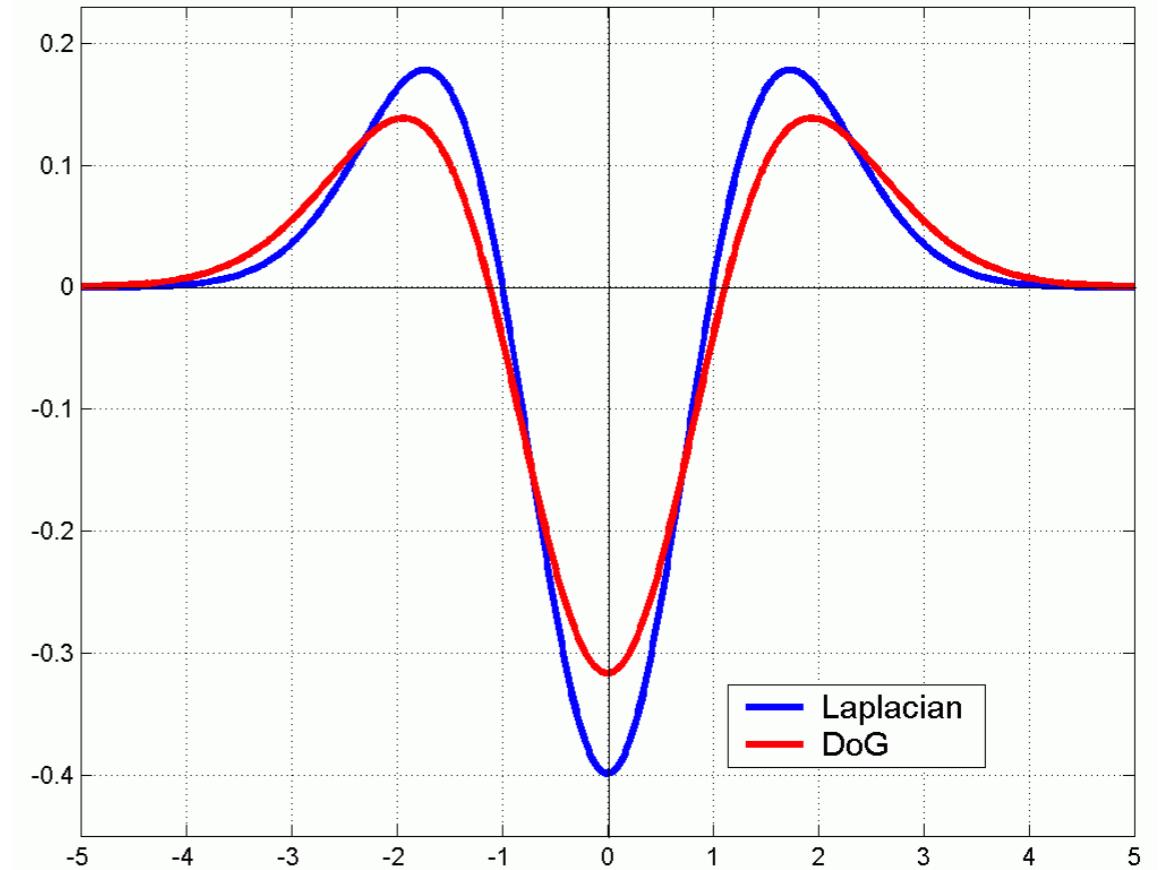
(Laplacian)

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)

where Gaussian

$$G(x, y, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

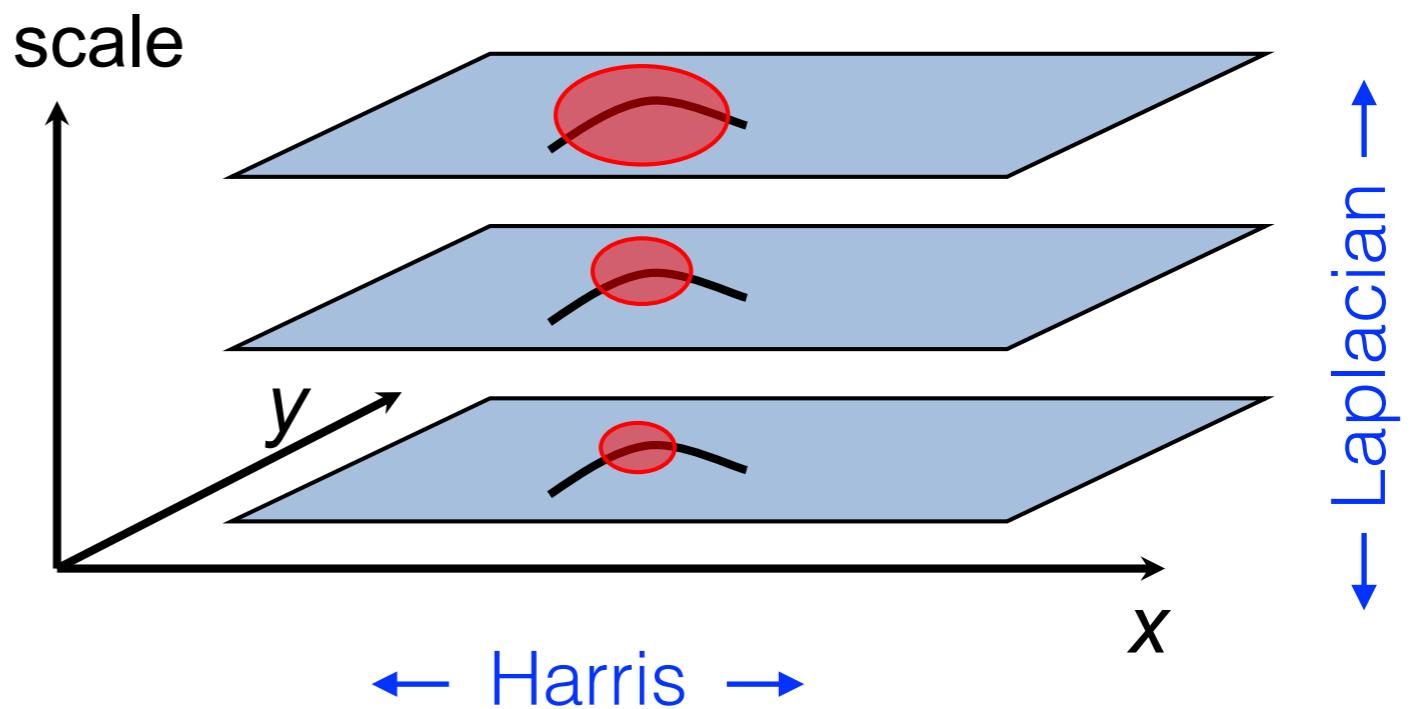


Note: both kernels are invariant to scale and rotation

# Scale invariant detection



- **Harris-Laplacian** detector: find local maximum of
  - Harris corner detector in space (image coordinates)
  - Laplacian in scale

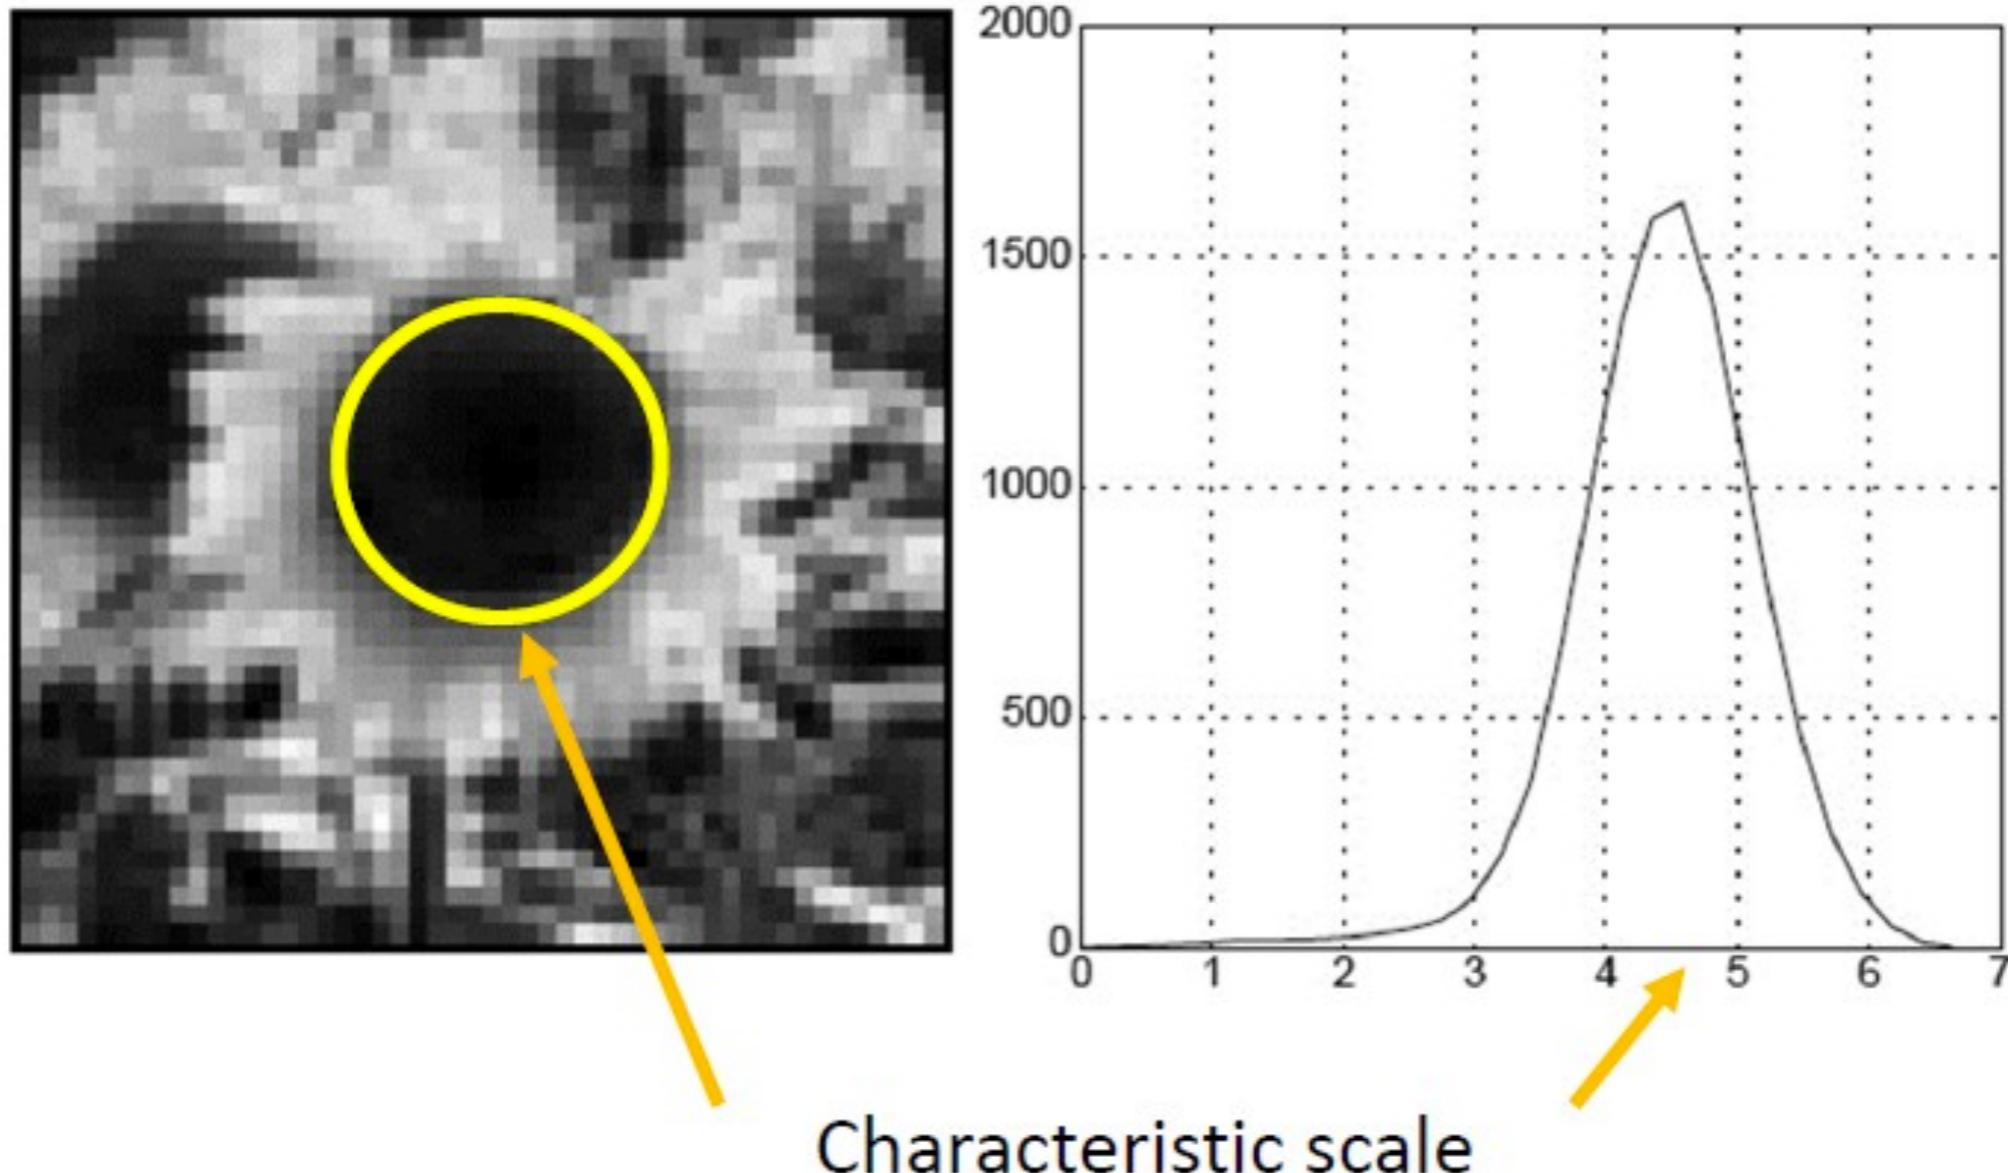


K.Mikolajczyk, C.Schmid, “Indexing Based on Scale Invariant Interest Points”, ICCV 2001

# Scale invariant detection



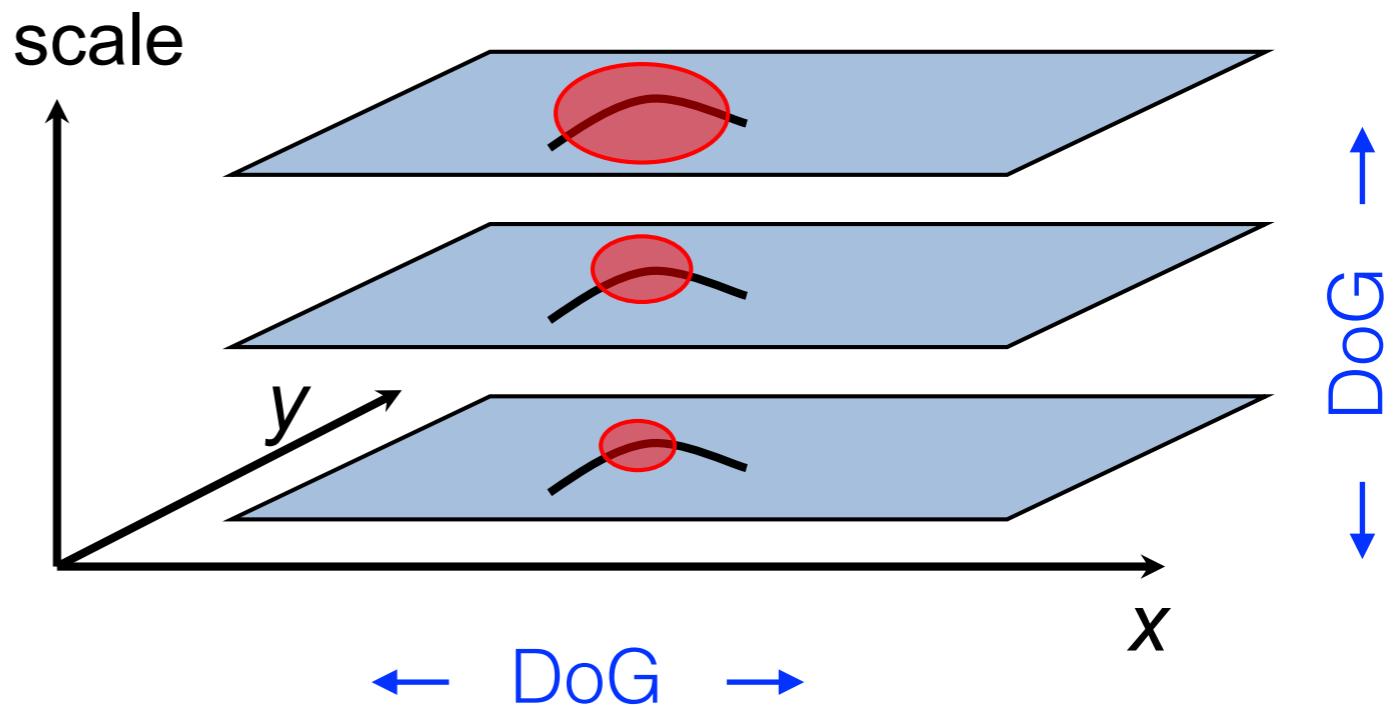
- An example (Laplacian):



# Scale invariant detection



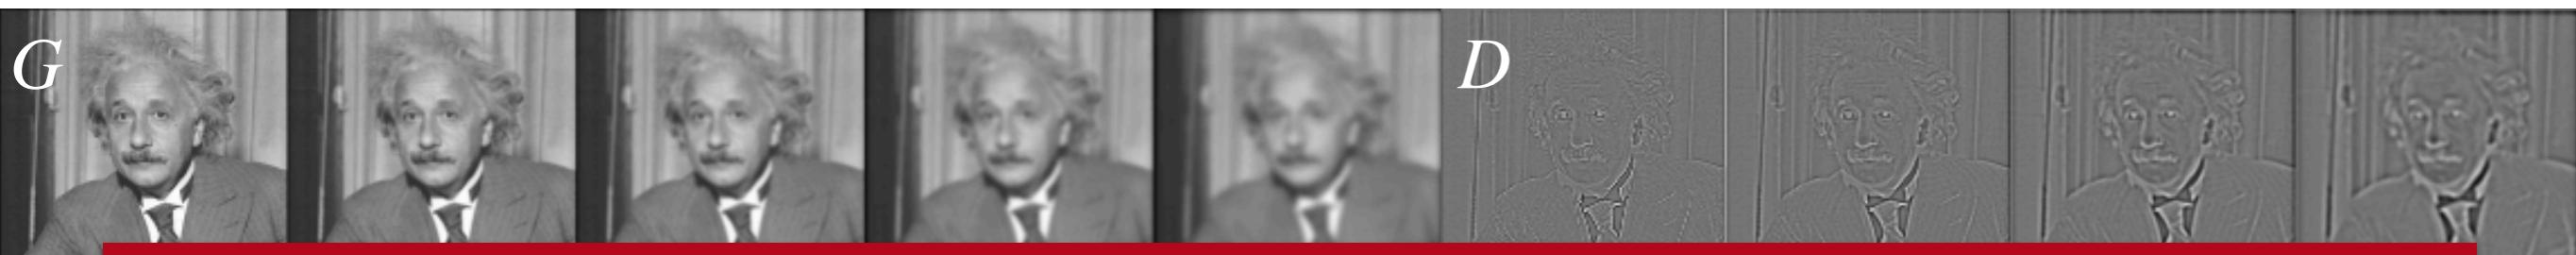
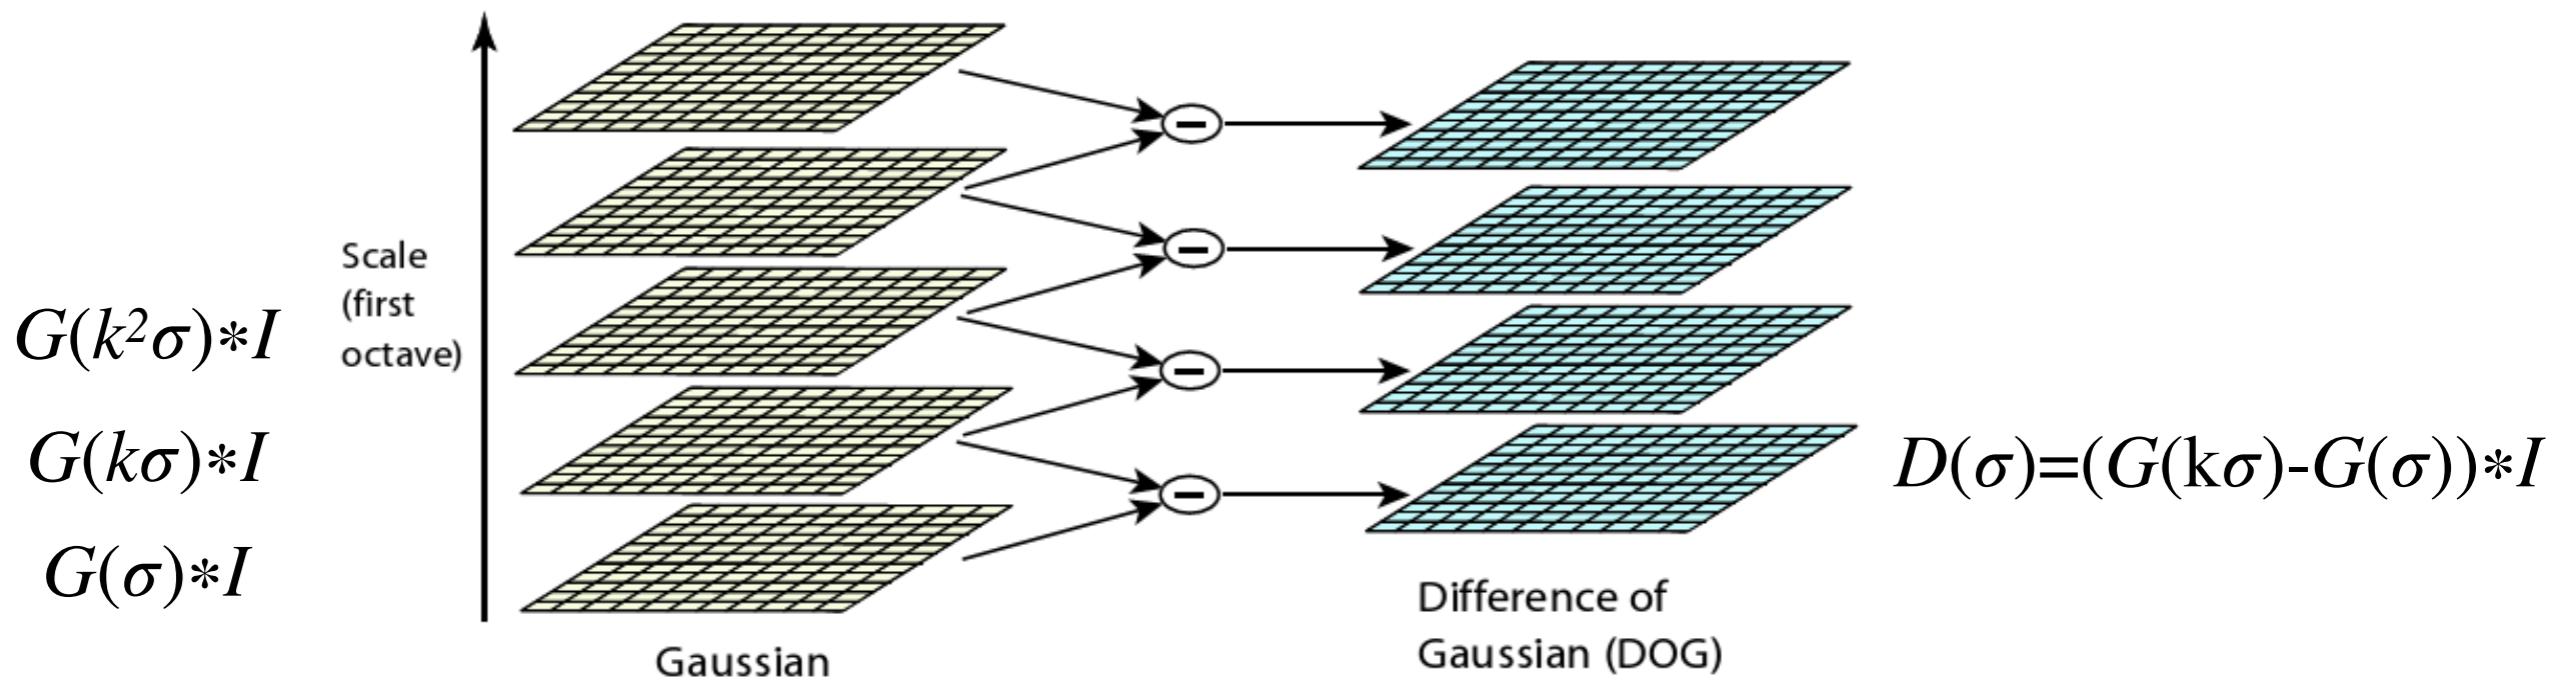
- **SIFT** detector: find local maximum of
  - ▶ Difference of Gaussians in space
  - ▶ Difference of Gaussians in scale



D. Lowe, “Distinctive image features from scale-invariant keypoints”, IJCV 2004

# Difference of Gaussians (SIFT detector)

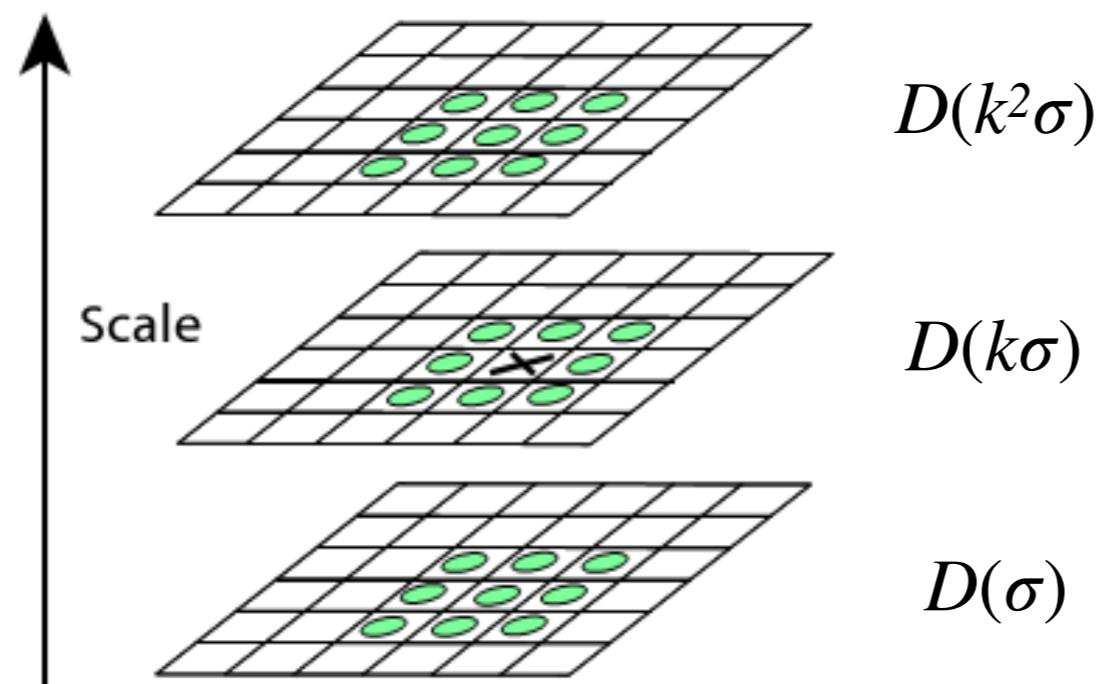
- Efficient approximation of LoG



Key idea: At different resolutions of kernel size, we see different fine details of the image. In other words, we can capture keypoints at varying scales.

# Difference of Gaussians (SIFT detector)

- Keypoint are selected as scale-space extrema
  - ▶ choose all extrema within  $3 \times 3 \times 3$  neighborhood



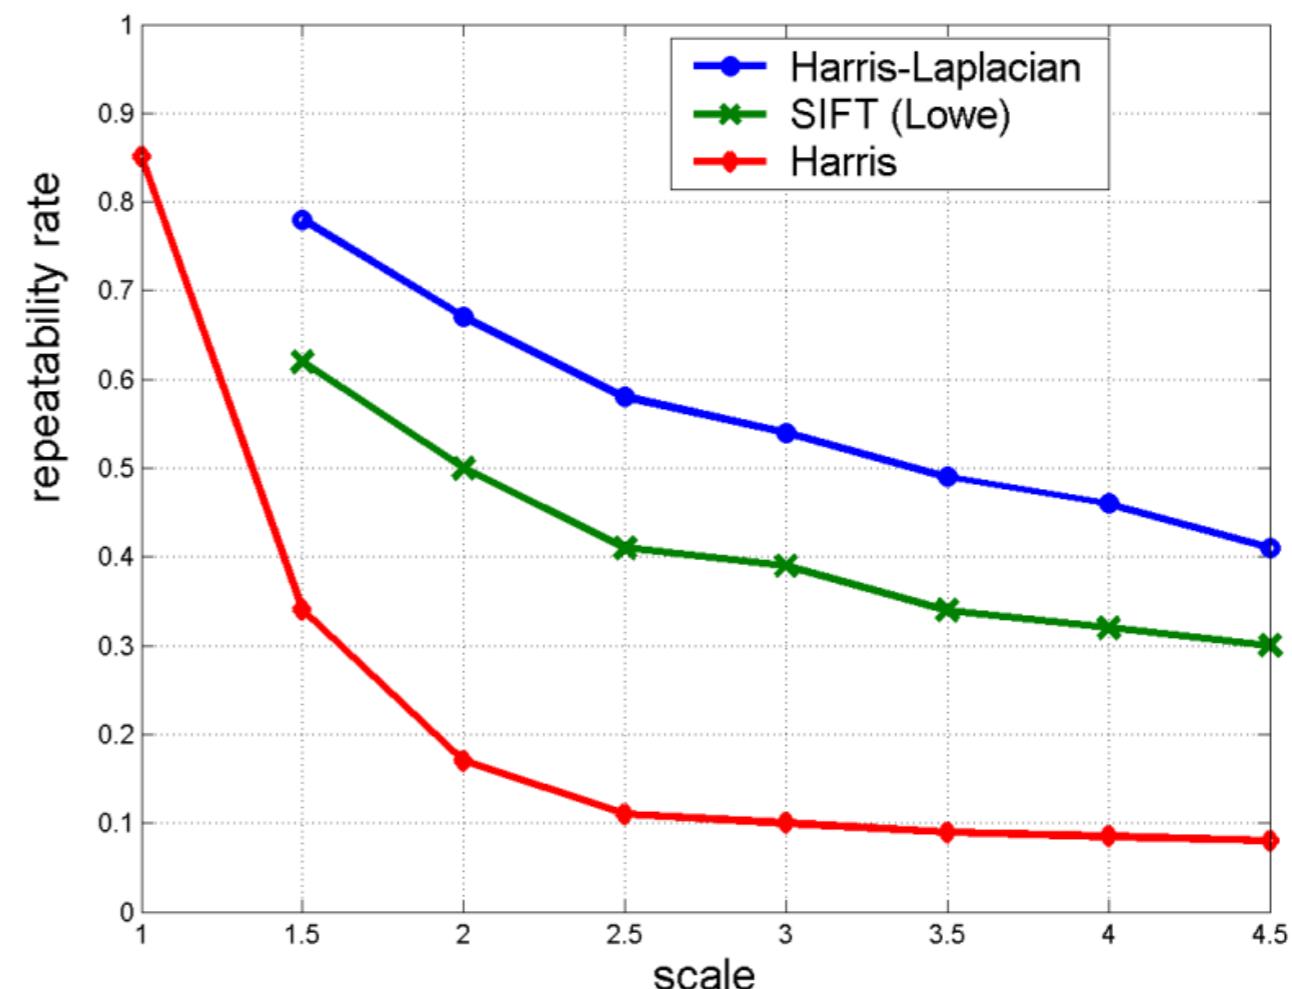
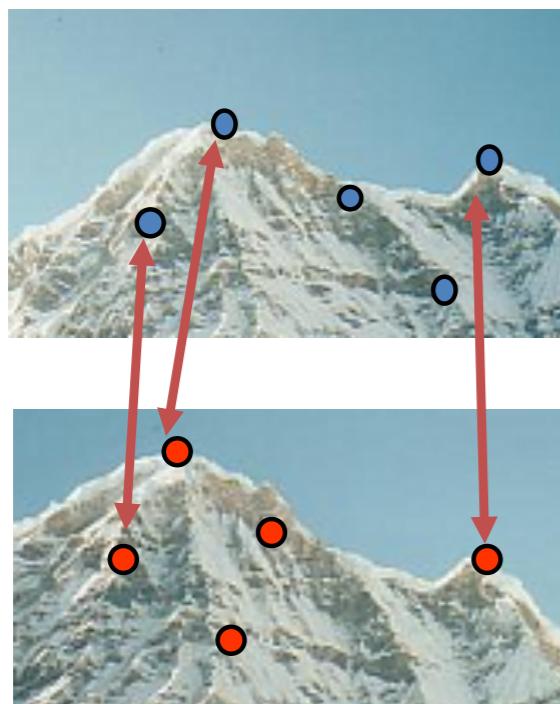
$X$  is selected if it is larger or smaller than all 26 neighbors

# Scale Invariant Detectors



- Experimental evaluation of detectors w.r.t. scale change

Repeatability rate:  
$$\frac{\# \text{ correspondences}}{\# \text{ possible correspondences}}$$



K. Mikolajczyk, C. Schmid, "Indexing Based on Scale Invariant Interest Points", ICCV 2001

# Scale Invariant Detectors

- Scale invariant detection summary:
  - **Given:** two images of the same scene with a large scale difference between them
  - **Goal:** find the same interest points independently in each image
  - **Solution:** search for maxima of suitable functions in scale and in space (over the image)

## Popular methods:

1. Harris-Laplacian [Mikolajczyk, Schmid]: maximize Laplacian of Gaussian over scale, Harris' measure of corner response over the image
2. SIFT [Lowe]: maximize Difference of Gaussians over scale and space

# Scale Invariant Detectors

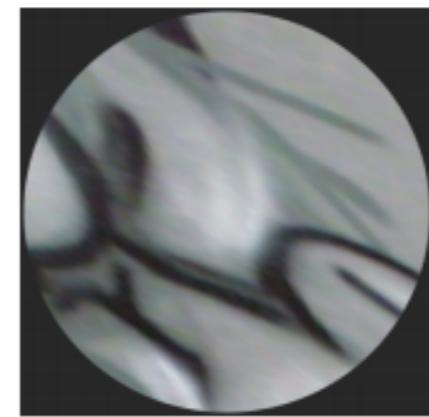
- Actually there are many other options...
- List/comparison of keypoint detectors:

Feature Detector	Corner	Blob	Region	Rotation invariant	Scale invariant	Affine invariant	Repeatability	Localization accuracy	Robustness	Efficiency
Harris	√			√			+++	+++	+++	++
Hessian		√		√			++	++	++	+
SUSAN	√			√			++	++	++	+++
Harris-Laplace	√	(√)		√	√		+++	+++	++	+
Hessian-Laplace	(√)	√		√	√		+++	+++	+++	+
DoG	(√)	√		√	√		++	++	++	++
SURF	(√)	√		√	√		++	++	++	+++
Harris-Affine	√	(√)		√	√	√	+++	+++	++	++
Hessian-Affine	(√)	√		√	√	√	+++	+++	+++	++
Salient Regions	(√)	√		√	√	(√)	+	+	++	+
Edge-based	√			√	√	√	+++	+++	+	+
MSER		√		√	√	√	+++	+++	++	+++
Intensity-based		√		√	√	√	++	++	++	++
Superpixels		√		√	(√)	(√)	+	+	+	+

# From feature detection to description

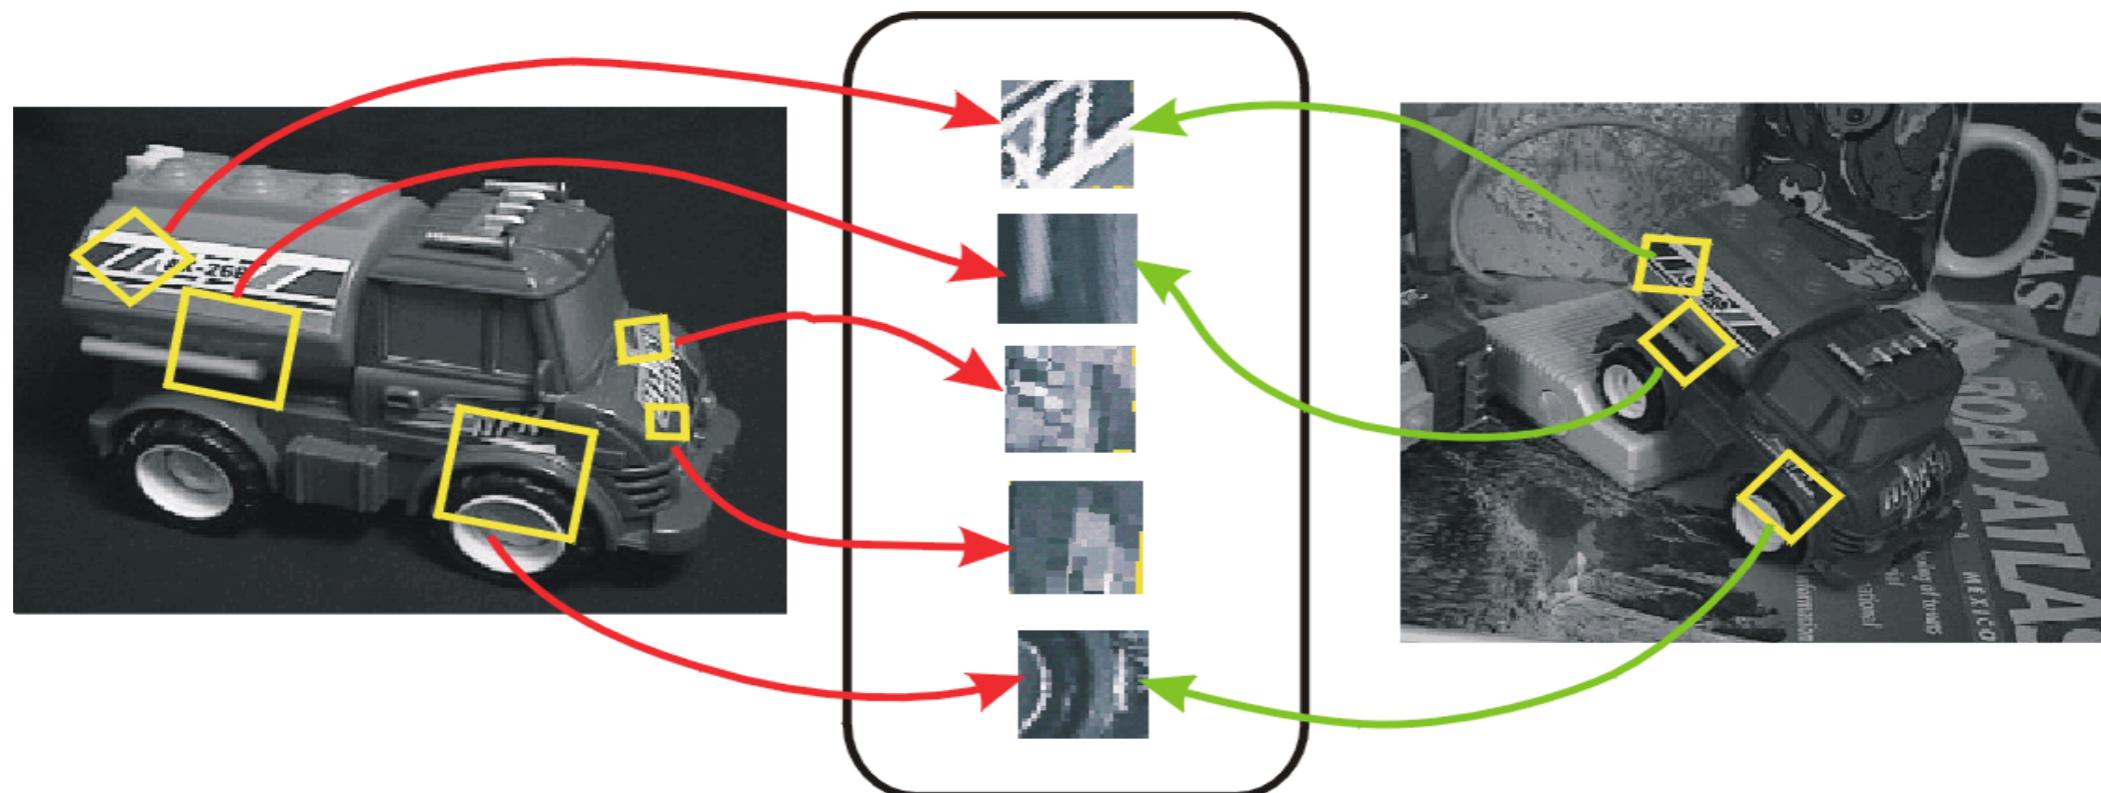


- Scaled and rotated versions of the same neighborhood will give rise to blobs that are related by the same transformation
- What to do if we want to compare the appearance of these image regions?
  - Normalization: transform these regions into same-size circles
  - Problem: **rotational ambiguity**



# Invariant local features

- Image content is transformed into local feature that are invariant to *translation*, *rotation*, *scale*, and other imaging parameters



**Reference:** CVPR 2003 Tutorial on “Recognition and Matching Based on Local Invariant Features” by David Lowe

# Desired properties of feature descriptors

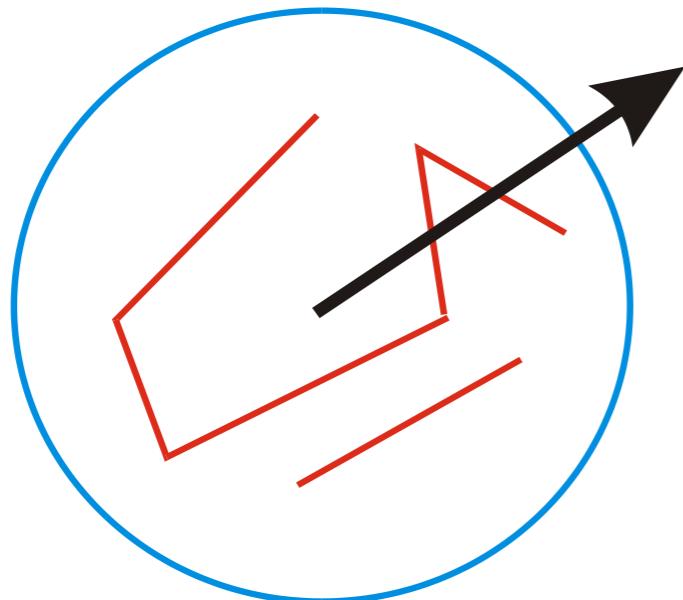
- Easily compared (compact, fixed-dimensional)
- Invariant:
  - Translation
  - Rotation
  - Scale
  - Change in image brightness (illumination)
  - *Change in perspective?*





# SIFT descriptors

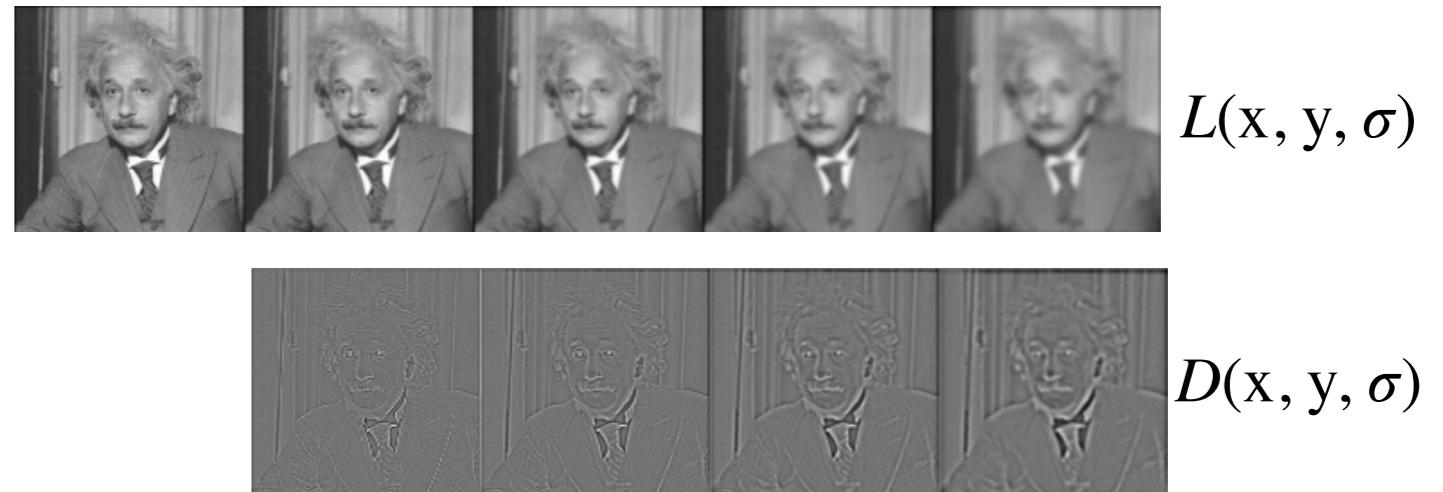
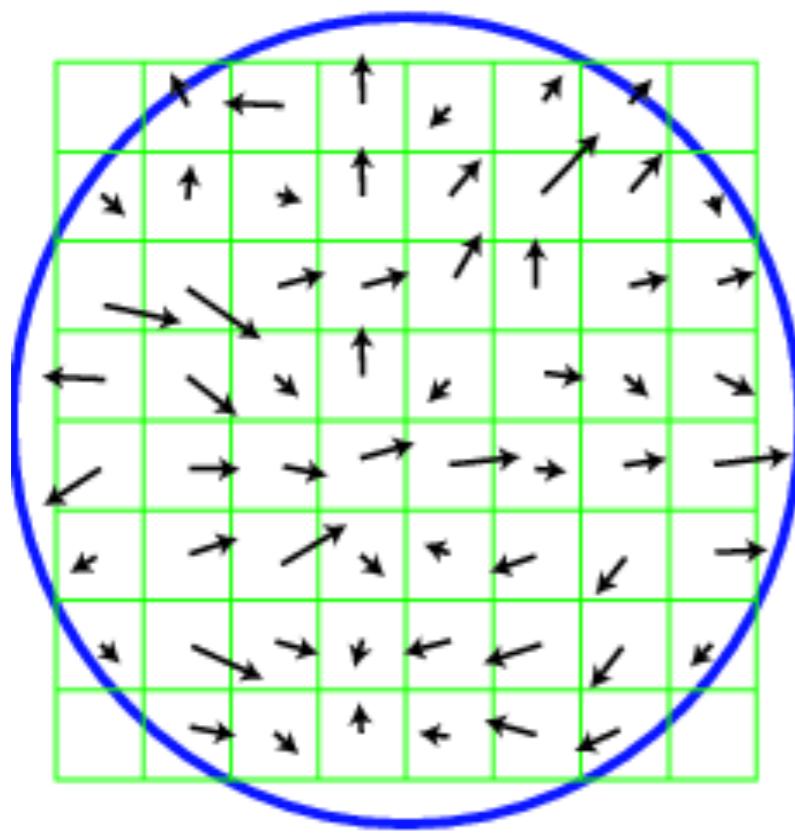
- We are given a keypoint and its scale from DoG
- To assign the orientation to circular image windows:
  - Create histogram of local gradient directions in the patch
  - Then we select a *characteristic orientation* for the keypoint based on the most prominent gradient (see next slides)
  - We will describe all features relative to this orientation



*If the keypoint appears rotated in another image, the features will be the same, because they're relative to the characteristic orientation*

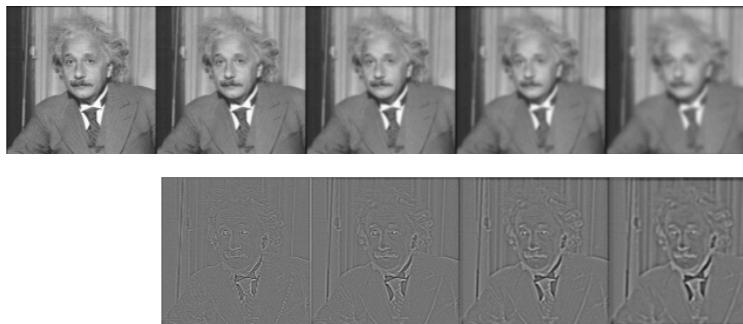
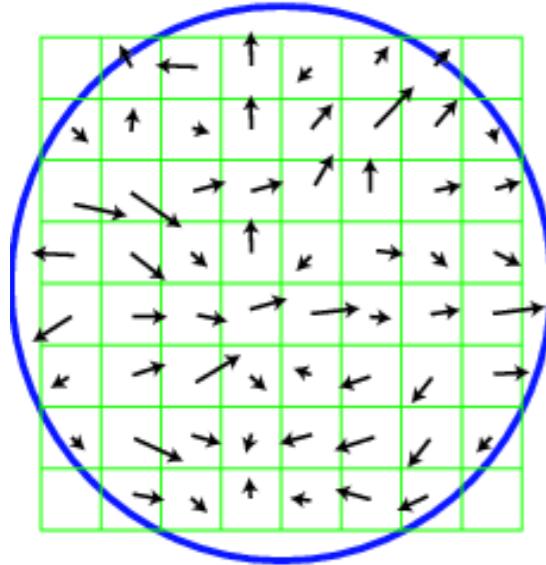


# SIFT descriptor formation



- Use the blurred image associated with the keypoint's scale
- Take image gradients over the keypoint neighborhood
- To become rotation invariant, rotate the gradient directions AND locations by keypoint orientation

# SIFT descriptor formation

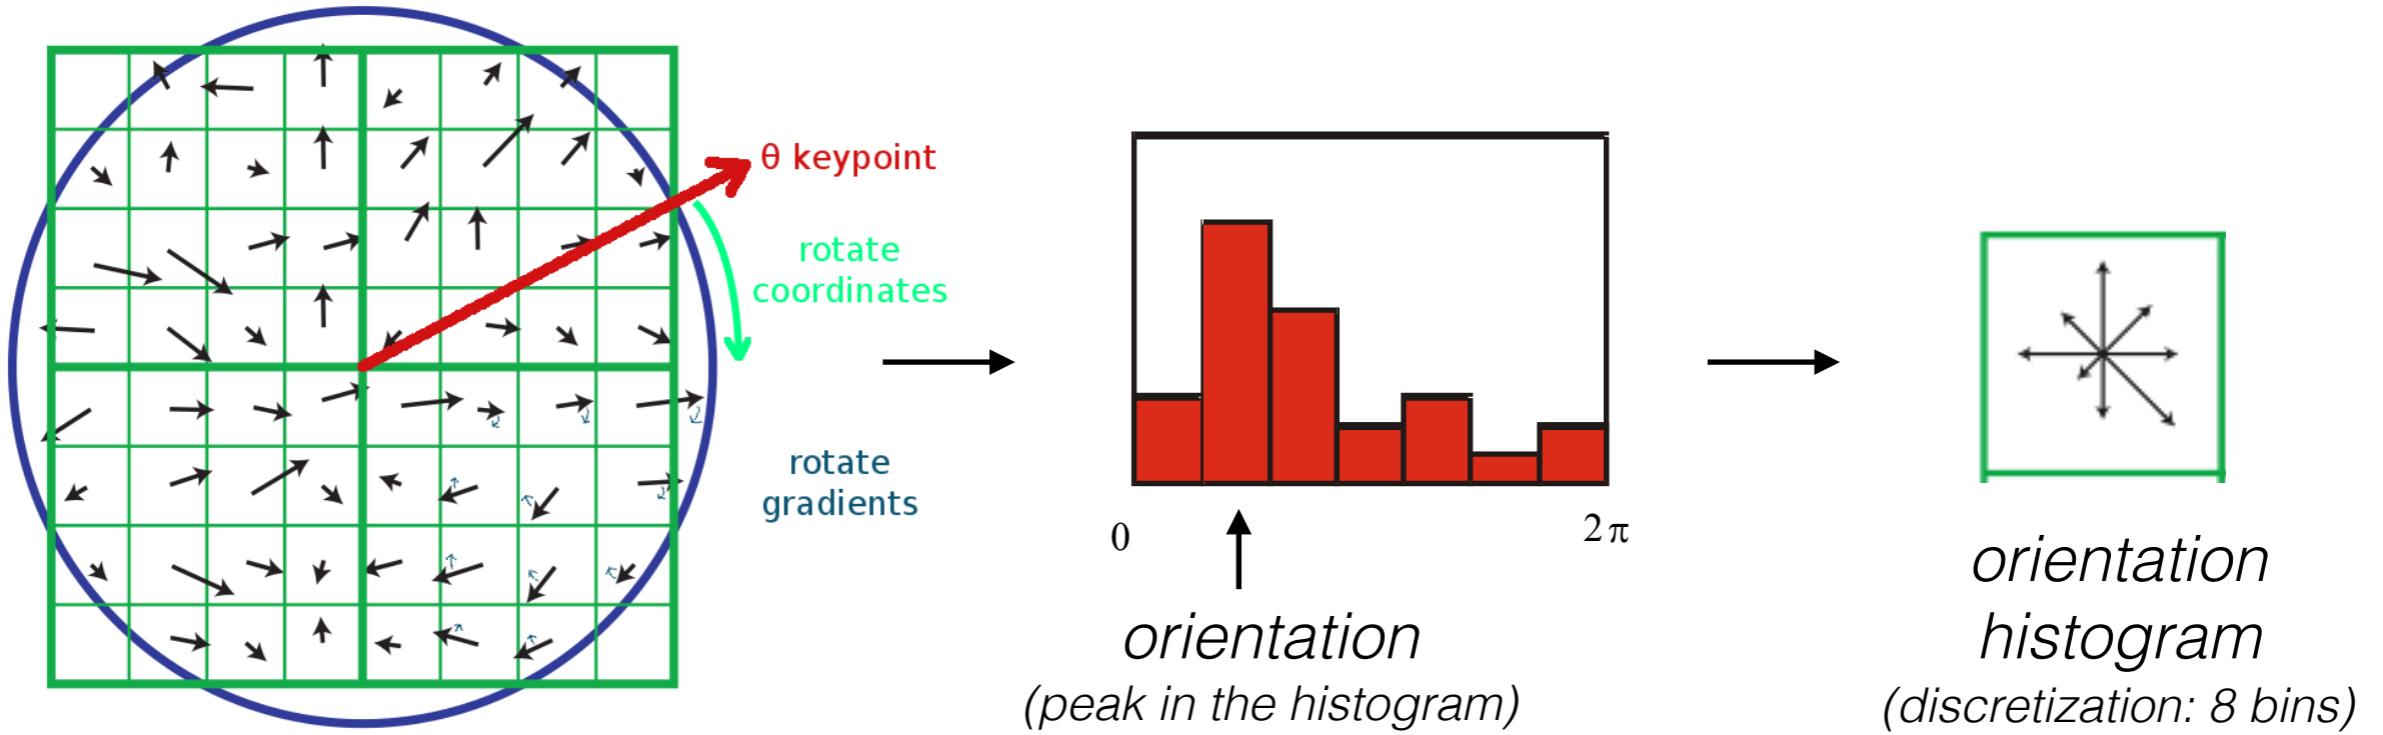


$$m(x, y) = \sqrt{(L(x + 1, y) - L(x - 1, y))^2 + (L(x, y + 1) - L(x, y - 1))^2}$$

$$\theta(x, y) = \text{atan2}(L(x, y + 1) - L(x, y - 1), L(x + 1, y) - L(x - 1, y))$$

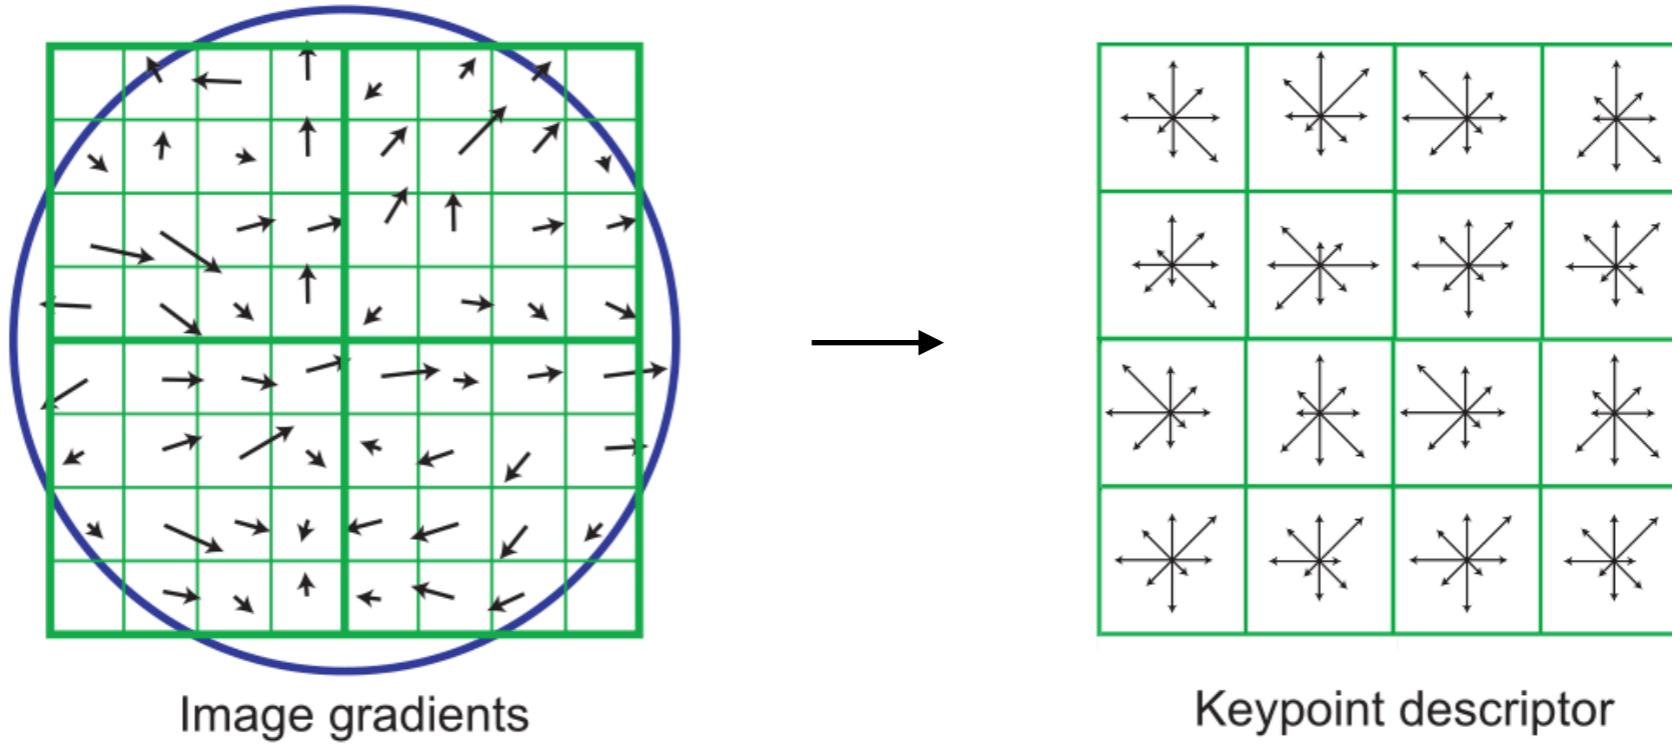
- Use the blurred image associated with the keypoint's scale
- Take image gradients over the keypoint neighborhood
- To become rotation invariant, rotate the gradient directions AND locations by keypoint orientation

# SIFT descriptor formation



- Use the blurred image associated with the keypoint's scale
- Take image gradients over the keypoint neighborhood
- To become rotation invariant, rotate the gradient directions AND locations by keypoint orientation

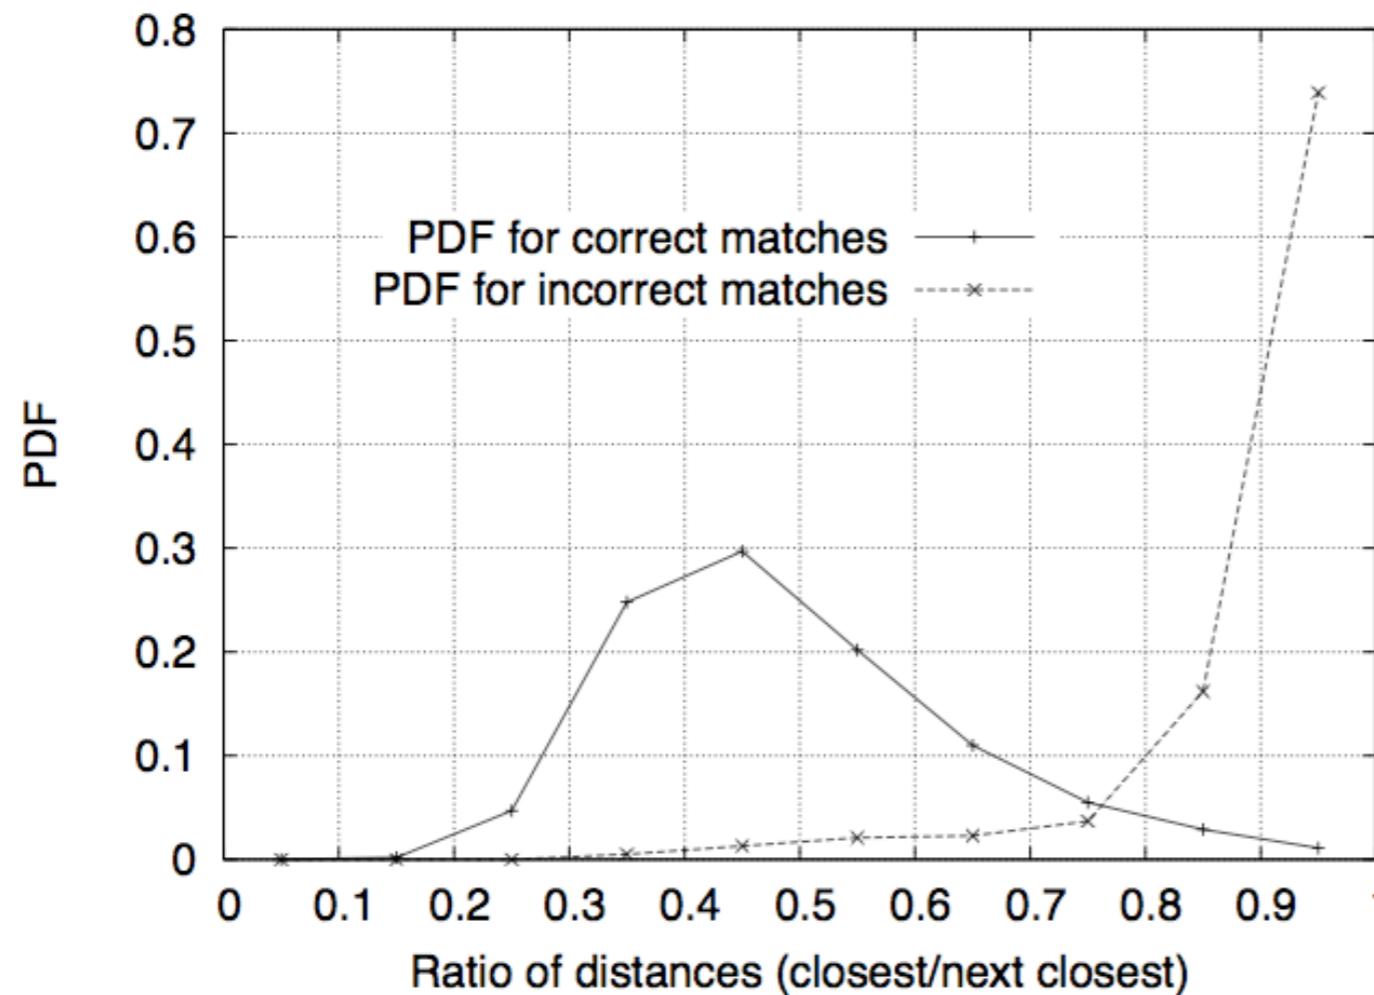
# SIFT descriptor formation



- 8 orientation bins per histogram and a  $4 \times 4$  histogram array yields  $8 \times 4 \times 4 = 128$  numbers
- So a SIFT descriptor is a 128-d vector that is invariant to rotation (because we rotated the descriptor) and scale (because of DoG)
- We can now compare each vector from image *A* to each vector from image *B* to find matching keypoints

# SIFT matching

- Euclidean distance between SIFT descriptor gives a good measure of keypoint similarity
- Ratio of distances is more reliable for matching 



The probability that a match is correct can be determined by taking the ratio of distance from the closest neighbor to the distance of the second closest  
*(plot: PDF obtained using a database of approx. 40,000 keypoints)*



# SIFT matching

- An example:



# Summary

- Edge detection
- Image gradients
- Local invariant features
- DoG and SIFT descriptors